


03-3 Toolchains

Embedded Linux system development

Cross-compiling toolchains

Thomas Petazzoni
Michael Opdenacker
Free Electrons

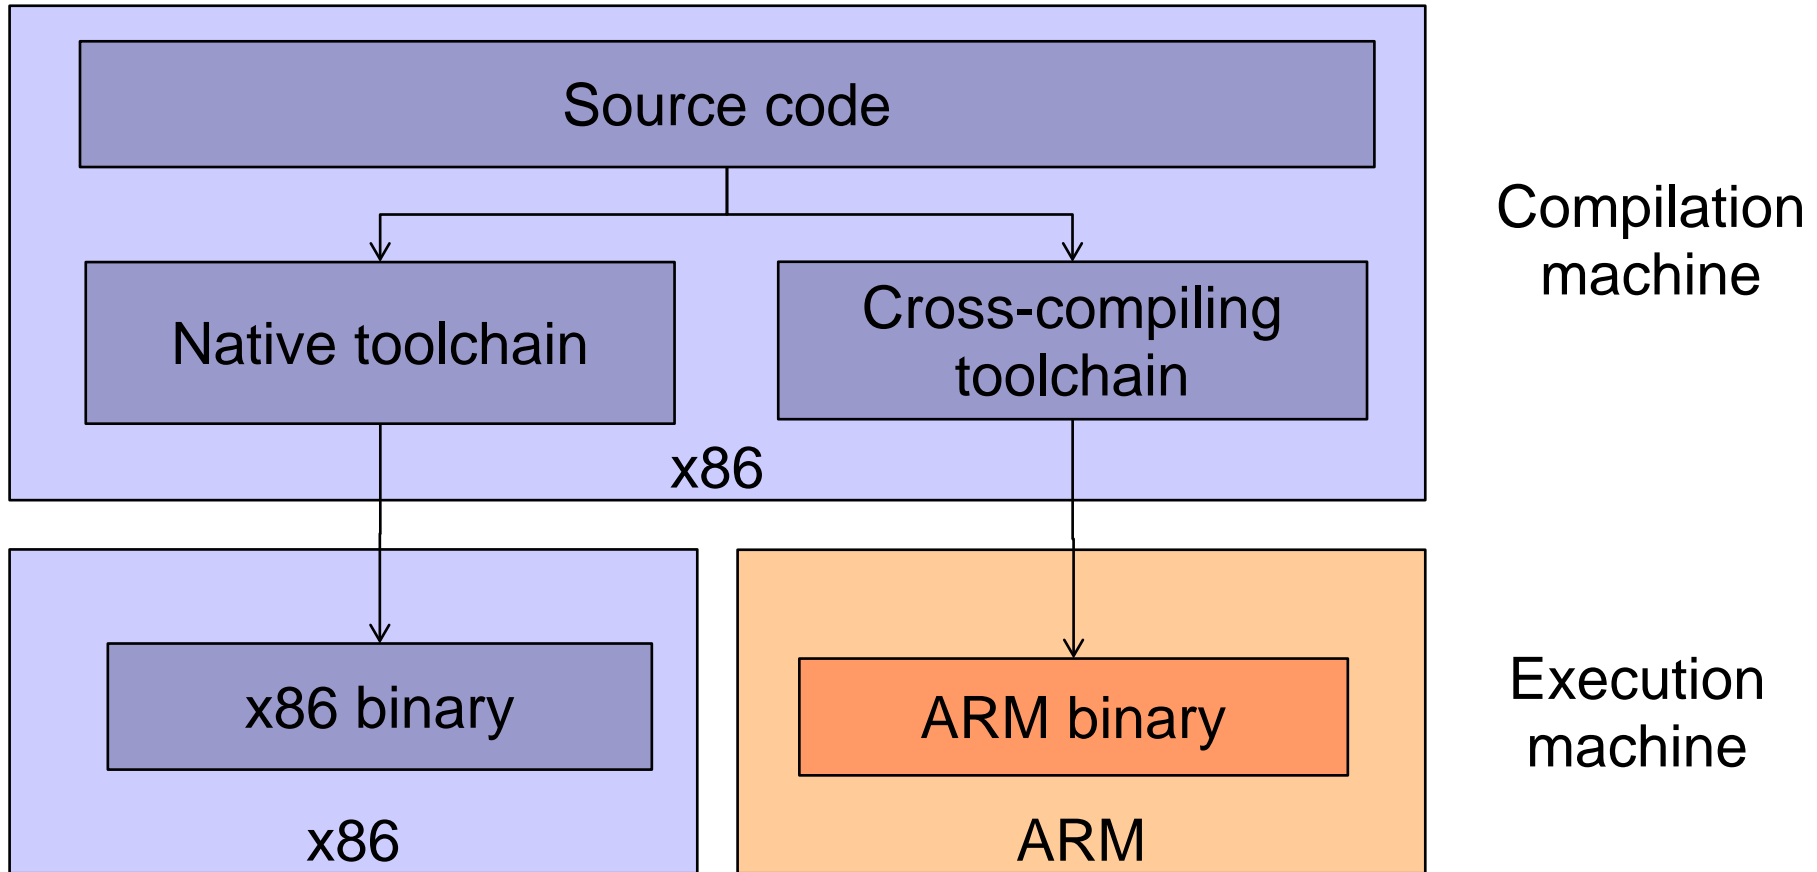


© Copyright 2004-2009, Free Electrons.
Creative Commons BY-SA 3.0 license
Latest update: 9/17/2018,
Document sources, updates and translations:
<http://free-electrons.com/docs/toolchains>
Corrections, suggestions, contributions and translations are
welcome!

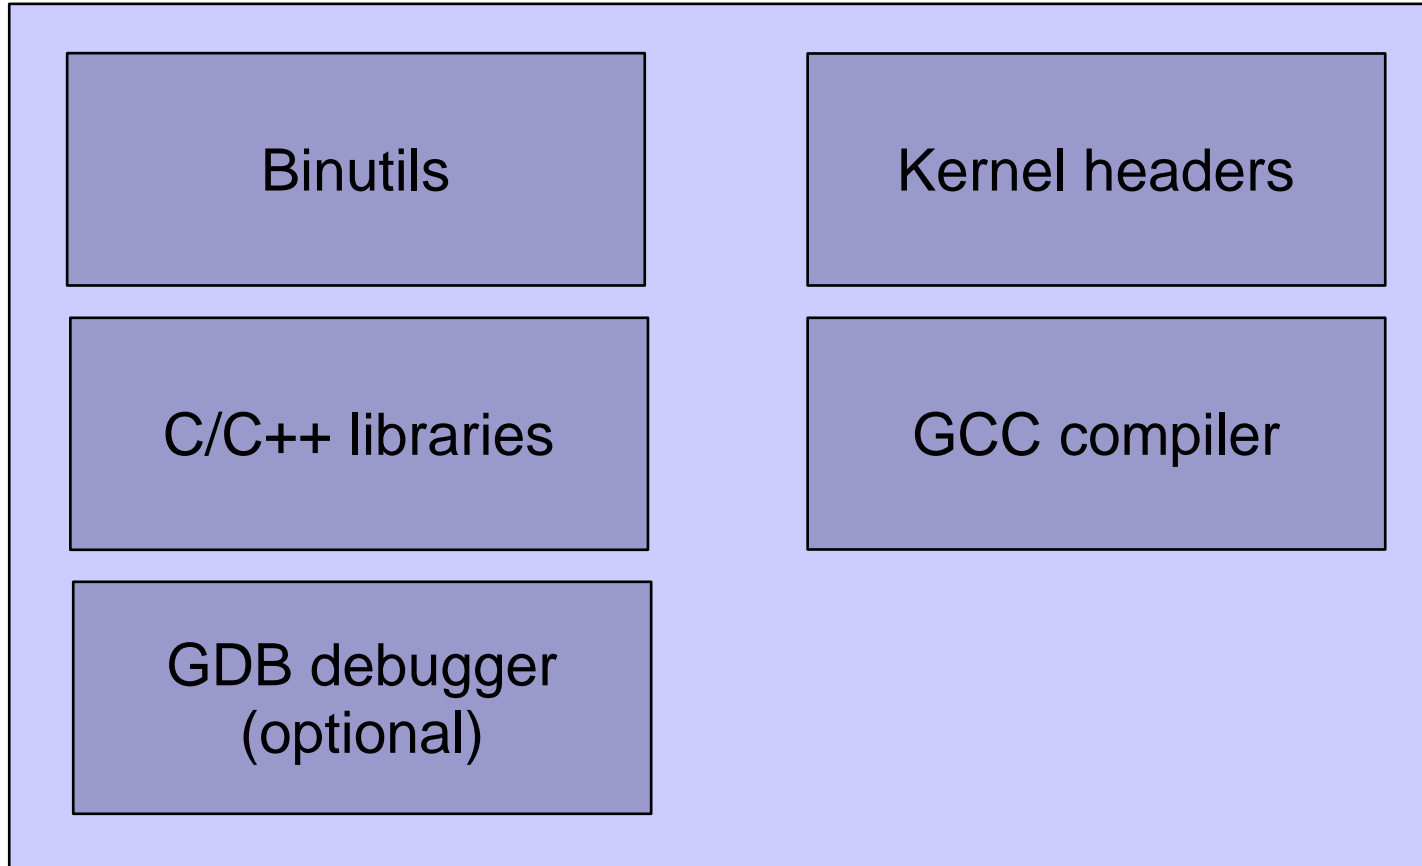
Definition (1)

- ▶ The usual development tools available on a GNU/Linux workstation is a **native toolchain**
- ▶ This toolchain runs on your workstation and generates code for your workstation, usually x86
- ▶ For embedded system development, it is usually impossible or not interesting to use a native toolchain
 - ▶ The target is too restricted in terms of storage and/or memory
 - ▶ The target is very slow compared to your workstation
 - ▶ You may not want to install all development tools on your target.
- ▶ Therefore, **cross-compiling toolchains** are generally used. They run on your workstation but generate code for your target.

Definition (2)



Components




binutils

- ▶ **Binutils** is a set of tools to generate and manipulate binaries for a given CPU architecture
- ▶ **as**, the assembler, that generates binary code from assembler source code
- ▶ **ld**, the linker
- ▶ **ar**, **ranlib**, to generate **.a** archives, used for libraries
- ▶ **objdump**, **readelf**, **size**, **nm**, **strings**, to inspect binaries. Very useful analysis tools !
- ▶ **strip**, to strip useless parts of binaries in order to reduce their size
- ▶ <http://www.gnu.org/software/binutils/>
- ▶ GPL license

Kernel headers (1)

- ▶ The C library and compiled programs need to interact with the kernel
 - ▶ Available system calls and their numbers
 - ▶ Constant definitions
 - ▶ Data structures, etc.
- ▶ Therefore, compiling the C library requires kernel headers, and many applications also require them.
- ▶ Available in `<linux/...>` and `<asm-generic/...>` and a few other directories corresponding to the ones visible in `/usr/include/` in the kernel sources



/usr/include

```
bone$ cd /usr/include
```

```
bone$ ls arm-linux-gnueabi/hf/asm
```

auxvec.h	hwcap.h	kvm_para.h	poll.h	setup.h	socket.h	termbits.h
bitsperlong.h	ioctl.h	mman.h	posix_types.h	shmbuf.h	sockios.h	termios.h
byteorder.h	ioctl.h	msgbuf.h	ptrace.h	sigcontext.h	statfs.h	types.h
errno.h	ipcbuf.h	param.h	resource.h	siginfo.h	stat.h	unistd.h
fcntl.h	kvm.h	perf_regs.h	sembuf.h	signal.h	swab.h	

```
bone$ ls asm-generic
```

auxvec.h	int-l64.h	kvm_para.h	poll.h	shmbuf.h	socket.h	termbits.h
bitsperlong.h	int-l164.h	mman-common.h	posix_types.h	shmparam.h	sockios.h	termios.h
errno-base.h	ioctl.h	mman.h	resource.h	siginfo.h	statfs.h	types.h
errno.h	ioctl.h	msgbuf.h	sembuf.h	signal-defs.h	stat.h	ucontext.h
fcntl.h	ipcbuf.h	param.h	setup.h	signal.h	swab.h	unistd.h

Kernel headers (2)

- ▶ System call numbers, in `</include/asm/unistd.h>`

```
#define __NR_exit      1
#define __NR_fork      2
#define __NR_read      3
```

- ▶ Constant definitions, here in `</include/asm-generic/fcntl.h>`, included from `</include/asm/fcntl.h>`, included from `</include/linux/fcntl.h>`

```
#define O_RDWR      00000002
```

- ▶ Data structures, here in `</include/asm/stat.h>`

```
struct stat {
    unsigned long  st_dev;
    unsigned long  st_ino;
    [...]
};
```

Kernel headers (3)

- ▶ The kernel-to-userspace ABI is backward compatible
- ▶ Binaries generated with a toolchain using kernel headers older than the running kernel will work without problem, but won't be able to use the new system calls, data structures, etc.
- ▶ Binaries generated with a toolchain using kernel headers newer than the running kernel might work on if they don't use the recent features, otherwise they will break
- ▶ Using the latest kernel headers is not necessary, unless access to the new kernel features is needed
- ▶ The kernel headers are extracted from the kernel sources using the `headers_install` kernel `Makefile` target.


GCC compiler

- ▶ GNU C Compiler, the famous free software compiler
- ▶ Can compile C, C++, Ada, Fortran, Java, Objective-C, Objective-C++, and
- ▶ Generate code for a large number of CPU architectures, including **ARM**, AVR, Blackfin, CRIS, FRV, M32, MIPS, MN10300, PowerPC, SH, v850, i386, x86_64, IA64, Xtensa, etc.
- ▶ <http://gcc.gnu.org/>
- ▶ Available under the GPL license, libraries under the LGPL.



C library

- ▶ The C library is an essential component of a Linux system
- ▶ Interface between the applications and the kernel
- ▶ Provides the well-known standard C API to ease application development
- ▶ Several C libraries are available: `glibc`, `uClibc`, `eglibc`, `dietlibc`, `newlib`, etc.
- ▶ The choice of the C library must be made at the time of the cross-compiling toolchain generation, as the GCC compiler is compiled against a specific C library.



glibc

<http://www.gnu.org/software/libc/>



- ▶ License: LGPL
- ▶ C library from the **GNU** project
- ▶ Designed for performance, standards compliance and portability
- ▶ Found on all **GNU / Linux** host systems
- ▶ Of course, actively maintained
- ▶ Quite big for small embedded systems: approx 2.5 MB on **arm** (version 2.9 - **libc**: 1.5 MB, **libm**: 750 KB)
- ▶ 2016-08-05: [glibc 2.24 released.](#)
- ▶ 2018-08-01: [glibc 2.28 released](#)

uClibc

<http://www.uclibc.org/> from CodePoet Consulting

- ▶ License: LGPL
- ▶ Lightweight C library for small embedded systems
- ▶ High configurability: many features can be enabled or disabled through a [menuconfig](#) interface
- ▶ Works only with Linux/uClinux, works on most embedded architectures
- ▶ No stable ABI, different ABI depending on the library configuration
- ▶ Focus on size rather than performance
- ▶ Small compile time

uClibc (2)

- ▶ Most of the applications compile with uClibc. This applies to all applications used in embedded systems.
- ▶ Size ([arm](#)): 4 times smaller than [glibc](#)!
[uClibc 0.9.30.1](#): approx. 600 KB ([libuClibc](#): 460 KB, [libm](#): 96KB)
[glibc 2.9](#): approx 2.5 MB
- ▶ Used on a large number of production embedded products, including consumer electronic devices
- ▶ Actively maintained, large developer and user base
- ▶ Now supported by [MontaVista](#), [TimeSys](#) and [Wind River](#).
- ▶ **15 May 2012, uClibc 0.9.33.2 Released**

Honey, I shrunk the programs!

<i>C program</i>	<i>Compiled with shared libraries</i>		<i>Compiled statically</i>	
	<i>glibc</i>	<i>uClibc</i>	<i>glibc</i>	<i>uClibc</i>
Plain “hello world” (stripped)	5.6 K (glibc 2.9)	5.4 K (uClibc 0.9.30.1)	472 K (glibc 2.9)	18 K (uClibc 0.9.30.1)
Busybox (stripped)	245 K (older glibc)	231 K (older uClibc)	843 K (older glibc)	311 K (older uClibc)

Executable size comparison on ARM

Installing and using a precompiled toolchain

- ▶ Follow the installation procedure proposed by the vendor
- ▶ Usually, it is simply a matter of extracting a tarball at the proper place
- ▶ Then, add the path to toolchain binaries in your **PATH**:
`export PATH=/path/to/toolchain/bin/:$PATH`

Or

```
host$ export ARCH=arm
host$ export CROSS_COMPILE= arm-linux-gnueabihf-
host$ PATH=$PATH:~/BeagleBoard/bb-
kernel/dl/gcc-linaro-5.3-2016.02-x86_64_arm-
linux-gnueabihf/bin
host$ ${CROSS_COMPILE}gcc helloWorld.c
```