

## 04-2 Adding to the Kernel, Kernel Initialization

# Adding to the Kernel

- Makefile Targets
- Kernel Configuration
- Custom Configuration Options
- Kernel Makefiles
- Kernel Documentation

# Composite Kernel Image

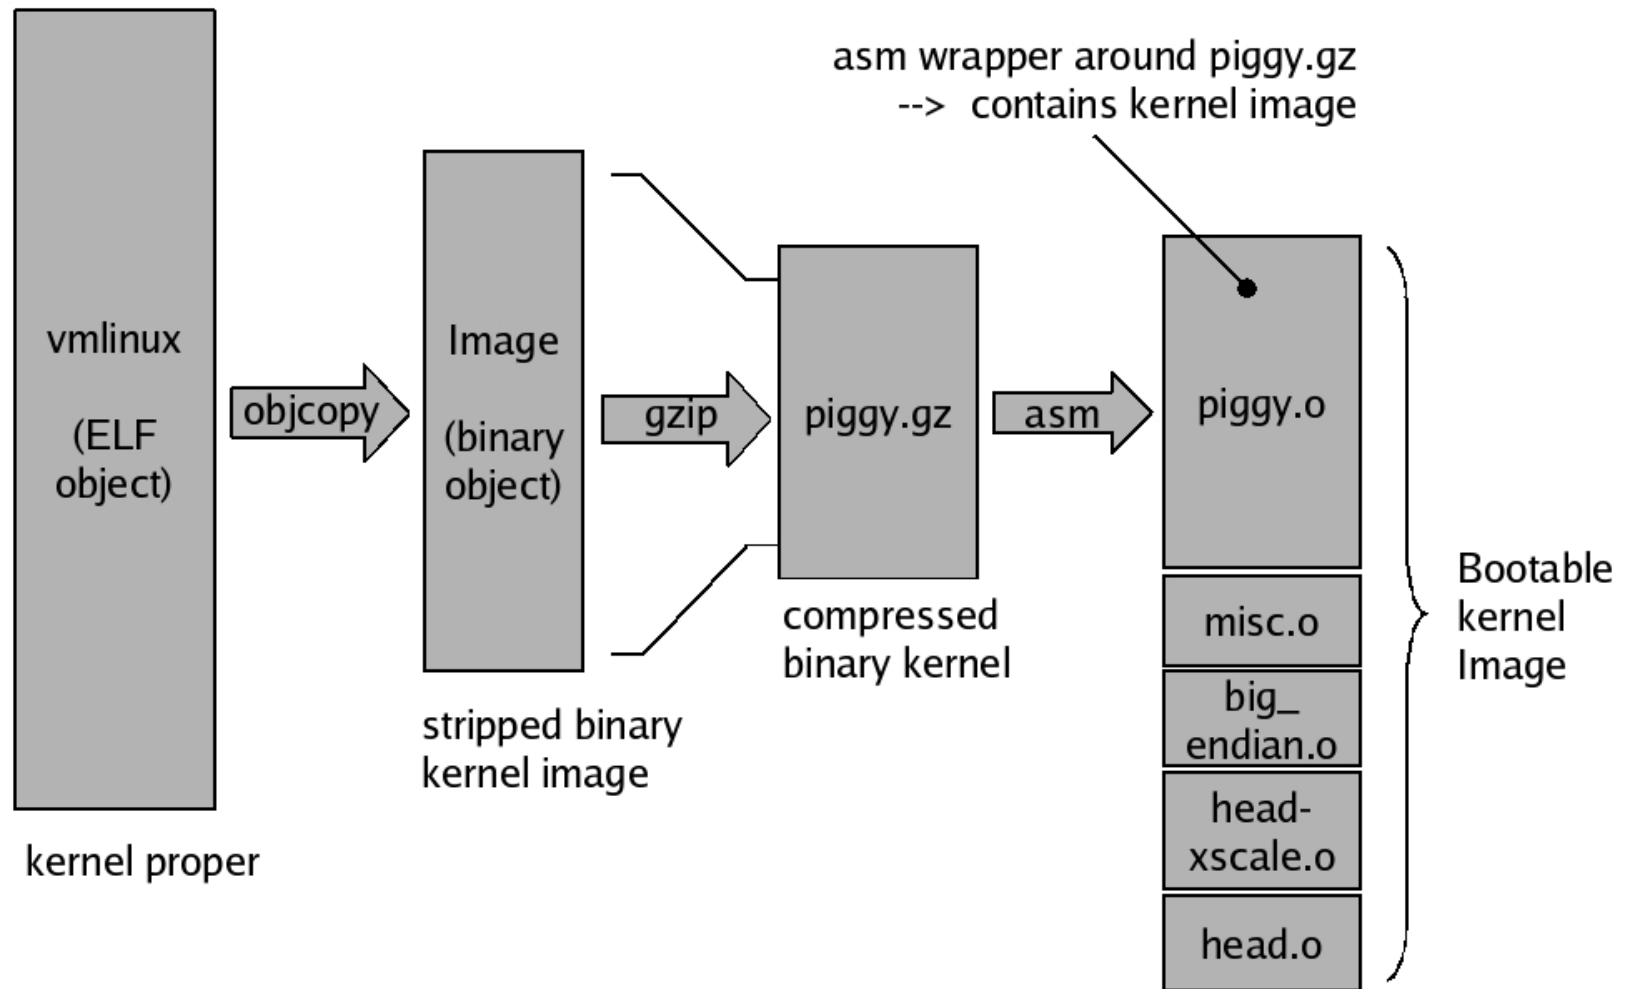


Figure 5.1 page 103

# piggy.S

3.2 kernel

```
.section .piggydata,#alloc
.globl    input_data
input_data:
.incbin   "arch/arm/boot/compressed/piggy.gz"
.globl    input_data_end
input_data_end:
```

How do you find this file?

```
host$ cd bb-kernel/KERNEL
host$ find . -iname piggy.s
```

```
.incbin "arch/arm/boot/compressed/piggy_data"
```

# Compiling Kernel

```
host$ source ~/crossCompileEnv.sh
host$ make -j3 uImage
... < many build steps omitted for clarity >
AS      arch/arm/boot/compressed/head.o
        XZKERN arch/arm/boot/compressed/piggy.xzkern
...
AS      arch/arm/boot/compressed/piggy.xzkern.o
LD      arch/arm/boot/compressed/vmlinux
OBJCOPY arch/arm/boot/zImage
Kernel: arch/arm/boot/zImage is ready
UIMAGE arch/arm/boot/uImage
Image Name:   Linux-3.8.13+
Created:      Thu Oct  3 17:13:18 2013
Image Type:   ARM Linux Kernel Image (uncompressed)
Data Size:    2898464 Bytes = 2830.53 kB = 2.76 MB
Load Address: 80008000
Entry Point:  80008000
Image arch/arm/boot/uImage is ready
```

# .../arch/arm/boot/compressed

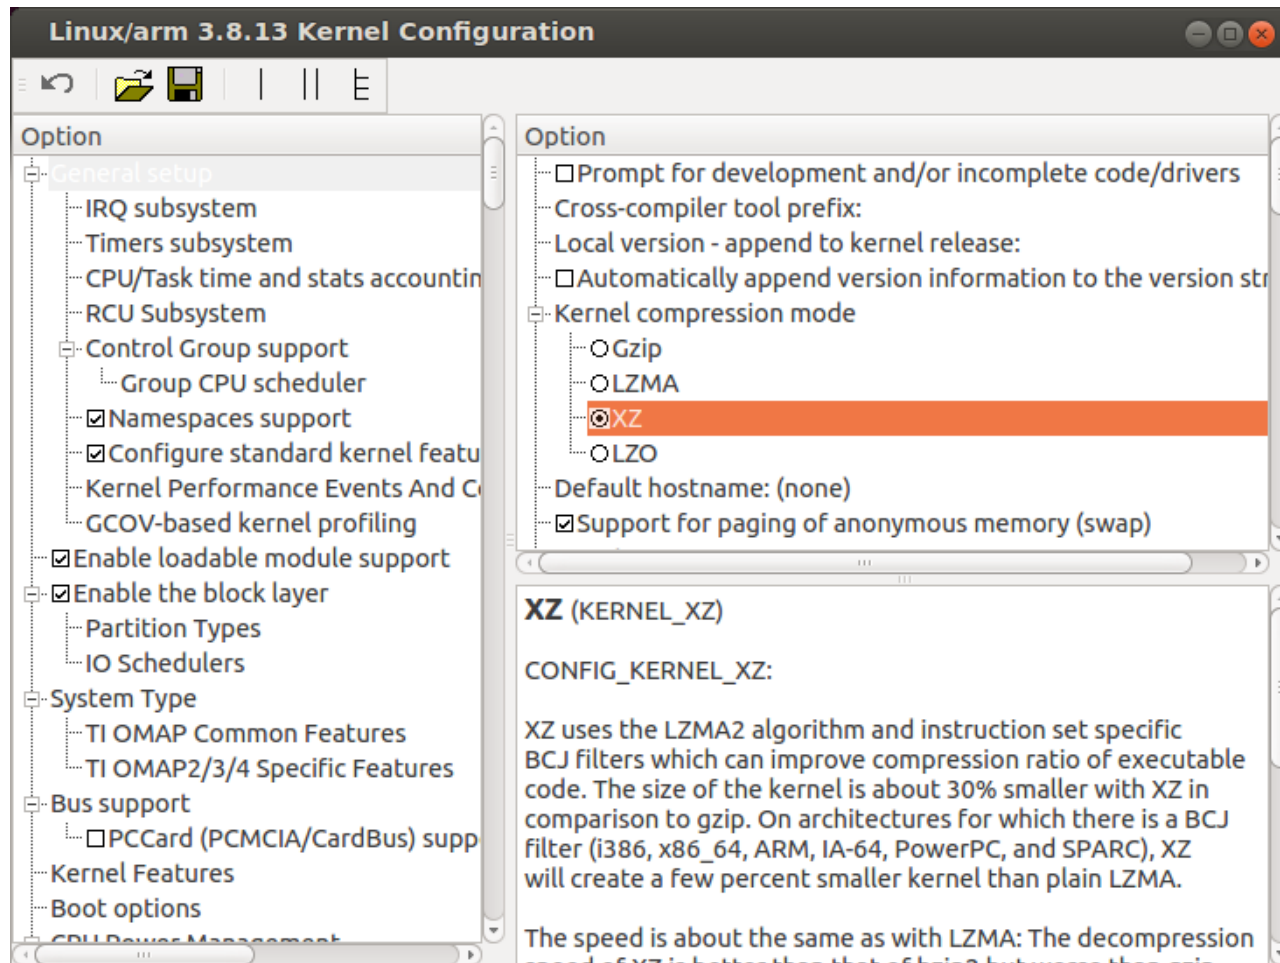
host\$ **ls**

ashldi3.o	hyp-stub.o	piggy.lzo.S
ashldi3.S	hyp-stub.S	<b>piggy.xzkern</b>
atags_to_fdt.c	liblfuncs.o	<b>piggy.xzkern.o</b>
big-endian.S	liblfuncs.S	<b>piggy.xzkern.S</b>
decompress.c	libfdt_env.h	sdhi-sh7372.c
decompress.o	ll_char_wr.S	sdhi-shmobile.c
head.o	Makefile	sdhi-shmobile.h
head.S	misc.c	string.c
head-sa1100.S	misc.o	string.o
head-shark.S	mmCIF-sh7372.c	vmlinux
head-sharpsl.S	ofw-shark.c	vmlinux.lds
head-shmobile.S	piggy.gzip.S	vmlinux.lds.in
head-xscale.S	piggy.lzma.S	

# piggy.xzkern.S

```
.section .piggydata,#alloc
.globl    input_data
input_data:
.incbin
    "arch/arm/boot/compressed/piggy.xzkern"
.globl    input_data_end
input_data_end:
```

# How does it know to use kernxz?





Linux/arm 4.12.0-rc4 Kernel Configuration

Option

(8) Maximum PAGE\_SIZE order of alignment for DMA IOMMU  
Patch physical to virtual translations at runtime

▼ General setup

IRQ subsystem  
Timers subsystem  
CPU/Task time and stats accounting  
RCU Subsystem

▼ Control Group support

CPU controller

☒ Namespaces support

Configure standard kernel features (expert users)

Kernel Performance Events And Counters

☐ GCC plugins

GCOV-based kernel profiling

☒ Enable loadable module support

▼ ☒ Enable the block layer

Partition Types

IO Schedulers

▼ System Type

Multiple platform selection

☐ Marvell Engineering Business Unit (MVEBU) SoCs

☐ Axis Communications ARM based ARTPEC SoCs

☐ Atmel SoCs

☐ Broadcom SoC Support

☐ Marvell Berlin SoCs

☐ Amlogic Meson SoCs

☐ Freescale i.MX family

☐ Mediatek MT65xx & MT81xx SoC

TI OMAP Common Features

▼ TI OMAP/AM/DM/DRA Family

TI OMAP2/3/4 Specific Features

Option

Cross-compiler tool prefix:  
☐ Compile also drivers which will not load  
Local version - append to kernel release:  
☐ Automatically append version information to the version

▼ Kernel compression mode

☐ Gzip

☐ LZMA

☐ XZ

☒ LZO

☐ LZ4

Default hostname: (none)

☒ Support for paging of anonymous memory (swap)

☒ System V IPC

☒ POSIX Message Queues

LZO (KERNEL\_LZO)

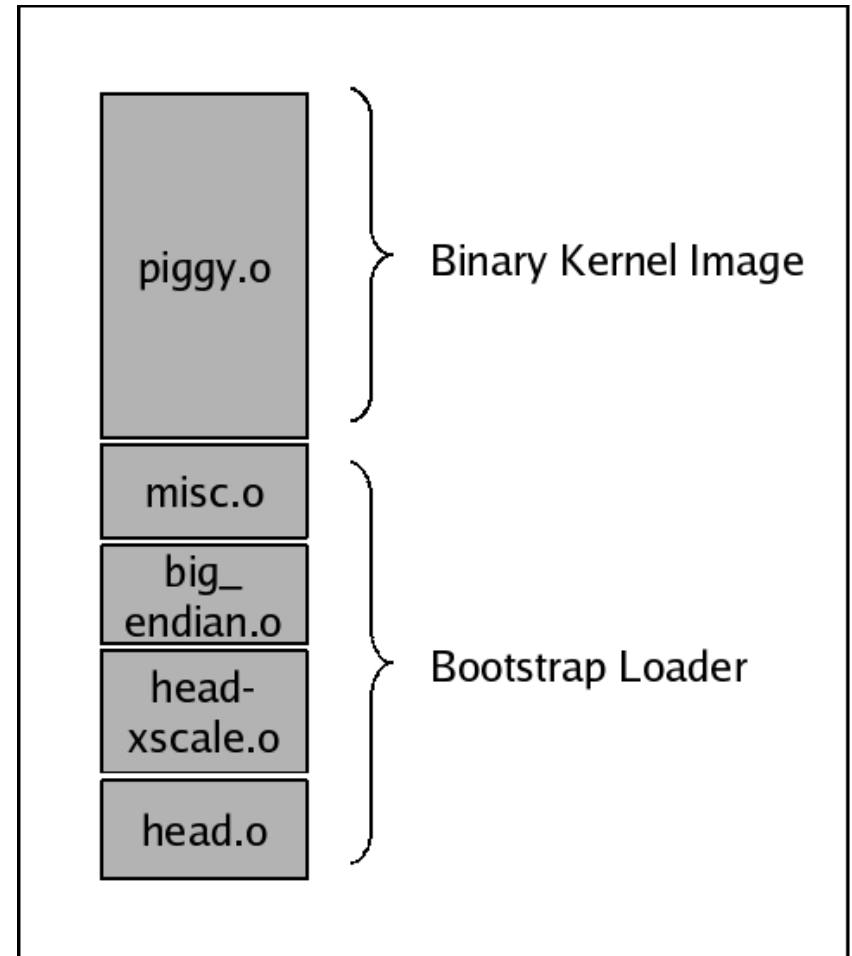
CONFIG\_KERNEL\_LZO:

Its compression ratio is the poorest among the choices. The kernel size is about 10% bigger than gzip; however its speed (both compression and decompression) is the fastest.

Symbol: KERNEL\_LZO [=y]  
Type : boolean  
Prompt: LZO  
Location:  
-> General setup  
-> Kernel compression mode (<choice> [=y])  
Defined at init/Kconfig:192  
Depends on: <choice> && HAVE\_KERNEL\_LZO [=y]

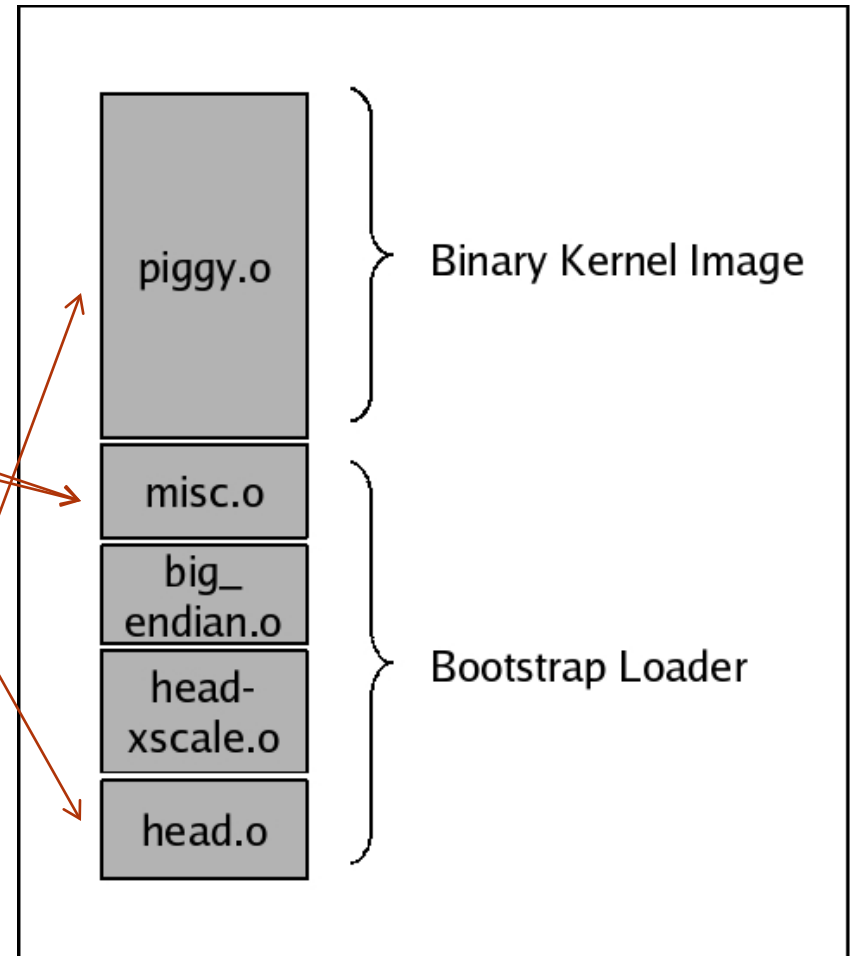
# Bootstrap Loader (not bootloader)

- Provide context for kernel
  - Enable instruction set
  - Data caches
  - Disable interrupt
  - C runtime environment
- Decompress (misc.o)
- Relocate kernel image



# Bootstrap Loader (not bootloader)

```
LD      vmlinux
        SORTEX vmlinux
sort done marker at 81c420
        SYSMAP  System.map
        OBJCOPY arch/arm/boot/Image
Kernel: arch/arm/boot/Image is ready
AS      arch/arm/boot/compressed/head.o
XZKERN  arch/arm/boot/compressed/piggy.xzkern
CC      arch/arm/boot/compressed/misc.o
CC      arch/arm/boot/compressed/decompress.o
CC      arch/arm/boot/compressed/string.o
SHIPPED arch/arm/boot/compressed/hyp-stub.S
AS      arch/arm/boot/compressed/lib1funcs.o
AS      arch/arm/boot/compressed/ashldi3.o
AS      arch/arm/boot/compressed/hyp-stub.o
AS      arch/arm/boot/compressed/piggy.xzkern.o
LD      arch/arm/boot/compressed/vmlinux
        OBJCOPY arch/arm/boot/zImage
Kernel: arch/arm/boot/zImage is ready
        UIMAGE arch/arm/boot/uImage
```

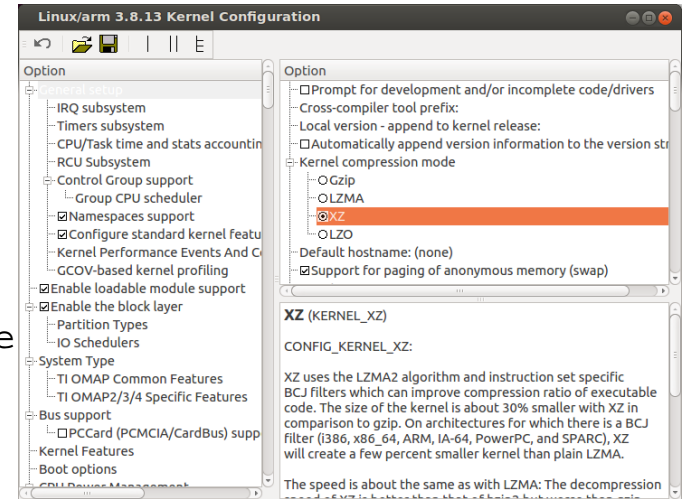


```
#ifndef CONFIG_KERNEL_GZIP
#include "../../../lib/decompress_inflate.c"
#endif

#ifdef CONFIG_KERNEL_LZO
#include "../../../lib/decompress_unlzo.c"
#endif

#ifdef CONFIG_KERNEL_LZMA
#include "../../../lib/decompress_unlzma.c"
#endif

#ifdef CONFIG_KERNEL_XZ
#define memmove memmove
#define memcpy memcpy
#include "../../../lib/decompress_unxz.c"
#endif
```



# Boot Messages

- See handout
- Note *kernel version string*
- Note *kernel command line*
- *EBC Boot Sequence* shows how to display the messages in the handout

```
bone$ cd /boot
```

```
bone$ ls -F
```

```
config-4.4.15-bone11      initrd.img-4.4.21-ti-r47  vmlinuz-4.4.15-bone11*  
config-4.4.19-ti-r41      SOC.sh                    vmlinuz-4.4.19-ti-r41*  
config-4.4.21-ti-r47      System.map-4.4.19-ti-r41  vmlinuz-4.4.21-ti-r47*  
config-4.4.22-bone13.1    System.map-4.4.21-ti-r47  vmlinuz-4.4.22-bone13.1*  
dtbs/                     uboot/  
initrd.img-4.4.19-ti-r41  uEnv.txt
```

```
bone$ cat uEnv.txt
```

```
#Docs: http://elinux.org/Beagleboard:U-boot\_partitioning\_layout\_2.0
```

```
uname_r=4.4.22-bone13.1
```

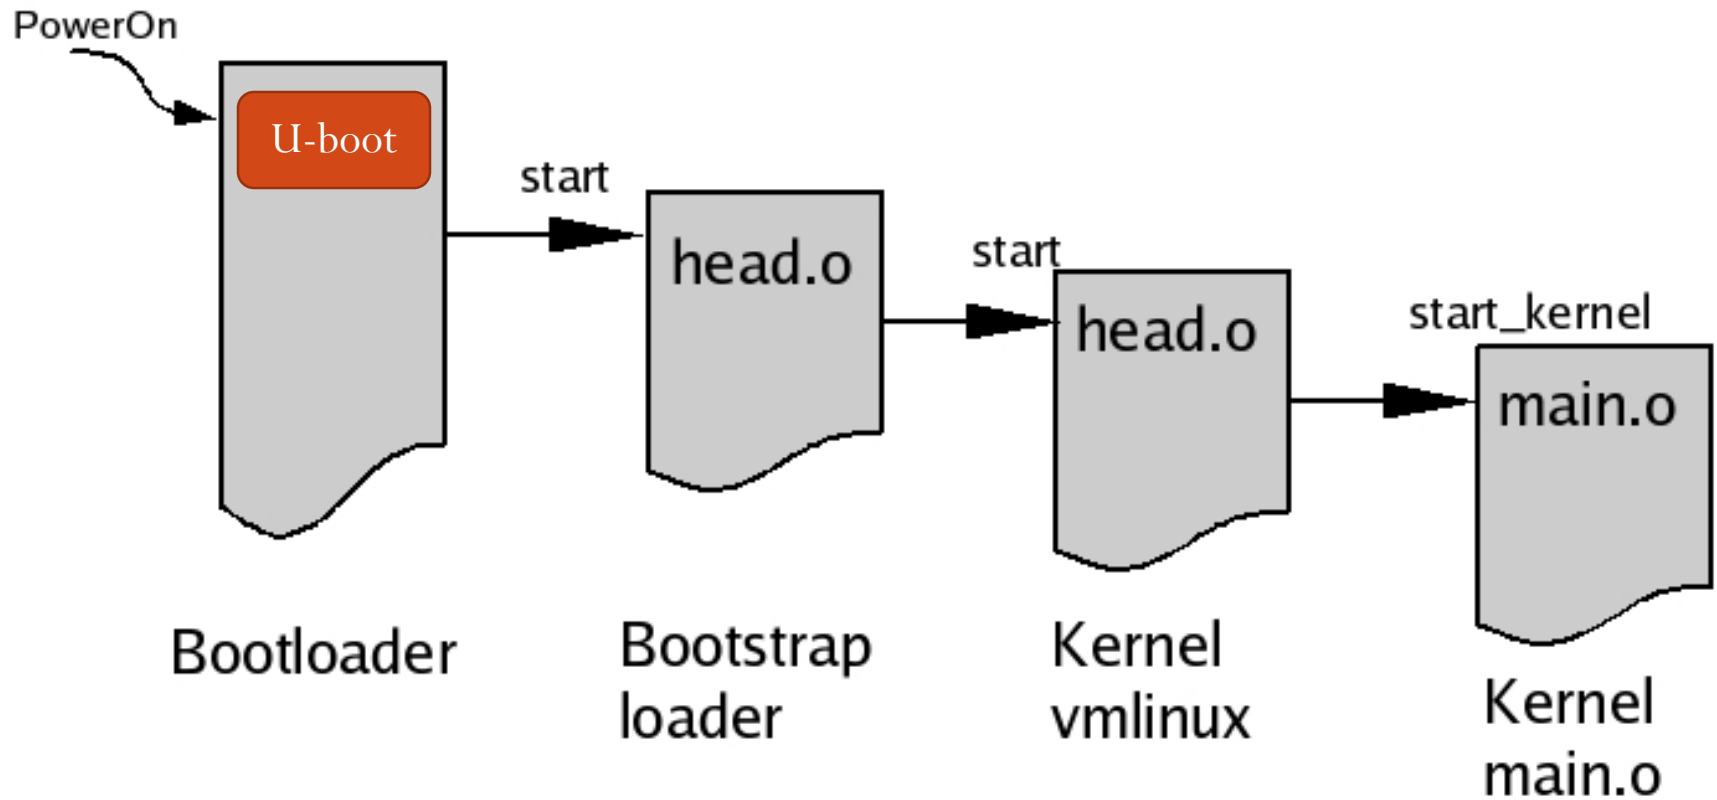
```
...
```

```
cmdline=coherent_pool=1M quiet cape_universal=enable
```





remove

## 5-3 ARM boot control flow



# Booting Linux – ROM to Kernel

MLO



# .../arch/arm/boot/compressed/head.S

```
#include <linux/linkage.h>

#ifdef DEBUG

#if defined(CONFIG_DEBUG_LL)
    .macro    loadsp, rb
    .endm

    .macro    writeb, ch, rb
    mcr      p14, 0, \ch, c0, c5, 0
    .endm
#else
    .macro    loadsp, rb
    .endm

    .macro    writeb, ch, rb
    mcr      p14, 0, \ch, c1, c0, 0
    .endm
#endif

    .macro    writeb,    ch, rb
    senduart \ch, \rb
    .endm

    .macro    loadsp, rb
    mov      \rb, #0x80000000    @
    physical base address
#endif

#include <mach/debug-macro.S>

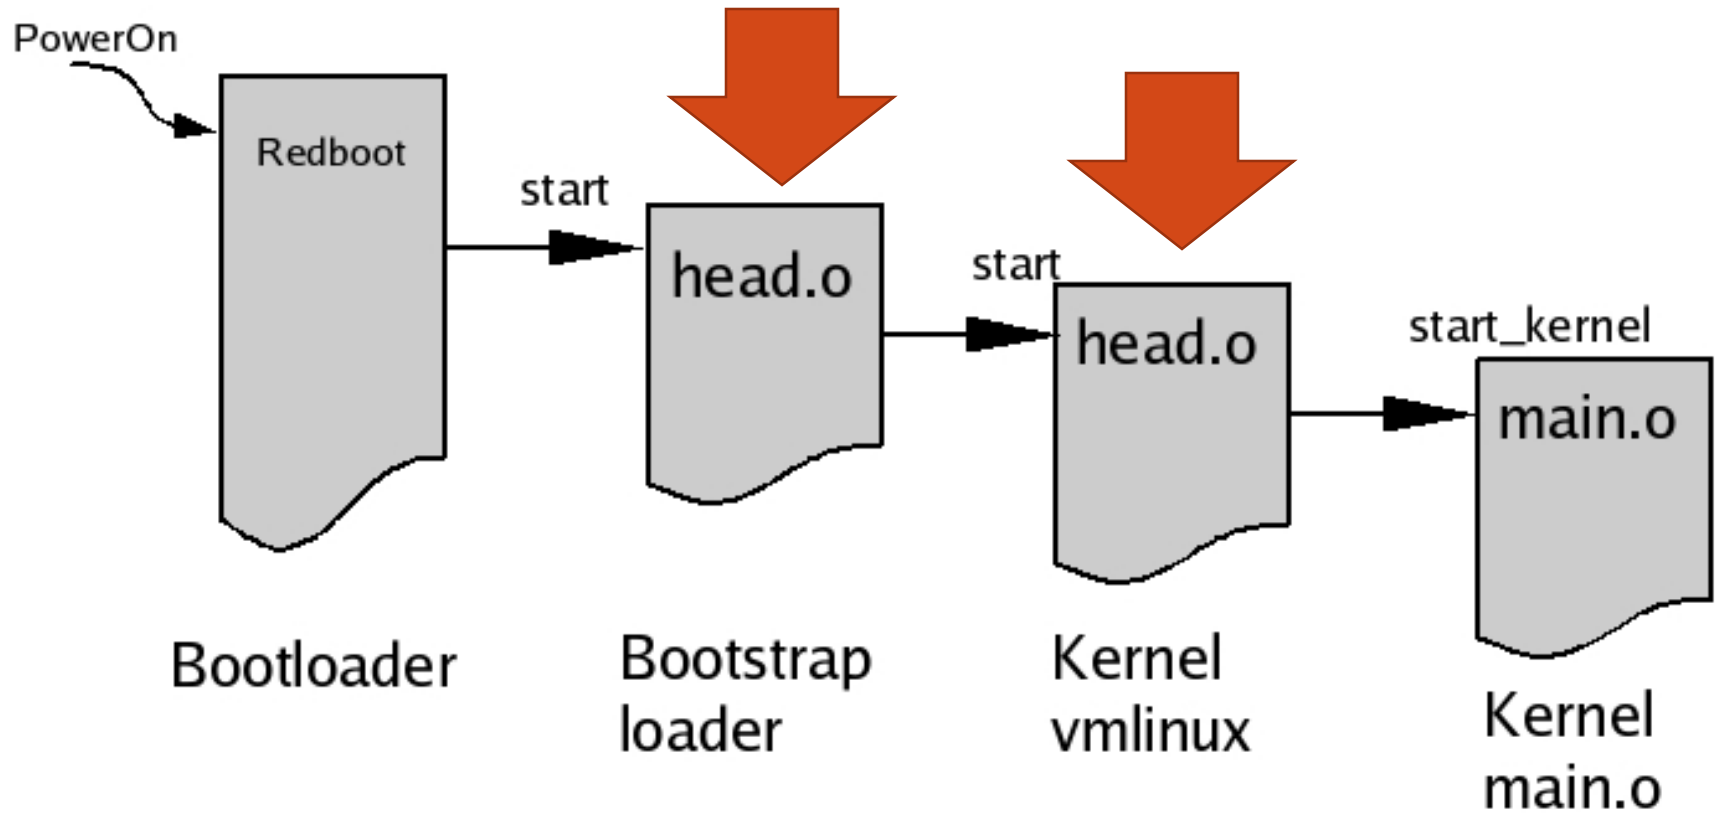
    .macro    writeb,    ch, rb
    senduart \ch, \rb
    .endm

    .macro    loadsp, rb
    mov      \rb, #0x80000000    @
    physical base address
#endif
```

How do you find the value?



# 2 head.o's



## .../arch/arm/kernel/head.S

1. Checks for valid processor and architecture
2. Creates initial page table entries
3. Enables the processor's memory management unit (MMU)
4. Establishes limited error detection and reporting
5. Jumps to the start of the kernel proper,  
**start\_kernel( )** in **main.c**.

Find these on the handout

# .../arch/arm/kernel/head.S

```
/*
 * Kernel startup entry point.
 * -----
 *
 * This is normally called from the decompressor code. The requirements
 * are: MMU = off, D-cache = off, I-cache = dont care, r0 = 0,
 * r1 = machine nr, r2 = atags or dtb pointer.
 *
 * This code is mostly position independent, so if you link the kernel at
 * 0xc0008000, you call this at __pa(0xc0008000).
 *
 * See linux/arch/arm/tools/mach-types for the complete list of machine
 * numbers for r1.
 *
 * We're trying to keep crap to a minimum; DO NOT add any machine specific
 * crap here - that's what the boot loader (or in extreme, well justified
 * circumstances, zImage) is for.
 */
```

# Kernel Startup

- **arch/arm/kernel/head.S**

**b start\_kernel**



# .../init/main.c

```
asmlinkage __visible void __init start_kernel(void)
{
    char *command_line;
    char *after_dashes;

    /*
     * Need to run as early as possible, to initialize the
     * lockdep hash:
     */
    lockdep_init();
    set_task_stack_end_magic(&init_task);
    smp_setup_processor_id();
    debug_objects_early_init();

    /*
     * Set up the the initial canary ASAP:
     */
    boot_init_stack_canary();

    cgroup_init_early();

    local_irq_disable();
    early_boot_irqs_disabled = true;
```

# Kernel Command Line Processing

- Kernel Command-Line Processing
- The **\_\_setup** macro

```
console=tty0 console=ttyO0,115200n8  
root=/dev/mmcblk0p1 rootfstype=ext4  
rootwait coherent_pool=1M  
cape_universal=enable
```

# Console Setup Code Snippet

```
/*
 * Setup a list of consoles. Called from init/main.c
 */

#ifdef CONFIG_SERIAL_OMAP
    if (!strncmp(str, "tty0", 4) && '0' <= str[4] && '9' >= str[4]) {
        str[3] = '0';
        pr_warn("We are opening your eyes, assuming you want to
use an OMAP based serial driver and not a zeroMAP based one! ;)\n");
        pr_warn("Which means 'tty0%s' was changed to 'tty0%s'
automagically for your pleasure.\n", str+4, str+4);
    }
#endif

...
return 1;
}

__setup("console=", console_setup);
```

Registration  
function

From .../kernel/printk/printk.c

# Console Setup Code Snippet

```
/*
 * Setup a list of consoles. Called from init/main.c
 */
static int __init console_setup(char *str)
{
    char buf[sizeof(console_cmdline[0].name) + 4]; /* 4 for "ttyS" */
    char *s, *options, *brl_options = NULL;
    int idx;
    if (_braille_console_setup(&str, &brl_options))
        return 1;

    <body omitted for clarity...>
    ...
    return 1;
}
__setup("console=", console_setup);
```

From .../kernel/printk/printk.c



New



# .../include/linux/init.h

```
/*
 * Only for really core code.  See moduleparam.h for the normal way.
 *
 * Force the alignment so the compiler doesn't space elements of the
 * obs_kernel_param "array" too far apart in .init.setup.
 */
#define __setup_param(str, unique_id, fn, early) \
    static char __setup_str_##unique_id[] __initdata __aligned(1) = str; \
    \
    static struct obs_kernel_param __setup_##unique_id \
        __used __section(.init.setup) \
        __attribute__((aligned((sizeof(long))))) \
        = { __setup_str_##unique_id, fn, early }

#define __setup(str, fn) \
    __setup_param(str, fn, fn, 0)
```

# \_\_setup

```
__setup("console=", console_setup);
```

- Expands to

```
static const char __setup_str_console_setup[] __initconst \
__aligned(1) = "console=";\nstatic struct obs_kernel_param __setup_console_setup __used \
__section(.init.setup) __attribute__\
((aligned((sizeof(long)))))) \
= { __setup_str_console_setup, console_setup, early};
```

- Which expands to

```
static struct obs_kernel_param __setup_console_setup \
__section(.init.setup) = { __setup_str_console_setup,\nconsole_setup, early};
```

- This stores the code in a table in section `.init.setup`.

# On initialization...

- The table in **.init.setup** has
  - Parameter string (“**console=**”) and
  - Pointer to the function that processes it.
- This way the initialization code can process everything on the command line without knowing at compile time where all the code is.
- See section 5.3 of *Embedded Linux Primer* for more details.