

06-1 The Kernel

It all started with...

From: torvalds@klaava.Helsinki.FI (Linus Benedict Torvalds)
Newsgroups: comp.os.minix
Subject: What would you like to see most in minix?
Summary: small poll for my new operating system
Message-ID: <1991Aug25.205708.9541@klaava.Helsinki.FI>
Date: 25 Aug 91 20:57:08 GMT
Organization: University of Helsinki

Hello everybody out there using minix -

I'm doing a (free) operating system (just a hobby, won't be big and professional like gnu) for 386(486) AT clones. This has been brewing since april, and is starting to get ready. I'd like any feedback on things people like/dislike in minix, as my OS resembles it somewhat (same physical layout of the file-system (due to practical reasons) among other things).


I've currently ported bash(1.08) and gcc(1.40), and things seem to work. This implies that I'll get something practical within a few months, and I'd like to know what features most people would want. Any suggestions are welcome, but I won't promise I'll implement them :-)

Linus (torvalds@kruuna.helsinki.fi)

Free Electrons

Linux kernel introduction

Michael Opdenacker
Thomas Petazzoni
Free Electrons

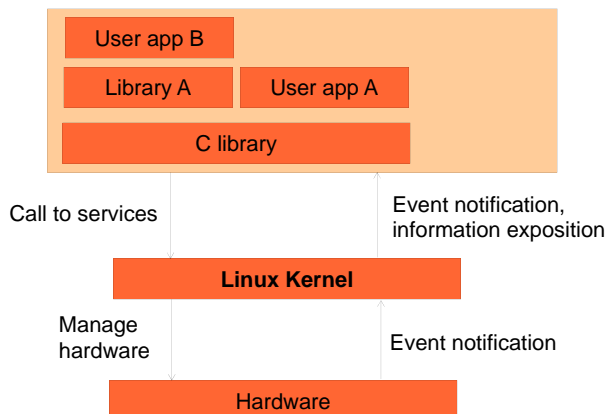


© Copyright 2004-2009, Free Electrons.
Creative Commons BY-SA 3.0 license
Latest update: 1/16/2012.
Document sources, updates and translations:
<http://free-electrons.com/docs/kernel-intro>
Corrections, suggestions, contributions and translations are welcome!

Embedded Linux driver development

Kernel overview
Linux features

Linux kernel in the system



History

- ▶ The Linux kernel is one component of a system, which also requires libraries and applications to provide features to end users
- ▶ The Linux kernel was created as a hobby in 1991 by a Finnish student, Linus Torvalds
- ▶ Linux quickly started to be used as the kernel for free software operating systems
- ▶ Linus Torvalds has been able to create a large and dynamic developer and user community around Linux
- ▶ Nowadays, hundreds of people contribute to each kernel release, individuals or companies big and small

Linux kernel key features

- ▶ Portability and hardware support. Runs on most architectures.
- ▶ Scalability
Can run on super computers as well as on tiny devices (4 MB of RAM is enough).
- ▶ Compliance to standards and interoperability.
- ▶ Exhaustive networking support.
- ▶ Security
It can't hide its flaws. Its code is reviewed by many experts.
- ▶ Stability and reliability.
- ▶ Modularity
Can include only what a system needs even at run time.
- ▶ Easy to program
You can learn from existing code. Many useful resources on the net.

Supported hardware architectures

2.6.31 status

What's the current version?

3.2.1

- ▶ See the [.../arch/](#) directory in the kernel sources
- ▶ Minimum: 32 bit processors, with or without MMU, and gcc support
- ▶ 32 bit architectures ([.../arch/](#) subdirectories)
[arm](#), [avr32](#), [blackfin](#), [cris](#), [frv](#), [h8300](#), [m32r](#), [m68k](#), [m68knommu](#), [microblaze](#), [mips](#), [mn10300](#), [parisc](#), [s390](#), [sparc](#), [um](#), [xtensa](#)
- ▶ 64 bit architectures:
[alpha](#), [ia64](#), [sparc64](#)
- ▶ 32/64 bit architectures
[powerpc](#), [x86](#), [sh](#)
- ▶ Find details in kernel sources: [arch/<arch>/Kconfig](#), [arch/<arch>/README](#), or [Documentation/<arch>/](#)

How did I find it?

kernel.org



System calls

- ▶ The main interface between the kernel and userspace is the set of system calls
- ▶ About ~300 system calls that provides the main kernel services
- ▶ This interface is stable over time: only new system calls can be added by the kernel developers
- ▶ This system call interface is wrapped by the C library, and userspace applications usually never make a system call directly but rather use the corresponding C library function

Virtual filesystems

- ▶ Linux makes system and kernel information available in user-space through virtual filesystems (virtual files not existing on any real storage). No need to know kernel programming to access such information!
- ▶ Mounting `/proc`:
`sudo mount -t proc none /proc`
- ▶ Mounting `/sys`:
`sudo mount -t sysfs none /sys`

Filesystem type Raw device or filesystem image
In the case of virtual filesystems, any string is fine Mount point

/proc details

A few examples:

- ▶ `/proc/cpuinfo`: processor information
- ▶ `/proc/meminfo`: memory status
- ▶ `/proc/version`: kernel version and build information
- ▶ `/proc/cmdline`: kernel command line
- ▶ `/proc/<pid>/environ`: calling environment
- ▶ `/proc/<pid>/cmdline`: process command line

... and many more! See by yourself!

Lots of details about the `/proc` interface are available in [Documentation/filesystems/proc.txt](#) (almost 2000 lines) in the kernel sources.

Embedded Linux usage

Kernel overview

Linux versioning scheme and development process

Changes since Linux 2.6 (1)


- ▶ Since **2.6.0**, kernel developers have been able to introduce lots of new features one by one on a steady pace, without having to make major changes in existing subsystems.
- ▶ Opening a new **Linux 2.7** (or **2.9**) development branch will be required only when **Linux 2.6** is no longer able to accommodate key features without undergoing traumatic changes.
- Thanks to this, more features are released to users at a faster pace.

Changes since Linux 2.6 (2)

Since 2.6.14, the kernel developers agreed on the following development model:

- ▶ After the release of a **2.6.x** version, a two-weeks merge window opens, during which major additions are merged.
- ▶ The merge window is closed by the release of test version **2.6.(x+1)-rc1**
- ▶ The bug fixing period opens, for 6 to 10 weeks.
- ▶ At regular intervals during the bug fixing period, **2.6.(x+1)-rcY** test versions are released.
- ▶ When considered sufficiently stable, kernel **2.6.(x+1)** is released, and the process starts again.

Merge and bug fixing windows



What's new in each Linux release?

commit 3c92c2ba33cd7d666c5f83cc32aa590e794e91b0
Author: Andi Kleen <ak@suse.de>
Date: Tue Oct 11 01:28:33 2005 +0200

[PATCH] i386: Don't discard upper 32bits of HWCR on K8

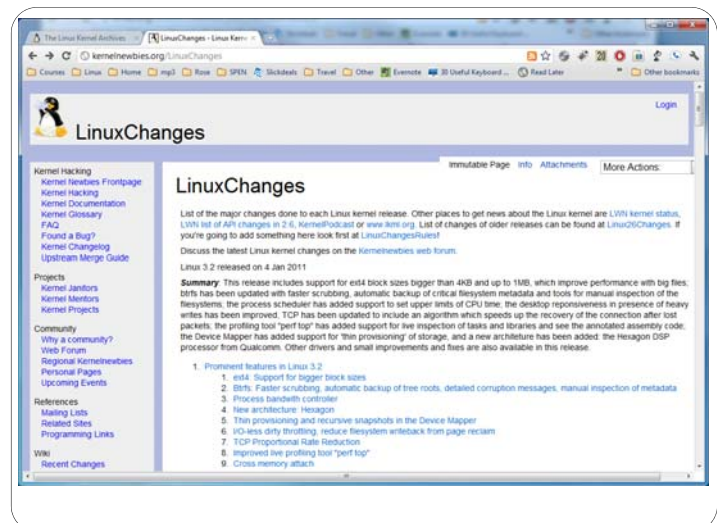
Need to use long long, not long when RMWing a MSR. I think it's harmless right now, but still should be better fixed if AMD adds any bits in the upper 32bit of HWCR.

Bug was introduced with the TLB flush filter fix for i386

Signed-off-by: Andi Kleen <ak@suse.de>
Signed-off-by: Linus Torvalds <torvalds@osdl.org>




- ▶ The official list of changes for each Linux release is just a huge list of individual patches!
- ▶ Very difficult to find out the key changes and to get the global picture out of individual changes.
- ▶ Fortunately, a summary of key changes with enough details is available on <http://wiki.kernelnewbies.org/LinuxChanges>



Embedded Linux kernel usage

Embedded Linux kernel usage

Michael Opdenacker
Thomas Petazzoni
Free Electrons



© Copyright 2004-2009, Free Electrons.
Creative Commons BY-SA 3.0 license
Latest update: 1/16/2012.
Document sources, updates and translations:
<http://free-electrons.com/docs/kernel-usage>
Corrections, suggestions, contributions and translations are welcome!

Embedded Linux usage

Compiling and booting Linux
Linux kernel sources

Location of kernel sources

- ▶ The official version of the Linux kernel, as released by Linus Torvalds is available at <http://www.kernel.org>
 - ▶ This version follows the well-defined development model of the kernel
- ▶ However, it may not contain the latest development from a specific area, due to the organization of the development model and because features in development might not be ready for mainline inclusion
- ▶ Many kernel sub-communities maintain their own kernel, with usually newer but less stable features
 - ▶ Architecture communities (ARM, MIPS, PowerPC, etc.), device drivers communities (I2C, SPI, USB, PCI, network, etc.), other communities (real-time, etc.)
 - ▶ They generally don't release official versions, only development trees are available

Like omap

Linux kernel size (1)


- ▶ Linux 2.6.31 sources:
Raw size: 350 MB (30,900 files, approx 12,000,000 lines)
[gzip](#) compressed tar archive: 75 MB
[bzip2](#) compressed tar archive: 59 MB (better)
[lzma](#) compressed tar archive: 49 MB (best)
- ▶ Minimum Linux 2.6.29 compiled kernel size with CONFIG_EMBEDDED, for a kernel that boots a QEMU PC (IDE hard drive, ext2 filesystem, ELF executable support):
532 KB (compressed), 1325 KB (raw)
- ▶ Why are these sources so big?
Because they include thousands of device drivers, many network protocols, support many architectures and filesystems...
- ▶ The Linux core (scheduler, memory management...) is pretty small!

Linux kernel size (1)

- ▶ Linux 2.6.31 sources:
Raw size: 350 MB (30,900 files, approx 12,000,000 lines)
[gzip](#) compressed tar archive: 75 MB
[bzip2](#) compressed tar archive: 59 MB (better)
[lzma](#) compressed tar archive: 49 MB (best)
- ▶ Linux 2.6.32 sources:
1.3G
- ▶ Linux 3.0.9 sources:
1.6G

Linux kernel size (2)

Size of Linux source directories (KB)



Getting Linux sources

- ▶ Full tarballs
 - ▶ Contain the complete kernel sources
 - ▶ Long to download and uncompress, but must be done at least once
 - ▶ Example:
<http://kernel.org/pub/linux/kernel/v2.6/linux-2.6.14.7.tar.bz2>

Getting Linux sources

- ▶ Incremental patches between versions
 - ▶ It assumes you already have a base version and you apply the correct patches in the right order
 - ▶ Quick to download and apply
 - ▶ Examples
<http://kernel.org/pub/linux/kernel/v2.6/patch-2.6.14.bz2> (2.6.13 to 2.6.14)
<http://kernel.org/pub/linux/kernel/v2.6/patch-2.6.14.7.bz2> (2.6.14 to 2.6.14.7)
- ▶ All previous kernel versions are available in
<http://kernel.org/pub/linux/kernel/>

Getting Linux sources

- ▶ git
 - ▶ `sudo apt-get install git-core6`
 - ▶ `git clone`
`git://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux-2.6.git` `linux-2.6`

Getting Linux sources

- ▶ bitbake
 - ▶ `cd ${OETREE}/build`
 - ▶ `bitbake -c clean linux-omap-2.6.28`
 - ▶ `bitbake -f -c compile linux-omap-2.6.28`

Top-Level Source Directory

<code>arch/</code>	<code>firmware/</code>	<code>kernel/</code>	<code>samples/</code>
<code>block/</code>	<code>fs/</code>	<code>lib/</code>	<code>scripts/</code>
<code>crypto/</code>	<code>include/</code>	<code>mm/</code>	<code>security/</code>
<code>Documentation/</code>	<code>init/</code>	<code>net/</code>	<code>sound/</code>
<code>drivers/</code>	<code>ipc/</code>	<code>patches/</code>	<code>usr/</code>
			<code>virt/</code>

Embedded Linux usage

Compiling and booting Linux
Kernel configuration

Kernel configuration

Defines what features to include in the kernel:

- ▶ Stored in the `.config` file at the root of kernel sources.
 - ▶ Simple text file
- ▶ Most useful commands to create this config file:
`make [xconfig|gconfig|menuconfig|oldconfig]`
- ▶ To modify a kernel in a GNU/Linux distribution:
the configuration files are usually released in `.../boot/`,
together with kernel images: `.../boot/config-2.6.17-11-generic`

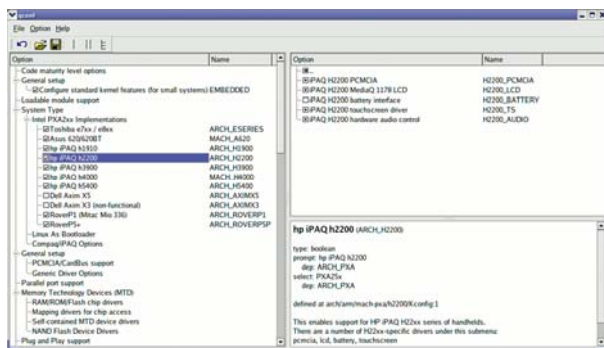
make xconfig

`make xconfig`

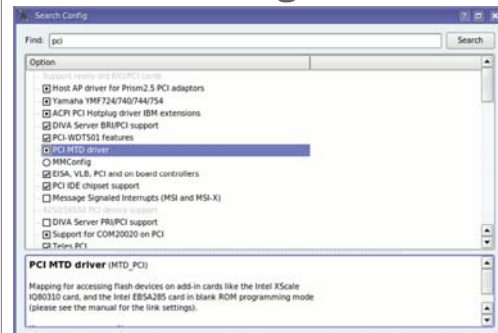
- ▶ The most common way to configure the kernel
- ▶ Make sure you read the help -> introduction: useful options!
- ▶ File browser: easier to load configuration files
- ▶ New search interface to look for parameters
- ▶ Required Debian / Ubuntu packages:
`libqt3-mt-dev, g++`



make xconfig screenshot



make xconfig search interface



Looks for a keyword in the description string

Allows to select or unselect found parameters.

Kernel configuration options

Compiled as a module (separate file)
`CONFIG_ISO9660_FS=m`

Driver options
`CONFIG_JOLIET=y` → ☒ Microsoft Joliet CDROM extensions
`CONFIG_ZISOFS=y` → ☒ Transparent decompression extension
`CONFIG_UDF_FS=y` → ☒ UDF file system support

Compiled statically into the kernel
`CONFIG_UDF_FS=y`

Corresponding .config file excerpt

```
#
# CD-ROM/DVD Filesystems
#
CONFIG_ISO9660_FS=m
CONFIG_JOLIET=y
CONFIG_ZISOFS=y
CONFIG_UDF_FS=y
CONFIG_UDF_NLS=y

#
# DOS/FAT/NT Filesystems
#
# CONFIG_MSDOS_FS is not set
# CONFIG_VFAT_FS is not set
CONFIG_NTFS_FS=m
# CONFIG_NTFS_DEBUG is not set
CONFIG_NTFS_RW=y
```

Section name
(helps to locate settings in the interface)

All parameters are prefixed with `CONFIG_`

Kernel option dependencies

- ▶ There are dependencies between kernel options
- ▶ For example, enabling a network driver requires the network stack to be enabled
- ▶ Two types of dependencies
 - ▶ *depends on* dependencies. In this case, option A that depends on option B is not visible until option B is enabled
 - ▶ *select* dependencies. In this case, with option A depending on option B, when option A is enabled, option B is automatically enabled
 - ▶ *make xconfig* allows to see all options, even those that cannot be selected because of missing dependencies. In this case, they are displayed in gray


make gconfig

`make gconfig`

New **GTK** based graphical configuration interface. Functionality similar to that of `make xconfig`.

Just lacking a search functionality.

Required Debian packages:
`libglade2-dev`



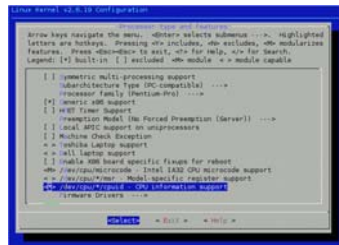
make menuconfig

`make menuconfig`

Useful when no graphics are available. Pretty convenient too!

Same interface found in other tools: `BusyBox`, `buildroot`...

Required Debian packages: `libncurses-dev`



make oldconfig

`make oldconfig`

- ▶ Needed very often!
- ▶ Useful to upgrade a `.config` file from an earlier kernel release
- ▶ Issues warnings for configuration parameters that no longer exist in the new kernel.
- ▶ Asks for values for new parameters

If you edit a `.config` file by hand, it's strongly recommended to run `make oldconfig` afterwards!

make allnoconfig

`make allnoconfig`

- ▶ Only sets strongly recommended settings to `y`.
- ▶ Sets all other settings to `n`.
- ▶ Very useful in embedded systems to select only the minimum required set of features and drivers.
- ▶ Much more convenient than unselecting hundreds of features one by one!

Undoing configuration changes

A frequent problem:

- ▶ After changing several kernel configuration settings, your kernel no longer works.
- ▶ If you don't remember all the changes you made, you can get back to your previous configuration:
`> cp .config.old .config`
- ▶ All the configuration interfaces of the kernel (`xconfig`, `menuconfig`, `allnoconfig`...) keep this `.config.old` backup copy.



make help

`make help`

- ▶ Lists all available `make` targets
- ▶ Useful to get a reminder, or to look for new or advanced options!

Embedded Linux usage

Compiling and installing the kernel
for the host system

Compiling and installing the kernel

Compiling step

▶ `make`

You can speed up compiling by running multiple compile jobs in parallel, especially if you have multiple CPU cores.

Example: `make -j 4`

MAY1

Slide 6/9

MAY1 How do you build for the target?
Mark A. Yoder, 12/22/2009

Kernel cleanup targets



- ▶ Clean-up generated files
(to force re-compiling drivers):
`make clean`
- ▶ Remove **all** generated files. Needed when switching from one architecture to another
Caution: also removes your `.config` file!
`make mrproper`
- ▶ Also remove editor backup and patch reject files:
(mainly to generate patches):
`make distclean`

Generated files

Created when you run the `make` command. The kernel is in fact a single binary image, nothing more !

- ▶ `.../vmlinux`
Raw Linux kernel image, non compressed.
- ▶ `.../arch/<arch>/boot/zImage` (default image on `arm`)
`zlib` compressed kernel image
- ▶ `.../arch/<arch>/boot/bzImage` (default image on `x86`)
Also a `zlib` compressed kernel image.
Caution: `bz` means “big zipped” but not “`bzip2` compressed”!

News: new compression formats are now available since 2.6.30:
lzma and bzip2. Free Electrons also contributed lzo support (very fast decompression).

Files created by make install

- ▶ `/boot/vmlinuz-<version>`
Compressed kernel image. Same as the one in `/arch/<arch>/boot`
- ▶ `/boot/System.map-<version>`
Stores kernel symbol addresses
- ▶ `/boot/config-<version>`
Kernel configuration for this version

Don't Use

Files created by make modules_install

`/lib/modules/<version>/:` Kernel modules + extras

- ▶ `kernel/`
Module `.ko` (Kernel Object) files, in the same directory structure as in the sources.
- ▶ `modules.alias`
Module aliases for module loading utilities. Example line:
`alias sound-service-?-0 snd_mixer_oss`
- ▶ `modules.dep`
Module dependencies
- ▶ `modules.symbols`
Tells which module a given symbol belongs to.

Don't Use

All the files in this directory are text files.

Don't hesitate to have a look by yourself!

The Details

To understand a system one must first understand it parts.


--Chris Hallinan

Link Stage: vmlinux

```
$ arm-linux-ld -EB -p --no-undefined -X -o vmlinux
-T arch/arm/kernel/vmlinux.lds \
arch/arm/kernel/head.o \
arch/arm/kernel/init_task.o \
init/built-in.o \
--start-group \
usr/built-in.o \
arch/arm/kernel/built-in.o \
arch/arm/mm/built-in.o \
arch/arm/common/built-in.o \
arch/arm/mach-ixp4xx/built-in.o \
arch/arm/nwfpe/built-in.o \
kernel/built-in.o \
mm/built-in.o \
fs/built-in.o \
ipc/built-in.o \
security/built-in.o \
crypto/built-in.o \
block/built-in.o \
arch/arm/lib/lib.a \
lib/lib.a \
arch/arm/lib/built-in.o \
lib/built-in.o \
drivers/built-in.o \
sound/built-in.o \
firmware/built-in.o \
net/built-in.o \
-end-group \
.tmp_kallsyms2.o \
```

Look in `/usr/local/xtools/arm-unknown-linux-uclibcgnueabi/bin`

vmlinux image components



Compare the two

```
$ arm-linux-ld -EB -p --no-undefined -X -o vmlinux
-T arch/arm/kernel/vmlinux.lds \
arch/arm/kernel/head.o \
arch/arm/kernel/init_task.o \
init/built-in.o \
--start-group \
usr/built-in.o \
arch/arm/kernel/built-in.o \
arch/arm/mm/built-in.o \
arch/arm/common/built-in.o \
arch/arm/mach-ixp4xx/built-in.o \
arch/arm/nwfpe/built-in.o \
kernel/built-in.o \
mm/built-in.o \
fs/built-in.o \
arch/arm/kernel/head.o \
arch/arm/kernel/init_task.o \
init \
usr/built-in.o \
arch/arm/kernel \
arch/arm/mm \
arch/arm/common \
arch/arm/mach-ixp4xx \
arch/arm/nwfpe \
kernel \
mm \
fs \
```

vmlinux Image Components Description

Table 4-1
vmlinux Image Components Description

Component	Description
arch/arm/kernel/head.o	Kernel architecture-specific startup code.
arch/arm/kernel/init_task.o	Initial thread and task structs required by kernel.
init/built-in.o	Main kernel initialization code. See Chapter 5.
usr/built-in.o	Built-in initramfs image. See Chapter 6.
arch/arm/kernel/built-in.o	Architecture-specific kernel code.
arch/arm/mm/built-in.o	Architecture-specific memory-management code.
arch/arm/common/built-in.o	Architecture-specific generic code. Varies by architecture.
arch/arm/mach-lpx4xx/built-in.o	Machine-specific code, usually initialization.
arch/arm/nvfp/built-in.o	Architecture-specific floating point-emulation code.
kernel/built-in.o	Common components of the kernel itself.
mm/built-in.o	Common components of memory manage-