

07-2 Device Driver Basics

Using kernel modules

Free Electrons

Loadable kernel modules

- ▶ Modules: add a given functionality to the kernel (drivers, filesystem support, and many others)
- ▶ Can be loaded and unloaded at any time, only when their functionality is need
- ▶ Useful to keep the kernel image size to the minimum (essential in GNU/Linux distributions for PCs)
- ▶ Also useful to reduce boot time: you don't spent time initializing devices and kernel features that you only need later
- ▶ Caution: once loaded, have full access to the whole kernel address space. No particular protection

Minimal Device Driver (Listing 8-1)

```
/* Example Minimal Character Device Driver */
#include <linux/module.h>
static int __init hello_init(void) {
    printk(KERN_INFO "Hello Example Init\n");
    return 0;
}
static void __exit hello_exit(void) {
    printk("Hello Example Exit\n");
}
module_init(hello_init);
module_exit(hello_exit);

MODULE_AUTHOR("Chris Hallinan");
MODULE_DESCRIPTION("Hello World Example");
MODULE_LICENSE("GPL");
```

Module Build Infrastructure

1. Starting from the top-level Linux source directory, create a directory under `.../drivers/char` called **examples**.
2. Add a menu item to the kernel configuration to enable building **examples** and to specify built-in or loadable kernel module. See Section 8.1.4, page 205
3. Add the new examples subdirectory to the `.../drivers/char/Makefile` conditional on the menu item created in step 2.
4. Create a **Makefile** for the new **examples** directory, and add the **hello.o** module object to be compiled conditional on the menu item created in step 2.
5. Finally, create the driver **hello1.c** source file from Listing 8-1.

Typo page 206

```
diff --git a/drivers/char/Kconfig b/drivers/char/Kconfig
index 6f31c94..0805290 100644
--- a/drivers/char/Kconfig
+++ b/drivers/char/Kconfig
@@ -4,6 +4,13 @@

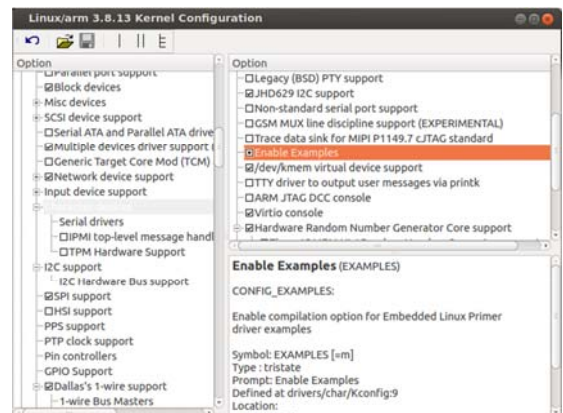
menu "Character devices"

+config EXAMPLES
+    tristate "Enable Examples"
+    default m
+    ---help---
+    Enable compilation option for Embedded Linux Primer
+    driver examples
+
config DEVKMEM
    bool "/dev/kmem virtual device support"
    default y
    help
```

Must be lower case

In new kernel

Check Config



Module Build Output

```
host$ time make modules
CHK    include/generated/uapi/linux/version.h
CHK    include/generated/utsrelease.h
make[1]: `include/generated/mach-types.h' is up to date.
CALL   scripts/checksyscalls.sh
CC [M]  drivers/char/examples/hello1.o
Building modules, stage 2.
MODPOST 1326 modules
LD [M]  drivers/char/examples/hello1.ko

real    0m33.706s
user    0m31.462s
sys     0m4.824s
```

First time takes much longer
Try make -jX modules

Once built... Option 1

On the Beagle... Two choices...

- Option 1:
`make INSTALL_MOD_PATH=~/.BeagleBoard modules_install`
- Will create **lib** directory in `~/BeagleBoard` with everything that goes in `/lib` on the Beagle

```
host$ ls -F ~/.BeagleBoard/lib/modules/3.8.13+/
```

```
build#          modules.dep.bin      modules.seriomap
kernel/         modules.devname    modules.softdep
modules.alias   modules.ieee1394map modules.symbols
modules.alias.bin modules.inputmap  modules.symbols.bin
modules.builtin modules.isapnpmap modules.usbmap
modules.builtin.bin modules.ofmap      source#
modules.cwmmap  modules.order
modules.dep      modules.pcimap
```

- Then
`host$ rm build source`
`host$ scp -r ~/.BeagleBoard/lib root@beagle:/lib`
- Could take a while to transfer

Once built... Option 2

- Just copy the new file you created

```
host$ scp .../drivers/char/examples/hello1.ko root@beagle:.
```

- On the Beagle

```
beagle$ cd /lib/modules/3.8.13/kernel/drivers/char/
beagle$ mkdir examples
beagle$ cd examples
beagle$ mv ~/hello1.ko .
```

Loading and Unloading a Module

```
beagle$ cd /lib/modules/3.8.13/kernel/drivers/char/examples
beagle$ insmod hello1.ko
beagle$ dmesg | tail -4
[  9.106206] snd-usb-audio 1-1:1.0: usb_probe_interface
[  9.106244] snd-usb-audio 1-1:1.0: usb_probe_interface - got id
[  9.813239] usbcore: registered new interface driver snd-usb-audio
[ 109.308551] Hello Example Init
beagle$ rmmod hello1
beagle$ dmesg | tail -4
[  9.106244] snd-usb-audio 1-1:1.0: usb_probe_interface - got id
[  9.813239] usbcore: registered new interface driver snd-usb-audio
[ 109.308551] Hello Example Init
[ 241.037368] Hello Example Exit
```

Module Utilities

```
$ insmod /lib/modules/`uname -r`/kernel/drivers/char/examples/hello1.ko
```

- No need to edit `modules.dep`

Example Driver with Parameter

```
/* Example Minimal Character Device Driver */
#include <linux/module.h>

static int debug_enable = 0;
module_param(debug_enable, int, 0);

MODULE_PARM_DESC(debug_enable, "Enable module debug mode.");

static int __init hello_init(void) {
    /* Now print value of new module parameter */
    printk("Hello Example Init - debug mode is %s\n",
           debug_enable ? "enabled" : "disabled");
    return 0;
}

static void __exit hello_exit(void) {
    printk("Hello Example Exit\n");
}

module_init(hello_init);
module_exit(hello_exit);

MODULE_AUTHOR("Chris Hallinan");
MODULE_DESCRIPTION("Hello World Example");
MODULE_LICENSE("GPL");
```

Passing Parameters to a Module

```
insmod /lib/modules/.../examples/hello1.ko
debug_enable=1
```

Hello Example Init - debug mode is *enabled*

```
insmod /lib/modules/.../examples/hello1.ko
```

Hello Example Init - debug mode is *disabled*

Other module commands

```
# /sbin/lsmmod
# /sbin/modinfo hello1
# /sbin/rmmod hello1
# /sbin/depmod (creates modules.dep.bin)
```

- Go play with them

Adding File System Ops to Hello.c

- Section 8.3, page 217 has a long example about adding file system operations to **hello.c**
- Look it over
- Creates a new device (**/dev/hello1**)
- You can read and write it
- Do it.

Driver File System Operations

- Once a device driver is loaded into the live kernel...
 - **open()** is used to prepare it for subsequent operations
 - **release()** is used to clean up
 - **ioctl()** is used for nonstandard communication
- Think in terms of reading and writing a file...

```
fd = open("file", ...
read(fd, ...
close(fd)
```

open/release additions to hello.c

```
#include <linux/fs.h>

#define HELLO_MAJOR 234
...
struct file_operations hello_fops;

static int hello_open(struct inode *inode, struct file *file) {
    printk("hello_open: successful\n");
    return 0;
}

static int hello_release(struct inode *inode, struct file *file) {
    printk("hello_release: successful\n");
    return 0;
}
```

read/write additions to hello.c

```
static ssize_t hello_read(struct file *file,
    char *buf, size_t count, loff_t *ptr) {
    printk("hello_read: returning zero bytes\n");
    return 0;
}

static ssize_t hello_write(struct file *file,
    const char *buf, size_t count, loff_t *ppos)
{
    printk("hello_read: accepting zero bytes\n");
    return 0;
}
```

ioctl additions to hello.c

```
static int hello_ioctl(struct inode *inode,
    struct file *file, unsigned int cmd,
    unsigned long arg) {
    printk("hello_ioctl: cmd=%ld, arg=%ld\n",
        cmd, arg);
    return 0;
}
```

init additions to hello.c

```
#define HELLO_MAJOR 234
static int __init hello_init(void)
{
    int ret;
    printk("Hello Example Init - debug mode is %s\n",
        debug_enable ? "enabled" : "disabled");
    ret = register_chrdev(HELLO_MAJOR, "hello1", &hello_fops);
    if (ret < 0) {
        printk("Error registering hello device\n");
        goto hello_fail1;
    }
    printk("Hello: registered module successfully!\n");

    /* Init processing here... */

    return 0;

hello_fail1:
    return ret;
}
```

Major number for device driver

- Every device has a major and minor number

```
$ ls -ls /dev/console
0 crw----- 1 yoder root 5, 1 2011-02-06 17:57 /dev/console
```

- Device numbers *used* to be statically assigned
- See .../Documentation/devices.txt

```
5 char  Alternate TTY devices
        0 = /dev/tty           Current TTY device
        1 = /dev/console System console
        2 = /dev/ptmx          PTY master multiplex
        64 = /dev/cua0         Callout device for ttyS0
```

- The text uses static assignment

```
234-239          UNASSIGNED
240-254 char     LOCAL/EXPERIMENTAL USE
```

Registering our functions

- Struct `file_operations` is used to bind our functions to the requests from the file system.

```
struct file_operations hello_fops = {
    owner:    THIS_MODULE,
    read:     hello_read,
    write:    hello_write,
    ioctl:    hello_ioctl,
    open:     hello_open,
    release:  hello_release,
};
```

init additions to hello.c

```
#define HELLO_MAJOR 234
static int __init hello_init(void)
{
    int ret;
    printk("Hello Example Init - debug mode is %s\n",
        debug_enable ? "enabled" : "disabled");
    ret = register_chrdev(HELLO_MAJOR, "hello1", &hello_fops);
    if (ret < 0) {
        printk("Error registering hello device\n");
        goto hello_fail1;
    }
    printk("Hello: registered module successfully!\n");
    /* Init processing here... */
    return 0;
hello_fail1:
    return ret;
}
```

Device Nodes and `mknod`

- Use `mknod` to create a new device

```
$ mknod /dev/hello1 c 234 0
```



- Then

```
$ ls -l /dev/hello1
crw-r--r-- 1 root root 234, 0 Apr 2 2011 /dev/hello1
```

Dynamic Major Number

- The above example uses the older *static* method to assign a device number
- Today dynamic allocation is preferred
- Here is how:

```
#include <linux/kdev_t.h>
dev_t dev;
```

- This declares **dev** to be a device number (both major and minor). Now assign it a value
- **dev = MKDEV(234, 0);**

Requesting a number

- Now request a number

```
#include <linux/fs.h>
```

```
int register_chrdev_region(dev, 4, "hello");
```

- This requests a device number starting with 234 (previous page)
- It asks for 4 minor numbers
- Uses the name "hello"
- When done with the device use:

```
void unregister_chrdev_region(dev, 4);
```

Using **mknod**

- If your major number is assigned dynamically, how do you use **mknod**? Try the following

```
module="hello"
device="hello"
mode="664"
# remove stale nodes
/sbin/insmod ./module.ko $* || exit 1
rm -f /dev/${device}0
major=$(awk "\$2==\"$module\" {print \$1} /proc/devices"
mknod /dev/${device}0 c $major 0
```

/proc/devices

Character devices:	89 i2c	Block devices:	70 sd
1 mem	90 mtd	1 ramdisk	71 sd
4 /dev/vc/0	116 alsa	259 blkext	128 sd
4 tty	128 ptm	7 loop	129 sd
4 ttys	136 pts	8 sd	130 sd
5 /dev/tty	153 spi	11 sr	131 sd
5 /dev/console	161 ircomm	31 mtblock	132 sd
5 /dev/ptmx	180 usb	65 sd	133 sd
7 vcs	189 usb_device	66 sd	134 sd
10 misc	216 rfcomm	67 sd	135 sd
13 input	247 bccat	68 sd	179 mmc
14 sound	248 pvrsrvkm	69 sd	
21 sg	249 rtc		
29 fb	250 ttySDIO		
81 video4linux	251 omap-resizer		
	252 omap-previewer		
	253 ushmon		
	254 bsg		

Assignment

- See http://elinux.org/EBC_Exercise_26_Device_Drivers

Module dependencies

- ▶ Some kernel modules can depend on other modules, which need to be loaded first
- ▶ Example: the **usb-storage** module depends on the **scsi_mod**, **libusual** and **usbcore** modules
- ▶ Dependencies are described in **/lib/modules/<kernel-version>/modules.dep**

/lib/modules/2.6.32/models.dep

```
kernel/drivers/char/examples/hello1.ko:
kernel/crypto/twofish_common.ko:
kernel/crypto/ctr.ko:
kernel/crypto/blowfish.ko:
kernel/crypto/ghash-generic.ko:
kernel/crypto/gf128mul.ko
kernel/crypto/xts.ko:
kernel/crypto/gf128mul.ko
kernel/crypto/gcm.ko:
kernel/crypto/cryptd.ko:
kernel/crypto/md4.ko:
kernel/crypto/lrw.ko:
kernel/crypto/gf128mul.ko
```

Kernel log

When a new module is loaded, related information is available in the kernel log

- ▶ The kernel keeps its messages in a circular buffer (so that it doesn't consume more memory with many messages)
- ▶ Kernel log messages are available through the **dmesg** command ("diagnostic message")
- ▶ Kernel log messages are also displayed in the system console (messages can be filtered by level using `/proc/sys/kernel/printk`)

printk

- `/proc/sys/kernel/printk`
- The four values in this file are
 - `console_loglevel`,
 - `default_message_loglevel`,
 - `minimum_console_level` and
 - `default_console_loglevel`.
- These values influence **printk()** behavior when printing or logging error messages
- Messages with a higher priority than `console_loglevel` will be printed to the console
- Messages without an explicit priority will be printed with priority `default_message_level`

<http://www.tin.org/bin/man.cgi?section=5&topic=proc>

Kernel log levels

0 (KERN_EMERG)	The system is unusable
1 (KERN_ALERT)	Actions that must be taken care of immediately
2 (KERN_CRIT)	Critical conditions
3 (KERN_ERR)	Noncritical error conditions
4 (KERN_WARNING)	Warning conditions that should be taken care of
5 (KERN_NOTICE)	Normal, but significant events
6 (KERN_INFO)	Informational messages that require no action
7 (KERN_DEBUG)	Kernel debugging messages, output by the

Module utilities (1)

- ▶ **modinfo <module_name>**
modinfo <module_path>.ko
Gets information about a module: parameters, license, description and dependencies.
Very useful before deciding to load a module or not.
- ▶ **sudo insmod <module_path>.ko**
Tries to load the given module. The full path to the module object file must be given.

Understanding module loading

- ▶ When loading a module fails, **insmod** often doesn't give you enough details!
- ▶ Details are often available in the kernel log
- ▶ Example:

```
beagle$ sudo insmod ./intr_monitor.ko
insmod: error inserting './intr_monitor.ko': -1
Device or resource busy
beagle$ dmesg
[17549774.552000] Failed to register handler for irq channel 2
```

Module utilities (2)

▶ `sudo modprobe <module_name>`

Most common usage of `modprobe`: tries to load all the modules the given module depends on, and then this module. Lots of other options are available. `modprobe` automatically looks in `/lib/modules/<version>/` for the object file corresponding to the given module name.

▶ `lsmod`

Displays the list of loaded modules
Compare its output with the contents of `/proc/modules`!

lsmod

```
beagle$ lsmod
Module                Size  Used by
bufferclass_ti        4768  0
omaplfb               8733  0
pvrsrvkm             154248  2 bufferclass_ti,omaplfb
rfcomm               33484  0
ircomm_tty           30305  0
ircomm               16429  1 ircomm_tty
irda                 162973  2 ircomm_tty,ircomm
ipv6                 249063  14
hidp                 11193  0
l2cap                 30104  4 rfcomm,hidp,l2cap
bluetooth            49221  3 rfcomm,hidp,l2cap
...
```

Module utilities (3)

▶ `sudo rmmod <module_name>`

Tries to remove the given module.
Will only be allowed if the module is no longer in use (for example, no more processes opening a device file)

▶ `sudo modprobe -r <module_name>`

Tries to remove the given module and all dependent modules (which are no longer needed after the module removal)

Passing parameters to modules

▶ Find available parameters: `modinfo snd-intel8x0m`

▶ Through `insmod`: `sudo insmod ./snd-intel8x0m.ko index=-2`

▶ Through `modprobe`: Set parameters in `/etc/modprobe.conf` or in any file in `/etc/modprobe.d/`: `options snd-intel8x0m index=-2`

▶ Through the kernel command line, when the module is built statically into the kernel: `snd-intel8x0m.index=-2`

module name ↑
module parameter name ↑
module parameter value —

Useful reading

Linux Kernel in a Nutshell, Dec 2006

▶ By Greg Kroah-Hartman, O'Reilly <http://www.kroah.com/lkn/>

▶ A good reference book and guide on configuring, compiling and managing the Linux kernel sources.

▶ Freely available on-line!

Great companion to the printed book for easy electronic searches!
Available as single PDF file on

<http://free-electrons.com/community/kernel/lkn/>

▶ In `exercises/pptx`



Useful reading too

Linux Device Drivers, Third Edition, February 2005

▶ By Jonathan Corbet, Alessandro Rubini, Greg Kroah-Hartman, O'Reilly <http://lwn.net/Kernel/LDD3/>

▶ Freely available on-line!

Great companion to the printed book for easy electronic searches!
Available as single PDF file

▶ LDD3 is current as of the 2.6.10 kernel (Old?)

▶ In `exercises/pptx`

