

07-2 Device Driver Basics

Using kernel modules

Free Electrons

Loadable kernel modules

- ▶ Modules: add a given functionality to the kernel (drivers, filesystem support, and many others)
- ▶ Can be loaded and unloaded at any time, only when their functionality is need
- ▶ Useful to keep the kernel image size to the minimum (essential in GNU/Linux distributions for PCs)
- ▶ Also useful to reduce boot time: you don't spent time initializing devices and kernel features that you only need later
- ▶ Caution: once loaded, have full access to the whole kernel address space. No particular protection

Minimal Device Driver (Listing 8-1)

```
/* Example Minimal Character Device Driver */
#include <linux/module.h>
static int __init hello_init(void) {
    printk(KERN_INFO "Hello Example Init\n");
    return 0;
}
static void __exit hello_exit(void) {
    printk("Hello Example Exit\n");
}
module_init(hello_init);
module_exit(hello_exit);

MODULE_AUTHOR("Chris Hallinan");
MODULE_DESCRIPTION("Hello World Example");
MODULE_LICENSE("GPL");
```

Module Build Infrastructure

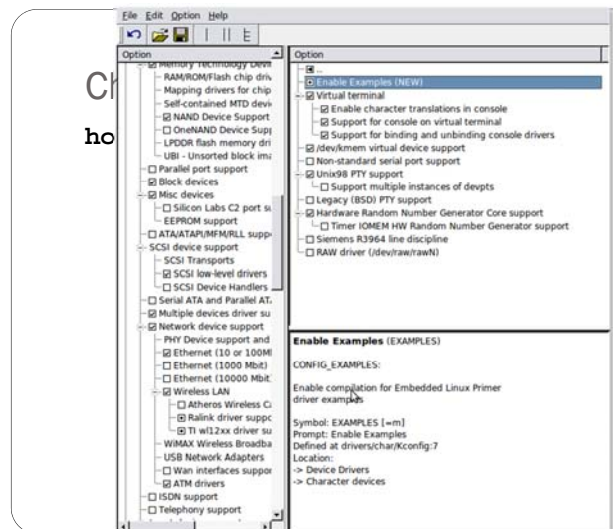
1. Starting from the top-level Linux source directory, create a directory under `.../drivers/char` called **examples**.
2. Add a menu item to the kernel configuration to enable building **examples** and to specify built-in or loadable kernel module. See Section 8.1.4, page 205
3. Add the new examples subdirectory to the `.../drivers/char/Makefile` conditional on the menu item created in step 2.
4. Create a **Makefile** for the new **examples** directory, and add the **hello1.o** module object to be compiled conditional on the menu item created in step 2.
5. Finally, create the driver **hello1.c** source file from Listing 8-1.

Typo page 206

```
diff --git a/drivers/char/Kconfig b/drivers/char/Kconfig
index 6f31c94..0805290 100644
--- a/drivers/char/Kconfig
+++ b/drivers/char/Kconfig
@@ -4,6 +4,13 @@
 menu "Character devices"

+config EXAMPLES
+    tristate "Enable Examples"
+    default m
+    ---help---
+    Enable compilation option for Embedded Linux Primer
+    driver examples
+
config VT
    bool "Virtual terminal" if EMBEDDED
    depends on !S390
```

Must be lower case



Module Build Output

```
host$ time make modules
CHK      include/linux/version.h
make[1]: `include/asm-arm/mach-types.h' is up to date.
CHK      include/linux/utsrelease.h
SYMLINK  include/asm -> include/asm-arm
CALL     scripts/checksyscalls.sh
<stdin>:1097:2: warning: #warning syscall fadvise64 not implemented
<stdin>:1265:2: warning: #warning syscall migrate_pages not implemented
CC [M]  drivers/char/examples/hello1.o
Building modules, stage 2.
MODPOST 691 modules
CC      drivers/char/examples/hello1.mod.o
LD [M]  drivers/char/examples/hello1.ko

real    0m41.559s
user    0m10.905s
sys     0m23.877s
```

Module First Build Output

```
host$ time make modules
HOSTLD  scripts/kconfig/conf
scripts/kconfig/conf -s arch/arm/Kconfig
*
* Restart config...
*
* Character devices
*
Enable Examples (EXAMPLES) [N/n/y/?] (NEW)
Virtual terminal (VT) [Y/n/?] y
  Enable character translations in console (CONSOLE_TRANSLATIONS) [Y/n/?] y
  Support for console on virtual terminal (VT_CONSOLE) [Y/n/?] y
  Support for binding and unbinding console drivers (VT_HW_CONSOLE_BINDING) [Y/n/?] y
/dev/kmem virtual device support (DEVKMEM) [Y/n/?] y
Non-standard serial port support (SERIAL_NONSTANDARD) [N/y/?] n
Unix98 PTY support (UNIX98_PTY) [Y/n/?] y
  Support multiple instances of devpts (DEVPTS_MULTIPLE_INSTANCES) [N/y/?] n
Legacy (BSD) PTY support (LEGACY_PTY) [N/y/?] n
Hardware Random Number Generator Core support (HW_RANDOM) [Y/n/m/?] y
  Timer IOMEM HW Random Number Generator support (HW_RANDOM_TIMERIOMEM) [N/m/y/?] n
Siemens R3964 line discipline (R3964) [N/m/y/?] n
RAM driver (/dev/raw/rawN) (RAM_DRIVER) [N/m/y/?] n
```

Module Build Output

```
#
# configuration written to .config
#
CHK      include/linux/version.h
make[1]: `include/asm-arm/mach-types.h' is up to date.
CHK      include/linux/utsrelease.h
SYMLINK  include/asm -> include/asm-arm
CALL     scripts/checksyscalls.sh
<stdin>:1523:2: warning: #warning syscall recvmsg not implemented
HOSTCC  scripts/genksyms/lex.o
HOSTLD  scripts/genksyms/genksyms
HOSTCC  scripts/mod/sumversion.o
HOSTLD  scripts/mod/modpost
HOSTCC  scripts/pnmtologo
Building modules, stage 2.
MODPOST 700 modules

real    1m25.549s
user    0m47.680s
sys     0m15.810s
```

Once built... Option 1

On the Beagle... Two choices....

- Option 1:
make INSTALL_MOD_PATH=~/BeagleBoard modules_install
- Will create **lib** directory in **~/BeagleBoard** with everything that goes in **/lib** on the Beagle
- Then
host\$ scp -r ~/BeagleBoard/lib root@beagle:/lib
- Could take a while to transfer

Once built... Option 2

- Just copy the new file you created
host\$ scp ../drivers/char/examples/hello1.ko root@beagle:..
- On the Beagle
beagle\$ mv hello1.ko
lib/modules/2.6.32/kernel/drivers/char/examples/
beagle\$ cd /lib/modules/2.6.32
beagle\$ mv modules.dep.bin modules.dep.bin.orig
beagle\$ edit modules.dep
- add:
kernel/drivers/char/examples/hello1.ko:

Loading and Unloading a Module

```
beagle$ /sbin/modprobe hello1
beagle$ dmesg | tail -4
[ 47.095764] OMAPFB: ioctl QUERY_PLANE
[ 47.095794] OMAPFB: ioctl GET_CAPS
[ 49.005889] eth0: no IPv6 routers present
[ 651.947784] Hello Example Init
beagle$ /sbin/modprobe -r hello1
beagle$ dmesg | tail -4
[ 47.095794] OMAPFB: ioctl GET_CAPS
[ 49.005889] eth0: no IPv6 routers present
[ 651.947784] Hello Example Init
[ 677.682769] Hello Example Exit
```

Module Utilities

```
$ insmod /lib/modules/`uname -r`/kernel/drivers/char/examples/hello1.ko
```

- No need to edit modules.dep

Example Driver with Parameter

```
/* Example Minimal Character Device Driver */
#include <linux/module.h>

static int debug_enable = 0;
module_param(debug_enable, int, 0);

MODULE_PARM_DESC(debug_enable, "Enable module debug mode.");

static int __init hello_init(void) {
    /* Now print value of new module parameter */
    printk("Hello Example Init - debug mode is %s\n",
           debug_enable ? "enabled" : "disabled");
    return 0;
}

static void __exit hello_exit(void) {
    printk("Hello Example Exit\n");
}

module_init(hello_init);
module_exit(hello_exit);

MODULE_AUTHOR("Chris Hallinan");
MODULE_DESCRIPTION("Hello World Example");
MODULE_LICENSE("GPL");
```

Passing Parameters to a Module

```
insmod /lib/modules/.../examples/hello1.ko
debug_enable=1
```

Hello Example Init - debug mode is *enabled*

```
insmod /lib/modules/.../examples/hello1.ko
```

Hello Example Init - debug mode is *disabled*

Other module commands

```
# /sbin/lsmmod
# /sbin/modinfo hello1
# /sbin/rmmod hello1
# /sbin/depmod (creates modules.dep.bin)
```

- Go play with them

Adding File System Ops to Hello.c

- Section 8.3, page 217 has a long example about adding file system operations to **hello.c**
- Look it over
- Creates a new device (**/dev/hello1**)
- You can read and write it
- Do it.

Driver File System Operations

- Once a device driver is loaded into the live kernel...
 - **open()** is used to prepare it for subsequent operations
 - **release()** is used to clean up
 - **ioctl()** is used for nonstandard communication
- Think in terms of reading and writing a file...

```
fd = open("file", ...
read(fd, ...
close(fd)
```

open/release additions to hello.c

```
#include <linux/fs.h>

#define HELLO_MAJOR 234
...
struct file_operations hello_fops;

static int hello_open(struct inode *inode, struct file *file) {
    printk("hello_open: successful\n");
    return 0;
}

static int hello_release(struct inode *inode, struct file *file) {
    printk("hello_release: successful\n");
    return 0;
}
```

read/write additions to hello.c

```
static ssize_t hello_read(struct file *file,
    char *buf, size_t count, loff_t *ptr) {
    printk("hello_read: returning zero bytes\n");
    return 0;
}

static ssize_t hello_write(struct file *file,
    const char *buf, size_t count, loff_t *ppos)
{
    printk("hello_read: accepting zero bytes\n");
    return 0;
}
```

ioctl additions to hello.c

```
static int hello_ioctl(struct inode *inode,
    struct file *file, unsigned int cmd,
    unsigned long arg) {
    printk("hello_ioctl: cmd=%ld, arg=%ld\n",
        cmd, arg);
    return 0;
}
```

init additions to hello.c

```
#define HELLO_MAJOR 234
static int __init hello_init(void)
{
    int ret;
    printk("Hello Example Init - debug mode is %s\n",
        debug_enable ? "enabled" : "disabled");
    ret = register_chrdev(HELLO_MAJOR, "hello", &hello_fops);
    if (ret < 0) {
        printk("Error registering hello device\n");
        goto hello_fail1;
    }
    printk("Hello: registered module successfully!\n");

    /* Init processing here... */

    return 0;

hello_fail1:
    return ret;
}
```

Major number for device driver

- Every device has a major and minor number
- ```
$ ls -ls /dev/console
0 crw----- 1 yoder root 5, 1 2011-02-06 17:57 /dev/console
```
- Device numbers used to be statically assigned
  - See .../Documentation/devices.txt
- ```
5 char  Alternate TTY devices
        0 = /dev/tty          Current TTY device
        1 = /dev/console System console
        2 = /dev/ptmx         PTY master multiplex
        64 = /dev/cua0        Callout device for ttyS0
```
- The text uses static assignment
- ```
234-239 UNASSIGNED
240-254 char LOCAL/EXPERIMENTAL USE
```

## Registering our functions

- Struct file\_operations is used bind our functions to the requests from the file system.
- ```
struct file_operations hello_fops = {
    owner:    THIS_MODULE,
    read:     hello_read,
    write:    hello_write,
    ioctl:    hello_ioctl,
    open:     hello_open,
    release:  hello_release,
};
```

init additions to hello.c

```
#define HELLO_MAJOR 234
static int __init hello_init(void)
{
    int ret;
    printk("Hello Example Init - debug mode is %s\n",
           debug_enable ? "enabled" : "disabled");
    ret = register_chrdev(HELLO_MAJOR, "hello1", &hello_fops);
    if (ret < 0) {
        printk("Error registering hello device\n");
        goto hello_fail;
    }
    printk("Hello: registered module successfully!\n");
    /* Init processing here... */
    return 0;
hello_fail:
    return ret;
}
```

Device Nodes and **mknod**

- Use **mknod** to create a new device

```
$ mknod /dev/hello1 c 234 0
```



- Then

```
$ ls -l /dev/hello1
```

```
crw-r--r--  1 root  root  234, 0 Apr 2 2011  /dev/hello1
```

Dynamic Major Number

- The above example uses the older *static* method to assign a device number
- Today dynamic allocation is preferred
- Here is how:

```
#include <linux/kdev_t.h>
```

```
dev_t dev;
```

- This declares **dev** to be a device number (both major and minor). Now assign it a value
- **dev = MKDEV(234, 0);**

Requesting a number

- Now request a number

```
#include <linux/fs.h>
```

```
int register_chrdev_region(dev, 4, "hello");
```

- This requests a device number starting with 234 (previous page)
- It asks for 4 minor numbers
- Uses the name "hello"

- When done with the device use:

```
void unregister_chrdev_region(dev, 4);
```

Using **mknod**

- If your major number is assigned dynamically, how do you use **mknod**? Try the following

```
module="hello"
```

```
device="hello"
```

```
mode="664"
```

```
/sbin/insmod ./module.ko $* || exit 1
```

```
# remove stale nodes
```

```
rm -f /dev/${device}0
```

```
major=`awk "\$2==\"$module\" {print \$1} /proc/devices`
```

```
mknod /dev/${device}0 c $major 0
```

Assignment

- See http://elinux.org/ECE497_Lab08_Device_Drivers

Module dependencies

- ▶ Some kernel modules can depend on other modules, which need to be loaded first.
- ▶ Example: the `usb-storage` module depends on the `scsi_mod`, `libusual` and `usbcore` modules.
- ▶ Dependencies are described in `/lib/modules/<kernel-version>/modules.dep`

Kernel log

When a new module is loaded, related information is available in the kernel log.

- ▶ The kernel keeps its messages in a circular buffer (so that it doesn't consume more memory with many messages)
- ▶ Kernel log messages are available through the `dmesg` command. (“**diagnostic message**”)
- ▶ Kernel log messages are also displayed in the system console (messages can be filtered by level using `/proc/sys/kernel/printk`)

printk

- `/proc/sys/kernel/printk`
- The four values in this file are
 - `console_loglevel`,
 - `default_message_loglevel`,
 - `minimum_console_level` and
 - `default_console_loglevel`.
- These values influence `printk()` behavior when printing or logging error messages.
- Messages with a higher priority than `console_loglevel` will be printed to the console.
- Messages without an explicit priority will be printed with priority `default_message_level`.

<http://www.tin.org/bin/man.cgi?section=5&topic=proc>

Kernel log levels

0 (KERN_EMERG)	The system is unusable.
1 (KERN_ALERT)	Actions that must be taken care of immediately.
2 (KERN_CRIT)	Critical conditions.
3 (KERN_ERR)	Noncritical error conditions.
4 (KERN_WARNING)	Warning conditions that should be taken care of.
5 (KERN_NOTICE)	Normal, but significant events.
6 (KERN_INFO)	Informational messages that require no action.
7 (KERN_DEBUG)	Kernel debugging messages, output by the

Module utilities (1)

- ▶ `modinfo <module_name>`
`modinfo <module_path>.ko`
Gets information about a module: parameters, license, description and dependencies.
Very useful before deciding to load a module or not.
- ▶ `sudo insmod <module_path>.ko`
Tries to load the given module. The full path to the module object file must be given.

Understanding module loading

- ▶ When loading a module fails, `insmod` often doesn't give you enough details!
- ▶ Details are often available in the kernel log.
- ▶ Example:

```
> sudo insmod ./intr_monitor.ko
insmod: error inserting './intr_monitor.ko': -1
Device or resource busy
> dmesg
[17549774.552000] Failed to register handler for
irq channel 2
```

Module utilities (2)

▶ `sudo modprobe <module_name>`

Most common usage of `modprobe`: tries to load all the modules the given module depends on, and then this module. Lots of other options are available. `modprobe` automatically looks in `/lib/modules/<version>/` for the object file corresponding to the given module name.

▶ `lsmod`

Displays the list of loaded modules
Compare its output with the contents of
`/proc/modules!`

Module utilities (3)

▶ `sudo rmmod <module_name>`

Tries to remove the given module.
Will only be allowed if the module is no longer in use
(for example, no more processes opening a device file)

▶ `sudo modprobe -r <module_name>`

Tries to remove the given module and all dependent modules (which are no longer needed after the module removal)

Passing parameters to modules

▶ Find available parameters: `modinfo snd-intel8x0m`

▶ Through `insmod`: `sudo insmod ./snd-intel8x0m.ko index=-2`

▶ Through `modprobe`: Set parameters in `/etc/modprobe.conf` or in any file in `/etc/modprobe.d/`: `options snd-intel8x0m index=-2`

▶ Through the kernel command line, when the module is built statically into the kernel: `snd-intel8x0m.index=-2`

module name ↑
module parameter name ↑
module parameter value —

Useful reading

Linux Kernel in a Nutshell, Dec 2006

▶ By Greg Kroah-Hartman, O'Reilly <http://www.kroah.com/lkn/>

▶ A good reference book and guide on configuring, compiling and managing the Linux kernel sources.

▶ **Freely available on-line!**

Great companion to the printed book
for easy electronic searches!

Available as single PDF file on
<http://free-electrons.com/community/kernel/lkn/>



Useful reading too

Linux Device Drivers, Third Edition, February 2005

▶ By Jonathan Corbet, Alessandro Rubini, Greg Kroah-Hartman, O'Reilly <http://lwn.net/Kernel/LDD3/>

▶ **Freely available on-line!**

Great companion to the printed book
for easy electronic searches!
Available as single PDF file

▶ LDD3 is current as of the 2.6.10 kernel (Old?)

