

## 07-2 Device Driver Basics

Using kernel modules

Free Electrons

## Loadable kernel modules

- ▶ Modules: add a given functionality to the kernel (drivers, filesystem support, and many others)
- ▶ Can be loaded and unloaded at any time, only when their functionality is needed
- ▶ Useful to keep the kernel image size to the minimum (essential in GNU/Linux distributions for PCs)
- ▶ Also useful to reduce boot time: you don't spend time initializing devices and kernel features that you only need later
- ▶ Caution: once loaded, have full access to the whole kernel address space. No particular protection

## Minimal Device Driver (Listing 8-1)

```
/* Example Minimal Character Device Driver */
#include <linux/module.h>
static int __init hello_init(void) {
    printk(KERN_INFO "Hello Example Init\n");
    return 0;
}
static void __exit hello_exit(void) {
    printk("Hello Example Exit\n");
}
module_init(hello_init);
module_exit(hello_exit);

MODULE_AUTHOR("Chris Hallinan");
MODULE_DESCRIPTION("Hello World Example");
MODULE_LICENSE("GPL");
```

## Module Build Infrastructure

1. Starting from the top-level Linux source directory, create a directory under `.../drivers/char` called **examples**.
2. Add a menu item to the kernel configuration to enable building **examples** and to specify built-in or loadable kernel module.  
See Section 8.1.4, page 205
3. Add the new examples subdirectory to the `.../drivers/char/Makefile` conditional on the menu item created in step 2.
4. Create a **Makefile** for the new **examples** directory, and add the **hello.o** module object to be compiled conditional on the menu item created in step 2.
5. Finally, create the driver **hello1.c** source file from Listing 8-1.

## Typo page 206

```
diff --git a/drivers/char/Kconfig b/drivers/char/Kconfig
index 6f31c94..0805290 100644
--- a/drivers/char/Kconfig
+++ b/drivers/char/Kconfig
@@ -4,6 +4,13 @@

```

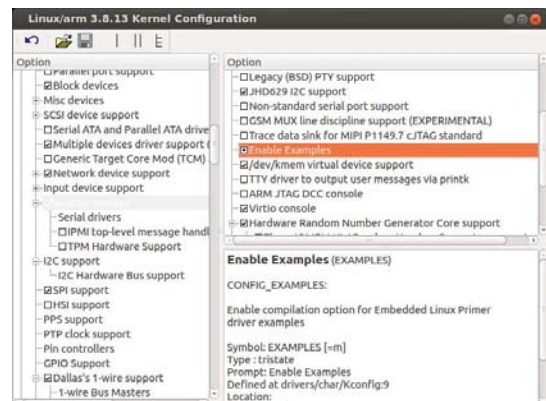
```
menu "Character devices"

+config EXAMPLES
+    tristate "Enable Examples"
+    default m
+    ---help---
+    Enable compilation option for Embedded Linux Primer
+    driver examples
+
config DEVMEM
    bool "/dev/kmem virtual device support"
    default y
    help
```

Must be lower case

In new kernel

## Check Config



## Module Build Output – In tree

```
host$ cd bb-kernel/KERNEL
host$ time make modules
CHK include/generated/uapi/linux/version.h
CHK include/generated/utsrelease.h
make[1]: `include/generated/mach-types.h' is up to date.
CALL scripts/checksyscalls.sh
CC [M] drivers/char/examples/hello1.o
Building modules, stage 2.
MODPOST 1326 modules
LD [M] drivers/char/examples/hello1.ko

real    0m33.706s
user    0m31.462s
sys     0m4.824s
```

First time takes much longer  
Try make -jX modules

## Module First Build Output

```
host$ time make modules
HOSTLD scripts/kconfig/conf
scripts/kconfig/conf -s arch/arm/Kconfig
*
* Restart config...
*
* Character devices
*
Enable Examples (EXAMPLES) [N/y/?] (NEW)
Virtual terminal (VT) [Y/n/?] y
  Enable character translations in console (CONSOLE_TRANSLATIONS) [Y/n/?] y
  Support for console on virtual terminal (VT_CONSOLE) [Y/n/?] y
  Support for binding and unbinding console drivers (VT_HW_CONSOLE_BINDING) [Y/n/?] n
/dev/kmem virtual device support (DEVKMEM) [Y/n/?] y
Non-standard serial port support (SERIAL_NONSTANDARD) [N/y/?] n
Unix98 PTY support (UNIX98_PTY) [Y/n/?] y
  Support multiple instances of devpts (DEVPTS_MULTIPLE_INSTANCES) [N/y/?] n
Legacy (BSD) PTY support (LEGACY_PTY) [N/y/?] n
Hardware Random Number Generator Core support (HW_RANDOM) [Y/n/m/?] y
  Timer IOMEM HW Random Number Generator support (HW_RANDOM_TIMERIOMEM) [N/m/y/?] n
Siemens R3964 line discipline (R3964) [N/m/y/?] n
RAW driver (/dev/raw/rawN) (RAW_DRIVER) [N/m/y/?] n
```

If you get this, run  
**make xconfig**  
and save

## Module Build Output – Out-of-tree

```
host$ cd ../drivers/char/examples
host$ vi Makefile
ifneq ($(KERNELRELEASE),)
obj-$(CONFIG_EXAMPLES) += hello1.o
else
KDIR := ~/BeagleBoard/bb-kernel/KERNEL
all:
$(MAKE) -C $(KDIR) M=$$PWD
endif
• Compile with
• host$ make
```

## Once built... Option 1

On the Beagle...Two choices....

- Option 1:  
make INSTALL\_MOD\_PATH=~/BeagleBoard modules\_install
- Will create lib directory in ~/BeagleBoard with everything that goes in /lib on the Beagle

```
host$ ls -F ~/BeagleBoard/lib/modules/3.8.13+/
```

builds	modules.dep.bin	modules.serimap
kernel/	modules.devname	modules.softdep
modules.alias	modules.ieee1394map	modules.symbols
modules.alias.bin	modules.inputmap	modules.symbols.bin
modules.builtin	modules.isapnpmap	modules.usbmap
modules.builtin.bin	modules.ofmap	sources
modules.covmap	modules.order	
modules.dep	modules.pcimap	

- Then  
host\$ rm build source  
host\$ scp -r ~/BeagleBoard/lib root@beagle:/lib
- Could take a while to transfer

## Once built... Option 2

- Just copy the new file you created  
host\$ scp ../drivers/char/examples/hello1.ko root@beagle:..
- On the Beagle  
bone\$ uname -r  
3.8.13-bone67  
bone\$ cd /lib/modules/\$(uname -r)/kernel/drivers/char/  
bone\$ mkdir examples  
bone\$ cd examples  
bone\$ mv ../hello1.ko .
- Now build a new dependencies file  
beagle\$ depmod -a

Note backquotes

## Loading and Unloading a Module

```
beagle$ modprob hello1
beagle$ dmesg | tail -4
[ 9.106206] snd-usb-audio 1-1:1.0: usb_probe_interface
[ 9.106244] snd-usb-audio 1-1:1.0: usb_probe_interface - got id
[ 9.813239] usbcore: registered new interface driver snd-usb-audio
[ 109.308551] Hello Example Init
beagle$ rmmod hello1
beagle$ dmesg | tail -4
[ 9.106244] snd-usb-audio 1-1:1.0: usb_probe_interface - got id
[ 9.813239] usbcore: registered new interface driver snd-usb-audio
[ 109.308551] Hello Example Init
[ 241.037368] Hello Example Exit
```

## Module Utilities

```
$ insmod /lib/modules/`uname -r`/kernel/drivers/char/examples/hello1.ko
```

- No need build dependencies file

## Example Driver with Parameter

```
/* Example Minimal Character Device Driver */
#include <linux/module.h>

static int debug_enable = 0;
module_param(debug_enable, int, 0);

MODULE_PARM_DESC(debug_enable, "Enable module debug mode.");

static int __init hello_init(void) {
    /* Now print value of new module parameter */
    printk("Hello Example Init - debug mode is %s\n",
           debug_enable ? "enabled" : "disabled");
    return 0;
}

static void __exit hello_exit(void) {
    printk("Hello Example Exit\n");
}

module_init(hello_init);
module_exit(hello_exit);

MODULE_AUTHOR("Chris Hallinan");
MODULE_DESCRIPTION("Hello World Example");
MODULE_LICENSE("GPL");
```

## Passing Parameters to a Module

```
beagle$ modprobe hello2 debug_enable=1
Hello Example Init - debug mode is enabled
beagle$ rmmod hello2
beagle$ modprobe hello2
Hello Example Init - debug mode is disabled
```

## Other module commands

```
# /sbin/lsmmod
# /sbin/modinfo hello1
# /sbin/rmmod hello1
# /sbin/depmod (creates modules.dep.bin)
```

- Go play with them

## Adding File System Ops to Hello.c

- Section 8.3, page 217 has a long example about adding file system operations to **hello.c**
- Look it over
- Creates a new device (**/dev/hello1**)
- You can read and write it
- Do it

## Driver File System Operations

- Once a device driver is loaded into the live kernel...
  - **open()** is used to prepare it for subsequent operations
  - **release()** is used to clean up
  - **ioctl()** is used for nonstandard communication
- Think in terms of reading and writing a file...

```
fd = open("file", ...
read(fd, ...
close(fd)
```

## open/release additions to hello.c

```
#include <linux/fs.h>

#define HELLO_MAJOR 234
...
struct file_operations hello_fops;

static int hello_open(struct inode *inode, struct file *file) {
    printk("hello_open: successful\n");
    return 0;
}

static int hello_release(struct inode *inode, struct file *file) {
    printk("hello_release: successful\n");
    return 0;
}
```

## read/write additions to hello.c

```
static ssize_t hello_read(struct file *file,
    char *buf, size_t count, loff_t *ptr) {
    printk("hello_read: returning zero bytes\n");
    return 0;
}

static ssize_t hello_write(struct file *file,
    const char *buf, size_t count, loff_t *ppos)
{
    printk("hello_read: accepting zero bytes\n");
    return 0;
}
```

## ioctl additions to hello.c

```
static long hello_ioctl(struct inode *inode,
    struct file *file, unsigned int cmd,
    unsigned long arg) {
    printk("hello_ioctl: cmd=%ld, arg=%ld\n",
        cmd, arg);
    return 0;
}
```

## init additions to hello.c

```
#define HELLO_MAJOR 234
static int __init hello_init(void)
{
    int ret;
    printk("Hello Example Init - debug mode is %s\n",
        debug_enable ? "enabled" : "disabled");
    ret = register_chrdev(HELLO_MAJOR, "hello!", &hello_fops);
    if (ret < 0) {
        printk("Error registering hello device\n");
        goto hello_fail1;
    }
    printk("Hello: registered module successfully!\n");

    /* Init processing here... */

    return 0;

hello_fail1:
    return ret;
}
```

## Major number for device driver

- Every device has a major and minor number

```
$ ls -ls /dev/console
0 crw----- 1 yoder root 5, 1 2011-02-06 17:57 /dev/console
```

- Device numbers *used* to be statically assigned
- See .../Documentation/devices.txt

```
5 char  Alternate TTY devices
      0 = /dev/tty           Current TTY device
      1 = /dev/console System console
      2 = /dev/ptmx         PTY master multiplex
      64 = /dev/cua0        Callout device for ttyS0
```

- The text uses static assignment

```
234-239          UNASSIGNED
240-254 char     LOCAL/EXPERIMENTAL USE
```

## Registering our functions

- Struct file\_operations is used bind our functions to the requests from the file system.

```
struct file_operations hello_fops = {
    owner:    THIS_MODULE,
    read:     hello_read,
    write:    hello_write,
    ioctl:    hello_ioctl,
    open:     hello_open,
    release:  hello_release,
};
```

## init additions to hello.c

```
#define HELLO_MAJOR 234
static int __init hello_init(void)
{
    int ret;
    printk("Hello Example Init - debug mode is %s\n",
           debug_enable ? "enabled" : "disabled");
    ret = register_chrdev(HELLO_MAJOR, "hello1", &hello_fops);
    if (ret < 0) {
        printk("Error registering hello device\n");
        goto hello_fail1;
    }
    printk("Hello: registered module successfully!\n");
    /* Init processing here... */
    return 0;
hello_fail1:
    return ret;
}
```

## Device Nodes and **mknod**

- Use **mknod** to create a new device

```
$ mknod /dev/hello1 c 234 0
```



- Then

```
$ ls -l /dev/hello1
```

```
crw-r--r-- 1 root root 234, 0 Apr 2 2011 /dev/hello1
```

## Dynamic Major Number

- The above example uses the older *static* method to assign a device number
- Today dynamic allocation is preferred
- Here is how:

```
#include <linux/kdev_t.h>
```

```
dev_t dev;
```

- This declares **dev** to be a device number (both major and minor). Now assign it a value
- dev = MKDEV(234, 0);**

## Requesting a number

- Now request a number

```
#include <linux/fs.h>
```

```
int register_chrdev_region(dev, 4, "hello");
```

- This requests a device number starting with 234 (previous page)
- It asks for 4 minor numbers
- Uses the name "hello"

- When done with the device use:

```
void unregister_chrdev_region(dev, 4);
```

## Using **mknod**

- If your major number is assigned dynamically, how do you use **mknod**? Try the following

```
module="hello"
device="hello"
mode="664"
# remove stale nodes
/sbin/insmod ./module.ko $* || exit 1
rm -f /dev/${device}0
major=`awk "\$2==\"$module\" {print \$1} /proc/devices`
mknod /dev/${device}0 c $major 0
```

## /proc/devices

Character devices:	89 i2c	Block devices:	70 sd
1 mem	90 mtd	1 ramdisk	71 sd
4 /dev/vc/0	116 alsa	259 blkext	128 sd
4 tty	128 ptm	7 loop	129 sd
4 ttyS	136 pts	8 sd	130 sd
5 /dev/tty	153 spi	31 mtblock	131 sd
5 /dev/console	161 ircomm	65 sd	132 sd
5 /dev/ptmx	166 ttyACM	66 sd	133 sd
7 vcs	180 usb	67 sd	134 sd
10 misc	189 usb_device	68 sd	135 sd
13 input	212 DVB	69 sd	179 mmc
14 sound	216 rfcomm		254 device-mapper
29 fb	226 drm		
81 video4linux	234 hello3		
	244 ttyGS		
	245 ttySDIO		
	246 ushmon		
	247 uio		
	248 ttyO		
	249 bsg		

## Assignment

- See [http://elinux.org/EBC\\_Device\\_Drivers](http://elinux.org/EBC_Device_Drivers)

## Module dependencies

- ▶ Some kernel modules can depend on other modules, which need to be loaded first
- ▶ Example: the `usb-storage` module depends on the `scsi_mod`, `libusual` and `usbcore` modules
- ▶ Dependencies are described in `/lib/modules/<kernel-version>/modules.dep`

## `/lib/modules/3.8.13/modules.dep`

```
kernel/lib/raid6/raid6_pq.ko:
kernel/lib/ts_kmp.ko:
kernel/lib/ts_bm.ko:
kernel/lib/ts_fsm.ko:
kernel/lib/notifier-error-inject.ko:
kernel/lib/pm-notifier-error-inject.ko:
kernel/lib/notifier-error-inject.ko
kernel/lib/mpi/mpi.ko:
kernel/lib/asnl_decoder.ko:
kernel/lib/oid_registry.ko:
kernel/drivers/char/examples/hello2.ko:
kernel/drivers/char/examples/hello3.ko:
kernel/drivers/char/examples/hello1.ko:
```

## Kernel log

When a new module is loaded, related information is available in the kernel log

- ▶ The kernel keeps its messages in a circular buffer (so that it doesn't consume more memory with many messages)
- ▶ Kernel log messages are available through the `dmesg` command ("diagnostic message")
- ▶ Kernel log messages are also displayed in the system console (messages can be filtered by level using `/proc/sys/kernel/printk`)

## `printk`

- `/proc/sys/kernel/printk`
- The four values in this file are
  - `console_loglevel`,
  - `default_message_loglevel`,
  - `minimum_console_level` and
  - `default_console_loglevel`.
- These values influence `printk()` behavior when printing or logging error messages
- Messages with a higher priority than `console_loglevel` will be printed to the console
- Messages without an explicit priority will be printed with priority `default_message_level`

<http://www.tin.org/bin/man.cgi?section=5&topic=proc>

## Kernel log levels

0 (KERN_EMERG)	The system is unusable
1 (KERN_ALERT)	Actions that must be taken care of immediately
2 (KERN_CRIT)	Critical conditions
3 (KERN_ERR)	Noncritical error conditions
4 (KERN_WARNING)	Warning conditions that should be taken care of
5 (KERN_NOTICE)	Normal, but significant events
6 (KERN_INFO)	Informational messages that require no action
7 (KERN_DEBUG)	Kernel debugging messages, output by the

## Module utilities (1)

- ▶ `modinfo <module_name>`  
`modinfo <module_path>.ko`  
Gets information about a module: parameters, license, description and dependencies.  
Very useful before deciding to load a module or not.
- ▶ `sudo insmod <module_path>.ko`  
Tries to load the given module. The full path to the module object file must be given.

## Understanding module loading

- ▶ When loading a module fails, `insmod` often doesn't give you enough details!
- ▶ Details are often available in the kernel log
- ▶ Example:  

```
beagle$ sudo insmod ./intr_monitor.ko
insmod: error inserting './intr_monitor.ko': -1
Device or resource busy
beagle$ dmesg
[17549774.552000] Failed to register handler for
irq channel 2
```

## Module utilities (2)

- ▶ `sudo modprobe <module_name>`  
Most common usage of `modprobe`: tries to load all the modules the given module depends on, and then this module. Lots of other options are available. `modprobe` automatically looks in `/lib/modules/<version>/` for the object file corresponding to the given module name.
- ▶ `lsmod`  
Displays the list of loaded modules  
Compare its output with the contents of `/proc/modules`!

## lsmod

```
beagle$ lsmod
Module                  Size  Used by
hello3                  1952  0
iptables_nat            2493  0
nf_conntrack_ipv4      13161  1
nf_defrag_ipv4         1378  1 nf_conntrack_ipv4
nf_nat_ipv4            3487  1 iptable_nat
nf_nat                 16075  2 nf_nat_ipv4, iptable_nat
nf_conntrack           80162  4 nf_nat, nf_nat_ipv4, iptable_nat, nf_conntrack_ipv4
ip_tables              11013  1 iptable_nat
x_tables               16848  1 ip_tables
g_multi                56699  2
libcomposite           17303  1 g_multi
rfcomm                 30945  0
ircomm_tty             17315  0
ircomm                 10198  1 ircomm_tty
irda                   108219  2 ircomm_tty, ircomm
hidp                   12626  0
bluetooth              190961  4 hidp, rfcomm
--
```

## Module utilities (3)

- ▶ `sudo rmmod <module_name>`  
Tries to remove the given module.  
Will only be allowed if the module is no longer in use (for example, no more processes opening a device file)
- ▶ `sudo modprobe -r <module_name>`  
Tries to remove the given module and all dependent modules (which are no longer needed after the module removal)

## Passing parameters to modules

- ▶ Find available parameters:  
`modinfo snd-intel8x0m`
- ▶ Through `insmod`:  
`sudo insmod ./snd-intel8x0m.ko index=-2`
- ▶ Through `modprobe`:  
Set parameters in `/etc/modprobe.conf` or in any file in `/etc/modprobe.d/`:  
`options snd-intel8x0m index=-2`
- ▶ Through the kernel command line,  
when the module is built statically into the kernel:  
`snd-intel8x0m.index=-2`

module name ↑  
module parameter name ↑  
module parameter value —

## Useful reading

Linux Kernel in a Nutshell, Dec 2006

- ▶ By Greg Kroah-Hartman, O'Reilly  
<http://www.kroah.com/lkn/>
- ▶ A good reference book and guide on configuring, compiling and managing the Linux kernel sources.
- ▶ **Freely available on-line!**  
Great companion to the printed book for easy electronic searches!  
Available as single PDF file on  
<http://free-electrons.com/community/kernel/lkn/>
- ▶ In `exercises/pptx`



## Useful reading too

Linux Device Drivers, Third Edition, February 2005

- ▶ By Jonathan Corbet, Alessandro Rubini, Greg Kroah-Hartman, O'Reilly  
<http://lwn.net/Kernel/LDD3/>
- ▶ **Freely available on-line!**  
Great companion to the printed book for easy electronic searches!  
Available as single PDF file
- ▶ LDD3 is current as of the 2.6.10 kernel (Old?)
- ▶ In `exercises/pptx`

