

Day 6-1

Assignment:

- HW6 due Thursday
- Project Proposals

Today's Topics:

- IFTTT
- Google Compute Engine
- Linux Kernel Module

07-2 Device Driver Basics

Using kernel modules

Free Electrons

Loadable kernel modules

- ▶ Modules: add a given functionality to the kernel (drivers, filesystem support, and many others)
- ▶ Can be loaded and unloaded at any time, only when their functionality is needed
- ▶ Useful to keep the kernel image size to the minimum (essential in GNU/Linux distributions for PCs)
- ▶ Also useful to reduce boot time: you don't spend time initializing devices and kernel features that you only need later
- ▶ Caution: once loaded, have full access to the whole kernel address space. No particular protection

Minimal Device Driver

```
/**
 * @file    hello.c
 * @author  Derek Molloy
 * @date    4 April 2015
 * @version 0.1
 * @brief   An introductory "Hello World!" loadable kernel module (LKM) that can display a message
 *          in the /var/log/kern.log file when the module is loaded and removed. The module can accept an
 *          argument when it is loaded -- the name, which appears in the kernel log files.
 * @see http://www.derekmolloy.ie/ for a full description and follow-up descriptions.
 */

#include <linux/init.h>          // Macros used to mark up functions e.g., __init __exit
#include <linux/module.h>        // Core header for loading LKMs into the kernel
#include <linux/kernel.h>        // Contains types, macros, functions for the kernel

MODULE_LICENSE("GPL");           ///< The license type -- this affects runtime behavior
MODULE_AUTHOR("Derek Molloy");   ///< The author -- visible when you use modinfo
MODULE_DESCRIPTION("A simple Linux driver for the BBB."); ///< The description -- see modinfo
MODULE_VERSION("0.1");           ///< The version of the module
```

Minimal Device Driver

```
static char *name = "world";      ///< An example LKM argument -- default value is "world"
module_param(name, charp, S_IRUGO); ///< Param desc. charp = char ptr, S_IRUGO can be read/not
changed
MODULE_PARAM_DESC(name, "The name to display in /var/log/kern.log"); ///< parameter description

/** @brief The LKM initialization function
 * The static keyword restricts the visibility of the function to within this C file. The __init
 * macro means that for a built-in driver (not a LKM) the function is only used at initialization
 * time and that it can be discarded and its memory freed up after that point.
 * @return returns 0 if successful
 */
static int __init helloBBB_init(void){
    printk(KERN_INFO "EBB: Hello %s from the BBB LKM!\n", name);
    return 0;
}
```

Minimal Device Driver

```
/** @brief The LKM cleanup function
 * Similar to the initialization function, it is static. The __exit macro notifies that if this
 * code is used for a built-in driver (not a LKM) that this function is not required.
 */
static void __exit helloBBB_exit(void){
    printk(KERN_INFO "EBB: Goodbye %s from the BBB LKM!\n", name);
}

/** @brief A module must use the module_init() module_exit() macros from linux/init.h, which
 * identify the initialization function at insertion time and the cleanup function (as
 * listed above)
 */
module_init(helloBBB_init);
module_exit(helloBBB_exit);
```

Module Build Output – Out-of-tree

- Load headers for current version of kernel

```
bone$ apt update
bone$ apt install linux-headers-`uname -r`
```

- Clone Molloy's examples

```
bone$ git clone https://github.com/derekmolloy/exploringBB.git
```

- Find hello world example

```
bone$ cd exploringBB/extras/kernel/hello
bone$ cat Makefile
```

```
obj-m+=hello.o
all:
    make -C /lib/modules/$(shell uname -r)/build/ M=$(PWD) modules
clean:
    make -C /lib/modules/$(shell uname -r)/build/ M=$(PWD) clean
```

- Compile with

```
bone$ make
```

Loading and Unloading a Module

```
bone$ insmod hello.ko
bone$ dmesg | tail -4
[  9.106206] snd-usb-audio 1-1:1.0: usb_probe_interface
[  9.106244] snd-usb-audio 1-1:1.0: usb_probe_interface - got id
[  9.813239] usbcore: registered new interface driver snd-usb-
audio
[Oct 7 14:20] EBB: Hello world from the BBB LKM!
bone$ rmmod hello
bone$ dmesg | tail -4
[  9.106244] snd-usb-audio 1-1:1.0: usb_probe_interface - got id
[  9.813239] usbcore: registered new interface driver snd-usb-
audio
[Oct 7 14:20] EBB: Hello world from the BBB LKM!
[ +20.535832] EBB: Goodbye world from the BBB LKM!
```

Example Driver with Parameter

```
///< An example LKM argument -- default value is "world"
static char *name = "world";

///< Param desc. charp = char ptr, S_IRUGO can be read/not changed
module_param(name, charp, S_IRUGO);

///< parameter description
MODULE_PARM_DESC(name, "The name to display in /var/log/kern.log");
```

Passing Parameters to a Module

```
bone$ insmod hello.ko name=Mark
[Oct 7 14:23] EBB: Hello Mark from the BBB LKM!
bone$ rmmod hello
[Oct 7 15:23] EBB: Goodbye Mark from the BBB LKM!
bone$ insmod hello.ko
[Oct 7 15:24] EBB: Hello world from the BBB LKM!
```

Other module commands

```
bone$ lsmod
bone$ modinfo.ko hello
bone$ depmod (creates modules.dep.bin)
```

- Go play with them

Day 6-2

Assignment:

- HW6 due today
- Project Proposals

Today's Topics:

- Linux Kernel Modules – file operations

Adding File System Ops to Hello.c

- <http://derekmolloy.ie/writing-a-linux-kernel-module-part-2-a-character-device/> has a long example about adding file system operations to **hello.c**
- Look it over
- Creates a new device (**/dev/ebbchar**)
- You can read and write it
- Do it

Major and Minor Number

- Every device has a major and minor number
- ```
$ ls -ls /dev/console
```

```
0 crw-rw-r-- 1 yoder root 5, 1 2011-02-06 17:57 /dev/console
```
- Used by the kernel to identify the correct device driver when the device is accessed
- Device numbers *used* to be statically assigned
- See **.../Documentation/devices.txt**

```
5 char Alternate TTY devices
 0 = /dev/tty Current TTY device
 1 = /dev/console System console
 2 = /dev/ptmx PTY master multiplex
 64 = /dev/cua0 Callout device for ttys0
```

- The text uses static assignment

```
234-239 UNASSIGNED
240-254 char LOCAL/EXPERIMENTAL USE
```

## Character Drivers

- Character devices are identified by a 'c'
- Block devices a 'b'

```
bone$ ls -l /dev
crw-rw-r-- 1 root i2c 89, 0 Oct 12 11:10 i2c-0
crw-rw-r-- 1 root i2c 89, 1 Oct 12 11:11 i2c-1
crw-rw-r-- 1 root i2c 89, 2 Oct 12 11:11 i2c-2
drwxr-xr-x 3 root root 100 Oct 12 11:11 input
crw-rw-r-- 1 root kmem 1, 2 Oct 12 11:11 kmem
crw-rw-r-- 1 root root 1, 11 Oct 12 11:11 kmsg
drwxr-xr-x 2 root root 60 Dec 31 1969 lightnvm
crw-rw-r-- 1 root disk 10, 237 Oct 12 11:10 loop-control
drwxr-xr-x 2 root root 60 Oct 12 11:10 mapper
crw-rw-r-- 1 root kmem 1, 1 Oct 12 11:11 mem
crw-rw-r-- 1 root root 10, 57 Oct 12 11:11 memory_bandwidth
brw-rw-r-- 1 root disk 179, 0 Oct 12 11:11 mmcblk0
brw-rw-r-- 1 root disk 179, 1 Oct 12 11:11 mmcblk0p1
brw-rw-r-- 1 root disk 179, 8 Oct 12 11:11 mmcblk1
brw-rw-r-- 1 root disk 179, 16 Oct 12 11:11 mmcblk1boot0
```

## Assigning Device Numbers

- You can manually create a device file and associate it with your device
- ```
bone$ mknod /dev/test c 92 1
```
- You have to make sure the device (92) isn't in use.
- Look in `/usr/src/linux/include/uapi/linux/major.h`
- But there is a better way...

File Operations Data Structure

- The `file_operations` data structure holds pointers to functions within a driver that allows you to define the behavior of certain file operations
- It is defined in `.../include/linux/fs.h`

```
1 // Note: __user refers to a user-space address.
2 struct file_operations {
3     struct module *owner;           // Pointer to the LKM that owns the structure
4     loff_t (*llseek) (struct file *, loff_t, int); // Change current read/write position in a file
5     ssize_t (*read) (struct file *, char __user *, size_t, loff_t *); // Used to retrieve data from the
6     ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *); // Used to send data to the
7     ssize_t (*read_iter) (struct kiocb *, const struct iovec *, unsigned long, loff_t *); // Asynchronous read
8     ssize_t (*write_iter) (struct kiocb *, const struct iovec *, unsigned long, loff_t *); // Asynchronous write
9     ssize_t (*read_iter) (struct kiocb *, struct iov_iter *); // possibly asynchronous read
10    ssize_t (*write_iter) (struct kiocb *, struct iov_iter *); // possibly asynchronous write
11    int (*iterate) (struct file *, struct dir_context *); // called when vfs needs to read
12    unsigned int (*poll) (struct file *, struct poll_table_struct *); // Does a read or write block?
13    long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long); // Called by the ioctl system call
14    long (*compat_ioctl) (struct file *, unsigned int, unsigned long); // Called by the ioctl system call
15    int (*mmap) (struct file *, struct vm_area_struct *); // Called by mmap system call
16    int (*mremap) (struct file *, struct vm_area_struct *); // Called by memory remap system call
17    int (*open) (struct inode *, struct file *); // first operation performed on a device file
18    int (*flush) (struct file *, fl_owner_t id); // called when a process closes its copy of a
19    int (*release) (struct inode *, struct file *); // called when a file structure is being refile
20    int (*fsync) (struct file *, loff_t, loff_t, int datasync); // notify device of change in its PASTBC
21    int (*fasync) (struct kiocb *, int datasync); // synchronous notify device of change in its
22    int (*fasync) (int, struct file *, int); // asynchronous notify device of change in its
23    int (*lock) (struct file *, int, struct file_lock *); // used to implement file locking
24};
```

Driver File System Operations

- Once a device driver is loaded into the live kernel...
 - **open()** is called each time the device is opened from user space
 - **read()** is called when data is sent from the device to user space
 - **write()** is called when data is sent from user space to the device
 - **release()** is called when the device is closed in user space

- Think in terms of reading and writing a file...

```
fd = open("file", ...
read(fd, ...
write(fd, ...
close(fd)
```

open/release additions to hello.c

```
static int    majorNumber;    ///< Stores the device number -- determined automatically
// The prototype functions for the character driver -- must come before the struct definition
static int    dev_open(struct inode *, struct file *);
static int    dev_release(struct inode *, struct file *);
static ssize_t dev_read(struct file *, char *, size_t, loff_t *);
static ssize_t dev_write(struct file *, const char *, size_t, loff_t *);

/** Devices are represented as file structure in the kernel. The file_operations structure from
 * /linux/fs.h lists the callback functions that you wish to associated with your file operations
 * using a C99 syntax structure. char devices usually implement open, read, write and release calls
 */
static struct file_operations fops =
{
    .open = dev_open,
    .read = dev_read,
    .write = dev_write,
    .release = dev_release,
};
```

From: [exploringBB/extras/kernel/ebbchar](#)

ebbchar_init

```
#define DEVICE_NAME "ebbchar"    ///< The device will appear at /dev/ebbchar using this value
#define CLASS_NAME  "ebb"    ///< The device class -- this is a character device driver

static int __init ebbchar_init(void){
    printk(KERN_INFO "EBBChar: Initializing the EBBChar LKM!\n");

    // Try to dynamically allocate a major number for the device -- more difficult but worth it
    majorNumber = register_chrdev(0, DEVICE_NAME, &fops);
    if (majorNumber<0){
        printk(KERN_ALERT "EBBChar failed to register a major number!\n");
        return majorNumber;
    }
    printk(KERN_INFO "EBBChar: registered correctly with major number %d\n", majorNumber);

    // Register the device class
    ebbcharClass = class_create(THIS_MODULE, CLASS_NAME);
    if (IS_ERR(ebbcharClass)){    // Check for error and clean up if there is
        unregister_chrdev(majorNumber, DEVICE_NAME);
        printk(KERN_ALERT "Failed to register device class!\n");
        return PTR_ERR(ebbcharClass);    // Correct way to return an error on a pointer
    }
    printk(KERN_INFO "EBBChar: device class registered correctly!\n");
}
```

ebbchar_init

```
// Register the device driver
ebbcharDevice = device_create(ebbcharClass, NULL,
    MKDEV(majorNumber, 0), NULL, DEVICE_NAME);
// Clean up if there is an error
if(IS_ERR(ebbcharDevice)){
    class_destroy(ebbcharClass);
    unregister_chrdev(majorNumber, DEVICE_NAME);
    printk(KERN_ALERT "Failed to create the device!\n");
    return PTR_ERR(ebbcharDevice);
}
printk(KERN_INFO "EBBChar: device class created
correctly!\n"); // Made it! device was initialized
return 0;
}
```

ebbchar_exit

```
static void __exit ebbchar_exit(void){
    // remove the device
    device_destroy(ebbcharClass, MKDEV(majorNumber, 0));
    // unregister the device class
    class_unregister(ebbcharClass);
    // remove the device class
    class_destroy(ebbcharClass);
    // unregister the major number
    unregister_chrdev(majorNumber, DEVICE_NAME);
    printk(KERN_INFO "EBBChar: Goodbye from the LKM!\n");
}
```

dev_open/dev_release

```
static int dev_open(struct inode *inodep, struct file *filep){
    numberOpens++;
    printk(KERN_INFO "EBBChar: Device has been opened %d time(s)\n", numberOpens);
    return 0;
}

static int dev_release(struct inode *inodep, struct file *filep){
    printk(KERN_INFO "EBBChar: Device successfully closed!\n");
    return 0;
}
```

dev_write

```
static ssize_t dev_write(struct file *filep,
    const char *buffer, size_t len, loff_t *offset){
    // appending received string with its length
    sprintf(message, "%s(%d letters)", buffer, len);
    // store the length of the stored message
    size_of_message = strlen(message);
    printk(KERN_INFO "EBBChar: Received %d characters from the
user\n", len);
    return len;
}
```

dev_read

```
static ssize_t dev_read(struct file *file, char *buffer, size_t len,
                       loff_t *offset){
    int error_count = 0;
    // copy_to_user has the format ( * to, *from, size) and returns 0 on success
    error_count = copy_to_user(buffer, message, size_of_message);

    if (error_count==0){           // if true then have success
        printk(KERN_INFO "EBBChar: Sent %d characters to the user\n",
               size_of_message);
        return (size_of_message); // clear the position to the start and return 0
    }
    else {
        printk(KERN_INFO "EBBChar: Failed to send %d characters to the user\n",
               error_count);
        return -EFAULT;          // Failed -- return a bad address message (i.e. -14)
    }
}
```

/proc/devices

Character devices:	116 alsa	Block devices:	128 sd
	128 ptm	259 blkext	129 sd
1 mem	136 pts	8 sd	130 sd
4 /dev/vc/0	153 spi	65 sd	131 sd
4 tty	180 usb	66 sd	132 sd
4 ttys	189 usb_device	67 sd	133 sd
5 /dev/tty	212 DVB	68 sd	134 sd
5 /dev/console	226 drm	69 sd	135 sd
5 /dev/ptmx	245 ebbchar	70 sd	179 mmc
7 vcs	246 uio	71 sd	
10 misc	247 ttyGS		
13 input	248 hidraw		
29 fb	249 bsg		
81 video4linux	250 watchdog		
89 i2c	251 ptp		
90 mtd	252 pps		
	253 media		
	254 rtc		

Module dependencies

- ▶ Some kernel modules can depend on other modules, which need to be loaded first
- ▶ Example: the `usb-storage` module depends on the `scsi_mod`, `libusual` and `usbcore` modules
- ▶ Dependencies are described in `/lib/modules/<kernel-version>/modules.dep`

/lib/modules/4.4.21-ti-r47/modules.dep

```
kernel/arch/arm/crypto/aes-arm.ko:
kernel/arch/arm/crypto/aes-arm-bs.ko:
    kernel/arch/arm/crypto/aes-arm.ko
    kernel/crypto/ablk_helper.ko
    kernel/crypto/cryptd.ko
kernel/arch/arm/crypto/sha1-arm.ko:
kernel/arch/arm/crypto/sha1-arm-neon.ko:
kernel/arch/arm/crypto/sha1-arm.ko
kernel/arch/arm/crypto/sha256-arm.ko:
kernel/arch/arm/crypto/sha512-arm.ko:
```

Kernel log

When a new module is loaded, related information is available in the kernel log

- ▶ The kernel keeps its messages in a circular buffer (so that it doesn't consume more memory with many messages)
- ▶ Kernel log messages are available through the `dmesg` command ("diagnostic message")
- ▶ Kernel log messages are also displayed in the system console (messages can be filtered by level using `/proc/sys/kernel/printk`)

printk

- `/proc/sys/kernel/printk`
- The four values in this file are
 - `console_loglevel`,
 - `default_message_loglevel`,
 - `minimum_console_level` and
 - `default_console_loglevel`.
- These values influence `printk()` behavior when printing or logging error messages
- Messages with a higher priority than `console_loglevel` will be printed to the console
- Messages without an explicit priority will be printed with priority `default_message_level`

<http://www.tin.org/bin/man.cgi?section=5&topic=proc>

Kernel log levels

0 (KERN_EMERG)	The system is unusable
1 (KERN_ALERT)	Actions that must be taken care of immediately
2 (KERN_CRIT)	Critical conditions
3 (KERN_ERR)	Noncritical error conditions
4 (KERN_WARNING)	Warning conditions that should be taken care of
5 (KERN_NOTICE)	Normal, but significant events
6 (KERN_INFO)	Informational messages that require no action
7 (KERN_DEBUG)	Kernel debugging messages, output by the

Useful reading

Linux Kernel in a Nutshell, Dec 2006

- ▶ By Greg Kroah-Hartman, O'Reilly
<http://www.kroah.com/lkn/>
- ▶ A good reference book and guide on configuring, compiling and managing the Linux kernel sources.
- ▶ **Freely available on-line!**
Great companion to the printed book for easy electronic searches!
Available as single PDF file on
<http://free-electrons.com/community/kernel/lkn/>
- ▶ In **exercises/pptx**



Useful reading too

Linux Device Drivers, Third Edition, February 2005

- ▶ By Jonathan Corbet, Alessandro Rubini, Greg Kroah-Hartman, O'Reilly
<http://lwn.net/Kernel/LDD3/>
- ▶ **Freely available on-line!**
Great companion to the printed book for easy electronic searches!
Available as single PDF file
- ▶ LDD3 is current as of the 2.6.10 kernel (Old?)
- ▶ In **exercises/pptx**

