

## 07-3 Device Driver Basics

Using kernel modules

Free Electrons

### Loadable kernel modules

- ▶ Modules: add a given functionality to the kernel (drivers, filesystem support, and many others)
- ▶ Can be loaded and unloaded at any time, only when their functionality is need
- ▶ Useful to keep the kernel image size to the minimum (essential in GNU/Linux distributions for PCs)
- ▶ Also useful to reduce boot time: you don't spent time initializing devices and kernel features that you only need later
- ▶ Caution: once loaded, have full access to the whole kernel address space. No particular protection

### Minimal Device Driver (Listing 8-1)

```
/* Example Minimal Character Device Driver */
#include <linux/module.h>
static int __init hello_init(void) {
    printk(KERN_INFO "Hello Example Init\n");
    return 0;
}
static void __exit hello_exit(void) {
    printk("Hello Example Exit\n");
}
module_init(hello_init);
module_exit(hello_exit);

MODULE_AUTHOR("Chris Hallinan");
MODULE_DESCRIPTION("Hello World Example");
MODULE_LICENSE("GPL");
```

### Module Build Infrastructure

1. Starting from the top-level Linux source directory, create a directory under `.../drivers/char` called **examples**.
2. Add a menu item to the kernel configuration to enable building **examples** and to specify built-in or loadable kernel module. See Section 8.1.4, page 205
3. Add the new examples subdirectory to the `.../drivers/char/Makefile` conditional on the menu item created in step 2.
4. Create a **Makefile** for the new **examples** directory, and add the **hello1.o** module object to be compiled conditional on the menu item created in step 2.
5. Finally, create the driver **hello1.c** source file from Listing 8-1.

### Typo page 206

```
diff --git a/drivers/char/Kconfig b/drivers/char/Kconfig
index 6f31c94..0805290 100644
--- a/drivers/char/Kconfig
+++ b/drivers/char/Kconfig
@@ -4,6 +4,13 @@
```

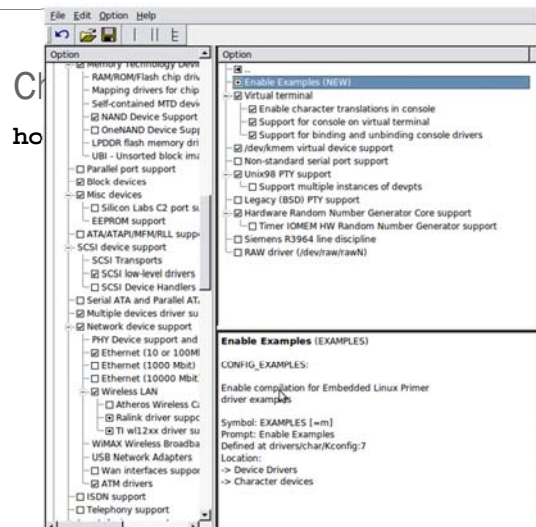
```
menu "Character devices"

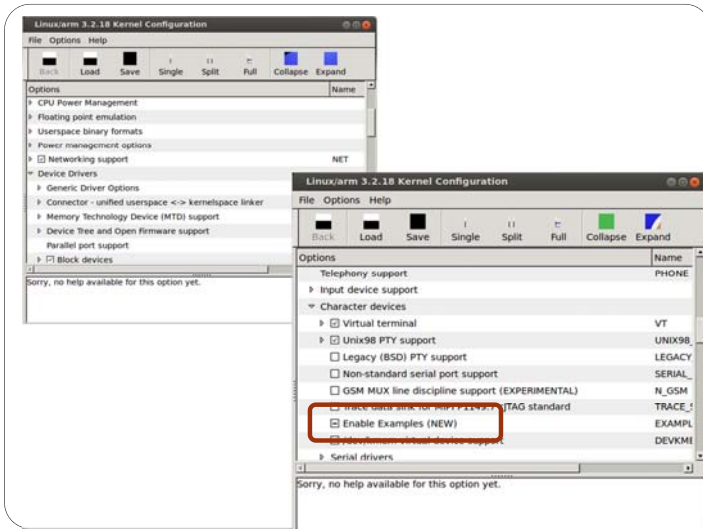
+config EXAMPLES
+    tristate "Enable Examples"
+    default m
+    ---help---
+    Enable compilation option for Embedded Linux Primer
+    driver examples

config DEVMEM
    bool "/dev/kmem virtual device support"
    default y
    help
```

Must be lower case

In new kernel





## Module Build Output

```
host$ time make modules
CHK    include/linux/version.h
CHK    include/generated/utsrelease.h
make[1]: `include/generated/mach-types.h' is up to date.
CALL   scripts/checksyscalls.sh
CC [M] drivers/char/examples/hello1.o
Building modules, stage 2.
MODPOST 1585 modules
WARNING: modpost: Found 1 section mismatch(es)
To see full details build your kernel with:
'make CONFIG_DEBUG_SECTION_MISMATCH=y'
LD [M] drivers/char/examples/hello1.ko

real    0m53.088s
user    0m36.258s
sys     0m7.724s
```

I tried this. No help!

First time takes much longer  
Try make -jX modules

## Once built... Option 1

On the Beagle... Two choices....

- Option 1:  
make INSTALL\_MOD\_PATH=~/BeagleBoard modules\_install
- Will create lib directory in ~/BeagleBoard with everything that goes in /lib on the Beagle

```
host$ ls -F ~/BeagleBoard/lib/modules/3.2.18/
```

```
build#          modules.dep.bin      modules.seriomap
kernel/         modules.devname      modules.softdep
modules.alias   modules.ieee1394map modules.symbols.bin
modules.alias.bin modules.inputmap  modules.symbols
modules.builtin modules.isapnpmap modules.usbmap
modules.builtin.bin modules.ofmap      source#
modules.cowmap  modules.order
modules.dep     modules.pcimap
```

- Then  
host\$ rm build source  
host\$ scp -r ~/BeagleBoard/lib root@beagle:/lib
- Could take a while to transfer

## Once built... Option 2

- Just copy the new file you created

```
host$ scp .../drivers/char/examples/hello1.ko root@beagle:.
```

- On the Beagle

```
beagle$ mv hello1.ko
/lib/modules/2.6.32/kernel/drivers/char/examples/
beagle$ cd /lib/modules/2.6.32
beagle$ mv modules.dep.bin modules.dep.bin.orig
beagle$ gedit modules.dep
```

- add:

```
kernel/drivers/char/examples/hello1.ko:
```

## Loading and Unloading a Module

```
beagle$ /sbin/modprobe hello1
beagle$ dmesg | tail -4
[ 47.095764] OMAPFB: ioctl QUERY_PLANE
[ 47.095794] OMAPFB: ioctl GET_CAPS
[ 49.005889] eth0: no IPv6 routers present
[ 651.947784] Hello Example Init
beagle$ /sbin/modprobe -r hello1
beagle$ dmesg | tail -4
[ 47.095794] OMAPFB: ioctl GET_CAPS
[ 49.005889] eth0: no IPv6 routers present
[ 651.947784] Hello Example Init
[ 677.682769] Hello Example Exit
```

## Module Utilities

```
$ insmod /lib/modules/`uname -r`/kernel/drivers/char/examples/hello1.ko
```

- No need to edit modules.dep

## Example Driver with Parameter

```
/* Example Minimal Character Device Driver */
#include <linux/module.h>

static int debug_enable = 0;
module_param(debug_enable, int, 0);

MODULE_PARM_DESC(debug_enable, "Enable module debug mode.");

static int __init hello_init(void) {
    /* Now print value of new module parameter */
    printk("Hello Example Init - debug mode is %s\n",
           debug_enable ? "enabled" : "disabled");
    return 0;
}

static void __exit hello_exit(void) {
    printk("Hello Example Exit\n");
}

module_init(hello_init);
module_exit(hello_exit);

MODULE_AUTHOR("Chris Hallinan");
MODULE_DESCRIPTION("Hello World Example");
MODULE_LICENSE("GPL");
```

## Passing Parameters to a Module

```
insmod /lib/modules/.../examples/hello1.ko
debug_enable=1
```

Hello Example Init - debug mode is enabled

```
insmod /lib/modules/.../examples/hello1.ko
```

Hello Example Init - debug mode is disabled

## Other module commands

```
# /sbin/lsmmod
# /sbin/modinfo hello1
# /sbin/rmmod hello1
# /sbin/depmod (creates modules.dep.bin)
```

- Go play with them

## Adding File System Ops to Hello.c

- Section 8.3, page 217 has a long example about adding file system operations to **hello.c**
- Look it over
- Creates a new device (**/dev/hello1**)
- You can read and write it
- Do it.

## Driver File System Operations

- Once a device driver is loaded into the live kernel...
  - **open()** is used to prepare it for subsequent operations
  - **release()** is used to clean up
  - **ioctl()** is used for nonstandard communication
- Think in terms of reading and writing a file...

```
fd = open("file", ...
read(fd, ...
close(fd)
```

## open/release additions to hello.c

```
#include <linux/fs.h>

#define HELLO_MAJOR 234
...
struct file_operations hello_fops;

static int hello_open(struct inode *inode, struct file *file) {
    printk("hello_open: successful\n");
    return 0;
}

static int hello_release(struct inode *inode, struct file *file) {
    printk("hello_release: successful\n");
    return 0;
}
```

## read/write additions to hello.c

```
static ssize_t hello_read(struct file *file,
    char *buf, size_t count, loff_t *ptr) {
    printk("hello_read: returning zero bytes\n");
    return 0;
}

static ssize_t hello_write(struct file *file,
    const char *buf, size_t count, loff_t * ppos)
{
    printk("hello_read: accepting zero bytes\n");
    return 0;
}
```

## ioctl additions to hello.c

```
static int hello_ioctl(struct inode *inode,
    struct file *file, unsigned int cmd,
    unsigned long arg) {
    printk("hello_ioctl: cmd=%ld, arg=%ld\n",
        cmd, arg);
    return 0;
}
```

## init additions to hello.c

```
#define HELLO_MAJOR 234
static int __init hello_init(void)
{
    int ret;
    printk("Hello Example Init - debug mode is %s\n",
        debug_enable ? "enabled" : "disabled");
    ret = register_chrdev(HELLO_MAJOR, "hello1", &hello_fops);
    if (ret < 0) {
        printk("Error registering hello device\n");
        goto hello_fail1;
    }
    printk("Hello: registered module successfully!\n");

    /* Init processing here... */

    return 0;

hello_fail1:
    return ret;
}
```

## Major number for device driver

- Every device has a major and minor number

```
$ ls -ls /dev/console
0 crw----- 1 yoder root 5, 1 2011-02-06 17:57 /dev/console
```

- Device numbers *used* to be statically assigned
- See [.../Documentation/devices.txt](#)

```
5 char  Alternate TTY devices
        0 = /dev/tty           Current TTY device
        1 = /dev/console System console
        2 = /dev/ptmx          PTY master multiplex
        64 = /dev/cua0         Callout device for ttys0
```

- The text uses static assignment

```
234-239          UNASSIGNED
240-254 char     LOCAL/EXPERIMENTAL USE
```

## Registering our functions

- Struct `file_operations` is used bind our functions to the requests from the file system.

```
struct file_operations hello_fops = {
    owner:    THIS_MODULE,
    read:     hello_read,
    write:    hello_write,
    ioctl:    hello_ioctl,
    open:     hello_open,
    release:  hello_release,
};
```

## init additions to hello.c

```
#define HELLO_MAJOR 234
static int __init hello_init(void)
{
    int ret;
    printk("Hello Example Init - debug mode is %s\n",
        debug_enable ? "enabled" : "disabled");
    ret = register_chrdev(HELLO_MAJOR, "hello1", &hello_fops);
    if (ret < 0) {
        printk("Error registering hello device\n");
        goto hello_fail1;
    }
    printk("Hello: registered module successfully!\n");
    /* Init processing here... */
    return 0;
hello_fail1:
    return ret;
}
```

## Device Nodes and **mknod**

- Use **mknod** to create a new device

```
$ mknod /dev/hello1 c 234 0
```



- Then

```
$ ls -l /dev/hello1  
crw-r--r--  1 root  root  234, 0 Apr 2 2011  /dev/hello1
```

## Dynamic Major Number

- The above example uses the older *static* method to assign a device number
- Today dynamic allocation is preferred
- Here is how:

```
#include <linux/kdev_t.h>  
dev_t dev;
```

- This declares **dev** to be a device number (both major and minor). Now assign it a value
- **dev = MKDEV(234, 0);**

## Requesting a number

- Now request a number
- ```
#include <linux/fs.h>  
int register_chrdev_region(dev, 4, "hello");
```
- This requests a device number starting with 234 (previous page)
  - It asks for 4 minor numbers
  - Uses the name "hello"
- When done with the device use:
- ```
void unregister_chrdev_region(dev, 4);
```

## Using **mknod**

- If your major number is assigned dynamically, how do you use **mknod**? Try the following
- ```
module="hello"  
device="hello"  
mode="664"  
# remove stale nodes  
/sbin/insmod ./module.ko $* || exit 1  
rm -f /dev/${device}0  
major=`awk "\\$2==\"$module\" {print \\$1} /proc/devices`  
mknod /dev/${device}0 c $major 0
```

## /proc/devices

|                    |                    |                |         |
|--------------------|--------------------|----------------|---------|
| Character devices: | 89 i2c             | Block devices: | 70 sd   |
| 1 mem              | 90 mtd             | 1 ramdisk      | 71 sd   |
| 4 /dev/vc/0        | 116 alsa           | 259 blkext     | 128 sd  |
| 4 tty              | 128 ptm            | 7 loop         | 129 sd  |
| 4 ttyS             | 136 pts            | 8 sd           | 130 sd  |
| 5 /dev/tty         | 153 spi            | 11 sr          | 131 sd  |
| 5 /dev/console     | 161 ircomm         | 31 mtblock     | 132 sd  |
| 5 /dev/ptmx        | 180 usb            | 65 sd          | 133 sd  |
| 7 vcs              | 189 usb_device     | 66 sd          | 134 sd  |
| 10 misc            | 216 rfcomm         | 67 sd          | 135 sd  |
| 13 input           | 247 bccat          | 68 sd          | 179 mmc |
| 14 sound           | 248 pvrsrvkm       | 69 sd          |         |
| 21 sg              | 249 rtc            |                |         |
| 29 fb              | 250 ttySDIO        |                |         |
| 81 video4linux     | 251 omap-resizer   |                |         |
|                    | 252 omap-previewer |                |         |
|                    | 253 ushmon         |                |         |
|                    | 254 bsg            |                |         |

## Assignment

- See [http://elinux.org/EBC\\_Exercise\\_26\\_Device\\_Drivers](http://elinux.org/EBC_Exercise_26_Device_Drivers)

## Module dependencies

- ▶ Some kernel modules can depend on other modules, which need to be loaded first
- ▶ Example: the `usb-storage` module depends on the `scsi_mod`, `libusual` and `usbcore` modules
- ▶ Dependencies are described in `/lib/modules/<kernel-version>/modules.dep`

## `/lib/modules/2.6.32/models.dep`

```
kernel/drivers/char/examples/hello1.ko:
kernel/crypto/twofish_common.ko:
kernel/crypto/ctr.ko:
kernel/crypto/blowfish.ko:
kernel/crypto/ghash-generic.ko:
kernel/crypto/gf128mul.ko:
kernel/crypto/xts.ko:
kernel/crypto/gf128mul.ko:
kernel/crypto/gcm.ko:
kernel/crypto/cryptd.ko:
kernel/crypto/md4.ko:
kernel/crypto/lrw.ko:
kernel/crypto/gf128mul.ko
```

## Kernel log

When a new module is loaded, related information is available in the kernel log

- ▶ The kernel keeps its messages in a circular buffer (so that it doesn't consume more memory with many messages)
- ▶ Kernel log messages are available through the `dmesg` command ("diagnostic message")
- ▶ Kernel log messages are also displayed in the system console (messages can be filtered by level using `/proc/sys/kernel/printk`)

## `printk`

- `/proc/sys/kernel/printk`
- The four values in this file are
  - `console_loglevel`,
  - `default_message_loglevel`,
  - `minimum_console_level` and
  - `default_console_loglevel`.
- These values influence `printk()` behavior when printing or logging error messages
- Messages with a higher priority than `console_loglevel` will be printed to the console
- Messages without an explicit priority will be printed with priority `default_message_level`

<http://www.tin.org/bin/man.cgi?section=5&topic=proc>

## Kernel log levels

|                  |                                                 |
|------------------|-------------------------------------------------|
| 0 (KERN_EMERG)   | The system is unusable                          |
| 1 (KERN_ALERT)   | Actions that must be taken care of immediately  |
| 2 (KERN_CRIT)    | Critical conditions                             |
| 3 (KERN_ERR)     | Noncritical error conditions                    |
| 4 (KERN_WARNING) | Warning conditions that should be taken care of |
| 5 (KERN_NOTICE)  | Normal, but significant events                  |
| 6 (KERN_INFO)    | Informational messages that require no action   |
| 7 (KERN_DEBUG)   | Kernel debugging messages, output by the        |

## Module utilities (1)

- ▶ `modinfo <module_name>`  
`modinfo <module_path>.ko`  
Gets information about a module: parameters, license, description and dependencies. Very useful before deciding to load a module or not.
- ▶ `sudo insmod <module_path>.ko`  
Tries to load the given module. The full path to the module object file must be given.

## Understanding module loading

- ▶ When loading a module fails, `insmod` often doesn't give you enough details!
- ▶ Details are often available in the kernel log
- ▶ Example:

```
beagle$ sudo insmod ./intr_monitor.ko
insmod: error inserting './intr_monitor.ko': -1
Device or resource busy
beagle$ dmesg
[17549774.552000] Failed to register handler for
irq channel 2
```

## Module utilities (2)

- ▶ `sudo modprobe <module_name>`  
Most common usage of `modprobe`: tries to load all the modules the given module depends on, and then this module. Lots of other options are available. `modprobe` automatically looks in `/lib/modules/<version>/` for the object file corresponding to the given module name.
- ▶ `lsmod`  
Displays the list of loaded modules  
Compare its output with the contents of `/proc/modules`!

## lsmod

```
beagle$ lsmod
Module                Size  Used by
bufferclass_ti        4768  0
omaplfb               8733  0
pvrsrvkm             154248  2 bufferclass_ti,omaplfb
rfcomm               33484  0
ircomm_tty           30305  0
ircomm               16429  1 ircomm_tty
irda                 162973  2 ircomm_tty,ircomm
ipv6                 249063  14
hidp                 11193  0
l2cap                30104  4 rfcomm,hidp,l2cap
bluetooth            49221  3 rfcomm,hidp,l2cap
...
```

## Module utilities (3)

- ▶ `sudo rmmod <module_name>`  
Tries to remove the given module.  
Will only be allowed if the module is no longer in use (for example, no more processes opening a device file)
- ▶ `sudo modprobe -r <module_name>`  
Tries to remove the given module and all dependent modules (which are no longer needed after the module removal)

## Passing parameters to modules

- ▶ Find available parameters:  
`modinfo snd-intel8x0m`
- ▶ Through `insmod`:  
`sudo insmod ./snd-intel8x0m.ko index=-2`
- ▶ Through `modprobe`:  
Set parameters in `/etc/modprobe.conf` or in any file in `/etc/modprobe.d/`:  
`options snd-intel8x0m index=-2`
- ▶ Through the kernel command line,  
when the module is built statically into the kernel:  
`snd-intel8x0m.index=-2`

↑ module name  
↑ module parameter name  
└─ module parameter value

## Useful reading

- Linux Kernel in a Nutshell, Dec 2006
- ▶ By Greg Kroah-Hartman, O'Reilly  
<http://www.kroah.com/lkn/>
- ▶ A good reference book and guide on configuring, compiling and managing the Linux kernel sources.
- ▶ **Freely available on-line!**  
Great companion to the printed book for easy electronic searches!  
Available as single PDF file on  
<http://free-electrons.com/community/kernel/lkn/>



## Useful reading too

Linux Device Drivers, Third Edition, February 2005

- ▶ By Jonathan Corbet, Alessandro Rubini, Greg Kroah-Hartman, O'Reilly

<http://lwn.net/Kernel/LDD3/>

- ▶ **Freely available on-line!**

Great companion to the printed book  
for easy electronic searches!

Available as single PDF file

- ▶ LDD3 is current as of the 2.6.10 kernel  
(Old?)

