




The Unix and GNU/Linux command line

The Unix and GNU/Linux command line

Michael Opendacker
Thomas Petazzoni
Free Electrons



© Copyright 2009, Free Electrons.
Creative Commons BY-SA 3.0 license
Latest update: 9/5/2017.
Document sources, updates and translations:
<http://free-electrons.com/docs/command-line>
Corrections, suggestions, contributions and translations are welcome!

1

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Command memento sheet



It is a useful companion to this presentation.

Examples for the most useful commands are given in just one sheet.

Suggestions for use
Stick this sheet on your wall, use it as desktop wallpaper, make it a mouse mat, print it on clothing, slice it into bookmarks...

Caution
Store away from mice!

Get it on

http://free-electrons.com/doc/legacy/command-line/command_memento.pdf

2

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Training Contents (1)

Shells, filesystem and file handling

- ▶ [Everything is a file](#)
- ▶ [GNU / Linux filesystem structure](#)
- ▶ [Command line interpreters](#)
- ▶ [Handling files and directories](#)
- ▶ [Displaying, scanning and sorting files](#)
- ▶ [Symbolic and hard link](#)
- ▶ [File access rights](#)

3

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Training contents (2)

Standard I/O, redirections, pipes

- ▶ [Standard input and output, redirecting to files](#)
- ▶ [Pipes: redirecting standard output to other commands](#)
- ▶ [Standard error](#)

4

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Training Contents (3)

Task control

- ▶ [Full control on tasks](#)
- ▶ [Executing in background, suspending, resuming and aborting](#)
- ▶ [List of active tasks](#)
- ▶ [Killing processes](#)
- ▶ [Environment variables](#)
- ▶ [PATH environment variables](#)
- ▶ [Shell aliases, .bashrc file](#)

5

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Training contents (4)

Miscellaneous

- ▶ [Text editors](#)
- ▶ [Compressing and archiving](#)
- ▶ [Printing files](#)
- ▶ [Comparing files and directories](#)
- ▶ [Looking for files](#)
- ▶ [Getting information about users](#)

6

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The Unix and GNU / Linux command line

Unix filesystem



Everything is a file

Almost everything in Unix is a file!

▶ Regular files

▶ Directories

Directories are just files listing a set of files

▶ Symbolic links

Files referring to the name of another file

▶ Devices and peripherals

Read and write from devices as with regular files

▶ Pipes

Used to cascade programs
`cat *.log | grep error`

▶ Sockets

Inter process communication

7

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

8

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



File names

File name features since the beginning of Unix

▶ Case sensitive

▶ No obvious length limit

▶ Can contain any character (including whitespace, except /).

File types stored in the file ("magic numbers").

File name extensions not needed and not interpreted. Just used for user convenience.

▶ File name examples:

`README` `.bashrc`

`Windows Buglist`

`index.htm`

`index.html`

`index.html.old`

9

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



File paths

A *path* is a sequence of nested directories with a file or directory at the end, separated by the `/` character

▶ Relative path: `documents/fun/microsoft_jokes.html`

Relative to the current directory

▶ Absolute path: `/home/bill/bugs/crash9402031614568`

▶ `/` : *root directory*.

Start of absolute paths for all files on the system (even for files on removable devices or network shared).

10

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



GNU / Linux filesystem structure (1)

Not imposed by the system. Can vary from one system to the other, even between two GNU/Linux installations!

<code>/</code>	Root directory
<code>/bin/</code>	Basic, essential system commands
<code>/boot/</code>	Kernel images, initrd and configuration files
<code>/dev/</code>	Files representing devices
	<code>/dev/hda</code> : first IDE hard disk
<code>/etc/</code>	System configuration files
<code>/home/</code>	User directories
<code>/lib/</code>	Basic system shared libraries

11

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



GNU / Linux filesystem structure (2)

`/lost+found`
`/media`
media:

Corrupt files the system tried to recover
Mount points for removable

`/media/usbdisk,`

`/media/cdrom`
`/mnt/`
filesystems
`/opt/`

Mount points for temporarily mounted

Specific tools installed by the sysadmin
`/usr/local/` often used

instead
`/proc/`

Access to system information
`/proc/cpuinfo,`

`/proc/version ...`
`/root/`
`/sbin/`
`/sys/`

root user home directory
Administrator-only commands
System and device controls

12

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



GNU / Linux filesystem structure (3)

<code>/tmp/</code>	Temporary files
<code>/usr/</code> system)	Regular user tools (not essential to the system)
<code>/usr/local/</code>	Specific software installed by the sysadmin (often preferred to <code>/opt/</code>)
<code>/var/</code>	Data used by the system or system servers (incoming print jobs)...
	<code>/usr/bin/</code> , <code>/usr/lib/</code> , <code>/usr/sbin/...</code> <code>/var/log/</code> , <code>/var/spool/mail</code> mail), <code>/var/spool/lpd</code>

The Unix filesystem structure is defined by the Filesystem Hierarchy Standard (FHS):
<http://www.pathname.com/fhs/>

13

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

The Unix and GNU / Linux command line

Shells and file handling

14

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

Command line interpreters

- ▶ Shells: tools to execute user commands
- ▶ Called “shells” because they hide the details on the underlying operating system under the shell's surface.
- ▶ Commands are input in a text terminal, either a window in a graphical environment or a text-only console.
- ▶ Results are also displayed on the terminal. No graphics are needed at all.
- ▶ Shells can be scripted: provide all the resources to write complex programs (variable, conditionals, iterations...)

15

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

Well known shells

Most famous and popular shells


- ▶ **sh**: The Bourne shell (obsolete)
Traditional, basic shell found on Unix systems, by Steve Bourne.
- ▶ **csh**: The C shell (obsolete)
Once popular shell with a C-like syntax
- ▶ **tcsh**: The TC shell (still very popular)
A C shell compatible implementation with evolved features (command completion, history editing and more...)
- ▶ **bash**: The Bourne Again shell (most popular)
An improved implementation of sh with lots of added features too.

16

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

fish: a great new shell

The Friendly Interactive SHell
<http://www.fishshell.org/>



- ▶ Standard features: history, command and file completion...
- ▶ Brand new features: command option completion, command completion with short description, syntax highlighting..
- ▶ Easier to open files: `open` built-in command.
- ▶ Much simpler and consistent syntax (not POSIX compliant)
Makes it easier to create shell scripts.

Command line beginners can learn much faster!
Even experienced users should find this shell very convenient.

17

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

ls command

Lists the files in the current directory, in alphanumeric order, except files starting with the “.” character.

- | | |
|--|--|
| ▶ ls -a (all)
Lists all the files (including . * files) | ▶ ls -S (size)
Lists the biggest files first |
| ▶ ls -l (long)
Long listing (type, date, size, owner, permissions) | ▶ ls -r (reverse)
Reverses the sort order |
| ▶ ls -t (time)
Lists the most recent files first | ▶ ls -ltr (options can be combined)
Long listing, most recent files at the end |

18

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



File name pattern substitutions

Better introduced by examples!

► `ls *txt`

The shell first replaces `*txt` by all the file and directory names ending by `txt` (including `.txt`), except those starting with `.`, and then executes the `ls` command line.

► `ls -d .*`

Lists all the files and directories starting with `.`

`-d` tells `ls` not to display the contents of directories.

► `cat ?.log`

Displays all the files which names start by 1 character and end by `.log`

19

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Special directories (1)

`./`

► The current directory. Useful for commands taking a directory argument. Also sometimes useful to run commands in the current directory (see later).

► So `./readme.txt` and `readme.txt` are equivalent.

`../`

► The parent (enclosing) directory. Always belongs to the `.` directory (see `ls -a`). Only reference to the parent directory.

► Typical usage:

`cd ..`

20

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Special directories (2)

`~/`

► Not a special directory indeed. Shells just substitute it by the home directory of the current user.

► Cannot be used in most programs, as it is not a real directory.

`~/sydney/`

► Similarly, substituted by shells by the home directory of the `sydney` user.

21

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The `cd` and `pwd` commands

► `cd <dir>`

Changes the current directory to `<dir>`.

► `cd -`

Gets back to the previous current directory.

► `pwd`

Displays the current directory ("working directory").

► `Pushd`

Pushes current directory on stack and `cd`'s to new directory.

► `popd`

22

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The `cp` command

► `cp <source_file> <target_file>`

Copies the source file to the target.

► `cp file1 file2 file3 ... dir`

Copies the files to the target directory (last argument).

► `cp -i` (interactive)

Asks for user confirmation if the target file already exists

► `cp -r <source_dir> <target_dir>` (recursive)

Copies the whole directory.

23

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



`mv` and `rm` commands

► `mv <old_name> <new_name>` (move)

Renames the given file or directory.

► `mv -i` (interactive)

If the new file already exists, asks for user confirm

► `rm file1 file2 file3 ...` (remove)

Removes the given files.

► `rm -i` (interactive)

Always ask for user confirm.

► `rm -r dir1 dir2 dir3` (recursive)

Removes the given directories with all their contents.

24

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Creating and removing directories

- ▶ `mkdir dir1 dir2 dir3 ...` (make dir)
Creates directories with the given names.
- ▶ `rmdir dir1 dir2 dir3 ...` (remove dir)
Removes the given directories
Safe: only works when directories are empty.
Alternative: `rm -r` (doesn't need empty directories).

25



Displaying file contents

Several ways of displaying the contents of files.

- ▶ `cat file1 file2 file3 ...` (concatenate)
Concatenates and outputs the contents of the given files.
- ▶ `more file1 file2 file3 ...`
After each page, asks the user to hit a key to continue.
Can also jump to the first occurrence of a keyword (`/` command).
- ▶ `less file1 file2 file3 ...`
Does more than `more` with less.
Doesn't read the whole file before starting.
Supports backward movement in the file (`?` command).

26



The head and tail commands

- ▶ `head [-<n>] <file>`
Displays the first <n> lines (or 10 by default) of the given file.
Doesn't have to open the whole file to do this!
- ▶ `tail [-<n>] <file>`
Displays the last <n> lines (or 10 by default) of the given file.
No need to load the whole file in RAM! Very useful for huge files.
- ▶ `tail -f <file>` (follow)
Displays the last 10 lines of the given file and continues to display new lines when they are appended to the file.
Very useful to follow the changes in a log file, for example.
- ▶ Examples
`head windows_bugs.txt`
`tail -f outlook_vulnerabilities.txt`

27



The grep command

- ▶ `grep <pattern> <files>`
Scans the given files and displays the lines which match the given pattern.
- ▶ `grep error *.log`
Displays all the lines containing `error` in the `*.log` files
- ▶ `grep -i error *.log`
Same, but case insensitive
- ▶ `grep -ri error .`
Same, but recursively in all the files in `.` and its subdirectories
- ▶ `grep -v info *.log`
Outputs all the lines in the files except those containing `info`.

28



The sort command

- ▶ `sort <file>`
Sorts the lines in the given file in character order and outputs them.
- ▶ `sort -r <file>`
Same, but in reverse order.
- ▶ `sort -ru <file>`
`u`: unique. Same, but just outputs identical lines once.
- ▶ More possibilities described later!

29



The sed command

- ▶ `sed` is a Stream Editor
- ▶ It parses text files and implements a programming language to apply transformations on the text.
- ▶ One of the most common usage of `sed` is text replacement, which relies on regular expressions
- ▶ `sed -e 's/abc/def/' testfile` will replace every string "abc" by "def" in the file `testfile` and display the result on the standard output.
- ▶ `sed 's/[\t]*//'` `testfile` will remove any tabulation or space at the beginning of a line
- ▶ `sed 's/^\([^\|]*\)\|([^\|]*)$/\1 -> \2/' testfile` replace lines like
|string1|string2|
by string1 -> string2

30



sed : regular expressions

- ▶ Regular expressions are useful in many Unix tools, not only sed. They allow to match the input text against an expression.
- ▶ `.` matches any character
- ▶ `[]` matches any character listed inside the brackets
- ▶ `[]^` matches any character not listed inside the brackets
- ▶ `^` matches the beginning of the line
- ▶ `$` matches the end of the line
- ▶ `*` matches the previous element zero or more times, `+` matches the previous element one or more times, `?` matches the previous element zero or one time
- ▶ `\(\)` defines a sub-expression that can be later recalled by using `\n`, where `n` is the number of the sub-expression in the regular expression

▶ More at <http://www.regular-expressions.info/>

31

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Symbolic links

A symbolic link is a special file which is just a reference to the name of another one (file or directory):

- ▶ Useful to reduce disk usage and complexity when 2 files have the same content.
- ▶ Example:
`anakin_skywalker_biography -> darth_vador_biography`
- ▶ How to identify symbolic links:
- ▶ `ls -l` displays `->` and the linked file name.
- ▶ GNU `ls` displays links with a different color.

32

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Creating symbolic links

- ▶ To create a symbolic link (same order as in `cp`):
`ln -s file_name link_name`
 - ▶ To create a link with to a file in another directory, with the same name:
`ln -s ../README.txt`
 - ▶ To create multiple links at once in a given directory:
`ln -s file1 file2 file3 ... dir`
 - ▶ To remove a link:
`rm link_name`
- Of course, this doesn't remove the linked file!

33

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Hard links

- ▶ The default behavior for `ln` is to create *hard links*
- ▶ A *hard link* to a file is a regular file with exactly the same physical contents
- ▶ While they still save space, hard links can't be distinguished from the original files.
- ▶ If you remove the original file, there is no impact on the hard link contents.
- ▶ The contents are removed when there are no more files (hard links) to them.

34

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>




Files names and inodes

Makes hard and symbolic (soft) links easier to understand!

Users

File name interface



Filesystem

35

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The Unix and GNU / Linux command line

Command documentation

36

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Command help

Some Unix commands and most GNU / Linux commands offer at least one help argument:

► `-h`

(`-` is mostly used to introduce 1-character options)

► `--help`

(`--` is always used to introduce the corresponding “long” option name, which makes scripts easier to understand)

You also often get a short summary of options when you input an invalid argument.

37



Manual pages

`man <keyword>`

Displays one or several manual pages for `<keyword>`

► `man man`

Most available manual pages are about Unix commands, but some are also about C functions, headers or data structures, or even about system configuration files!

► `man stdio.h`

► `man fstab` (for `/etc/fstab`)

Manual page files are looked for in the directories specified by the `MANPATH` environment variable.

38



Info pages

► In GNU, man pages are being replaced by info pages. Some manual pages even tell to refer to info pages instead.

`info <command>`

► info features:

► Documentation structured in sections (“nodes”) and subsections (“subnodes”)

► Possibility to navigate in this structure: top, next, prev, up

► Info pages generated from the same texinfo source as the HTML documentation pages

39



Searching the Internet for resources (2)

Looking for documentation

► Look for `<tool>` or `<tool>` page to find the tool or project home page and then find the latest documentation resources.

► Look for `<tool>` documentation or `<tool>` manual in your favorite search engine.

Looking for generic technical information

► Wikipedia: <http://wikipedia.org>

Lots of useful definitions in computer science. A real encyclopedia! Open to anyone's contributions.

40



Searching the Internet for resources (1)

Investigating issues

► Most forums and mailing list archives are public, and are indexed on a very frequent basis by [Google](http://www.google.com).

► If you investigate an error message, copy it verbatim in the search form, enclosed in double quotes (“error message”). Lots of chances that somebody else already faced the same issue.

► Don't forget to use Google Groups:

<http://groups.google.com/>

This site indexes more than 20 years of newsgroups messages.

41



The Unix and GNU / Linux command line

Users and permissions

42



File access rights

Use `ls -l` to check file access rights

3 types of access rights

- ▶ Read access (`r`)
- ▶ Write access (`w`)
- ▶ Execute rights (`x`)

3 types of access levels

- ▶ User (`u`): for the owner of the file
- ▶ Group (`g`): each file also has a "group" attribute, corresponding to a given list of users
- ▶ Others (`o`): for all other users

43



Access right constraints

▶ `x` is sufficient to execute binaries

Both `x` and `r` are required for shell scripts.

▶ Both `r` and `x` permissions needed in practice for directories: `r` to list the contents, `x` to access the contents.

▶ You can't rename, remove, copy files in a directory if you don't have `w` access to this directory.

▶ If you have `w` access to a directory, you CAN remove a file even if you don't have write access to this file (remember that a directory is just a file describing a list of files). This even lets you modify (remove + recreate) a file even without `w` access to it.

44



Access rights examples

▶ `-rw-r--r--`

Readable and writable for file owner, only readable for others

▶ `-rw-r-----`

Readable and writable for file owner, only readable for users belonging to the file group.

▶ `drwx-----`

Directory only accessible by its owner

▶ `-----r-x`

File executable by others but neither by your friends nor by yourself. Nice protections for a trap...



45



chmod: changing permissions

▶ `chmod <permissions> <files>`

2 formats for permissions:

▶ Octal format (`abc`):

`a,b,c = r*4+w*2+x` (`r, w, x`: booleans)

Example: `chmod 644 <file>`

(`rw` for `u`, `r` for `g` and `o`)

▶ Or symbolic format. Easy to understand by examples:

`chmod go+r`: add read permissions to group and others.

`chmod u-w`: remove write permissions from user.

`chmod a-x`: (`a`: all) remove execute permission from all.

46



More chmod (1)

`chmod -R a+rX linux/`

Makes `linux` and everything in it available to everyone!

▶ `R`: apply changes recursively

▶ `X`: `x`, but only for directories and files already executable

Very useful to open recursive access to directories, without adding execution rights to all files.

47



More chmod (2)

`chmod a+t /tmp`

▶ `t` (sticky). Special permission for directories, allowing only the directory and file owner to delete a file in a directory.

▶ Useful for directories with write access to anyone, like `/tmp`.

▶ Displayed by `ls -l` with a `t` character.

48



File ownership

Particularly useful in (embedded) system development when you create files for another system.

▶ `chown -R sco /home/linux/src` (-R: recursive)
Makes user `sco` the new owner of all the files in `/home/linux/src`.

▶ `chgrp -R empire /home/askywalker`
Makes `empire` the new group of everything in `/home/askywalker`.

▶ `chown -R borg:aliens usss_entreprise/`
`chown` can be used to change the owner and group at the same time.

49

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Beware of the dark side of root

▶ root user privileges are only needed for very specific tasks with security risks: mounting, creating device files, loading drivers, starting networking, changing file ownership, package upgrades...

▶ Even if you have the `root` password, your regular account should be sufficient for 99.9 % of your tasks (unless you are a system administrator).

▶ In a training session, it is acceptable to use `root`. In real life, you may not even have access to this account, or put your systems and data at risk if you do.



50

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Using the root account

In case you really want to use `root`...

▶ If you have the `root` password:
`su -` (switch user)

▶ In modern distributions, the `sudo` command gives you access to some `root` privileges with your own user password.
Example: `sudo mount /dev/hda4 /home`

51

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The Unix and GNU / Linux command line

Standard I/O, redirections, pipes

52

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Standard output

More about command output

▶ All the commands outputting text on your terminal do it by writing to their *standard output*.

▶ Standard output can be written (redirected) to a file using the `>` symbol

▶ Standard output can be appended to an existing file using the `>>` symbol

53

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Standard output redirection examples

▶ `ls ~saddam/* > ~gwb/weapons_mass_destruction.txt`

▶ `cat obiwan_kenobi.txt > starwars_biographies.txt`
`cat han_solo.txt >> starwars_biographies.txt`

▶ `echo "README: No such file or directory" > README`
Useful way of creating a file without a text editor.
Nice Unix joke too in this case. 😊

54

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Standard input

More about command input

- ▶ Lots of commands, when not given input arguments, can take their input from *standard input*.

▶ sort

windows

linux

[Ctrl][D]

linux

windows

sort takes its input from the standard input: in this case, what you type in the terminal (ended by [Ctrl][D])

▶ sort < participants.txt

The standard input of sort is taken from the given file.

55

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Pipes

- ▶ Unix pipes are very useful to redirect the standard output of a command to the standard input of another one.

▶ Examples

▶ cat *.log | grep -i error | sort

▶ grep -ri error . | grep -v "ignored" | sort -u \> serious_errors.log

▶ cat /home/*/*homework.txt | grep mark | more

- ▶ This one of the most powerful features in Unix shells!

56

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The tee command

tee [-a] file

- ▶ The **tee** command can be used to send standard output to the screen and to a file simultaneously.

▶ make | tee build.log

Runs the **make** command and stores its output to **build.log**.

▶ make install | tee -a build.log

Runs the **make install** command and appends its output to **build.log**.

57

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Standard error

- ▶ Error messages are usually output (if the program is well written) to *standard error* instead of standard output.

- ▶ Standard error can be redirected through **2>** or **2>>**

▶ Example:

cat f1 f2 nofile > newfile 2> errfile

- ▶ Note: 1 is the descriptor for standard output, so **1>** is equivalent to **>**.

- ▶ Can redirect both standard output and standard error to the same file using **&>** :

cat f1 f2 nofile &> wholefile

58

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The yes command

Useful to fill standard input with always the same string.

▶ yes <string> | <command>

Keeps filling the standard input of **<command>** with **<string>** (y by default).

▶ Examples

yes | rm -r dir/

bank> yes no | credit_applicant

yes "" | make oldconfig

(equivalent to hitting [Enter] to accept all default settings)

59

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Special devices (1)

Device files with a special behavior or contents

▶ /dev/null

The data sink! Discards all data written to this file.

Useful to get rid of unwanted output, typically log information: **mplayer black_adder_4th.avi &> /dev/null**

▶ /dev/zero

Reads from this file always return \0 characters

Useful to create a file filled with zeros:

dd if=/dev/zero of=disk.img bs=1k count=2048

See **man null** or **man zero** for details

60

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Special devices (2)

▶ /dev/random

Returns random bytes when read. Mainly used by cryptographic programs. Uses interrupts from some device drivers as sources of true randomness ("entropy"). Reads can be blocked until enough entropy is gathered.

▶ /dev/urandom

For programs for which pseudo random numbers are fine. Always generates random bytes, even if not enough entropy is available (in which case it is possible, though still difficult, to predict future byte sequences from past ones).

See [man random](#) for details.

61

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Special devices (3)

▶ /dev/full

Mimics a full device. Useful to check that your application properly handles this kind of situation.

See [man full](#) for details.

62

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The Unix and GNU / Linux command line

Task control



Full control on tasks

- ▶ Since the beginning, Unix supports true preemptive multitasking.
- ▶ Ability to run many tasks in parallel, and abort them even if they corrupt their own state and data.
- ▶ Ability to choose which programs you run.
- ▶ Ability to choose which input your programs takes, and where their output goes.

63

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

64

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Processes

"Everything in Unix is a file
Everything in Unix that is not a file is a process"

Processes

- ▶ Instances of a running programs
- ▶ Several instances of the same program can run at the same time
- ▶ Data associated to processes:
Open files, allocated memory, stack, process id, parent, priority, state...

65

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Running jobs in background

Same usage throughout all the shells

- ▶ Useful
- ▶ For command line jobs which output can be examined later, especially for time consuming ones.
- ▶ To start graphical applications from the command line and then continue with the mouse.
- ▶ Starting a task: add **&** at the end of your line:
`find_prince_charming --cute --clever --rich &`

66

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Background job control

▶ jobs

Returns the list of background jobs from the same shell

```
[1]-  Running ~/bin/find_meaning_of_life --without-god &
[2]+  Running make mistakes &
```

▶ fg

fg %<n>

Puts the last / nth background job in foreground mode

▶ Moving the current task in background mode:

[Ctrl] Z

bg

▶ kill %<n>

Aborts the nth job.

67

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Job control example

```
> jobs
```

```
[1]-  Running ~/bin/find_meaning_of_life --without-god &
[2]+  Running make mistakes &
```

```
> fg
```

```
make mistakes
```

```
> [Ctrl] Z
```

```
[2]+  Stopped make mistakes
```

```
> bg
```

```
[2]+  make mistakes &
```

```
> kill %1
```

```
[1]+  Terminated ~/bin/find_meaning_of_life --without-god
```

68

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Listing all processes

... whatever shell, script or process they are started from

▶ ps -aux

Lists all the processes running on the system

```
ps -aux | grep bart | grep bash
USER  PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
bart   3039  0.0  0.2 5916 1380 pts/2    S   14:35   0:00 /bin/bash
bart   3134  0.0  0.2 5388 1380 pts/3    S   14:36   0:00 /bin/bash
bart   3190  0.0  0.2 6368 1360 pts/4    S   14:37   0:00 /bin/bash
bart   3416  0.0  0.0    0 0 pts/2    RW   15:07   0:00 [bash]
```

▶ PID: Process id
VSZ: Virtual process size (code + data + stack)
RSS: Process resident size: number of KB currently in RAM
TTY: Terminal
STAT: Status: R (Runnable), S (Sleep), W (paging), Z (Zombie)...

69

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Live process activity

▶ top - Displays most important processes, sorted by cpu percentage (**htop**)

```
top - 15:44:33 up 1:11, 5 users, load average: 0.98, 0.61, 0.59
Tasks: 81 total, 5 running, 76 sleeping, 0 stopped, 0 zombie
Cpu(s): 92.7% us, 5.3% sy, 0.0% ni, 0.0% id, 1.7% wa, 0.3% hi, 0.0% si
Mem: 515344k total, 512384k used, 2960k free, 20464k buffers
Swap: 1044184k total, 0k used, 1044184k free, 277660k cached
```

```
PID USER  PR  NI  VIRT  RES  SHR  S %CPU %MEM    TIME+  COMMAND
3809 jdoe   25   0 6256 3932 1312 R 93.8  0.8   0:21.49 bunzip2
2769 root    16   0 157m 80m 90m R  2.7 16.0   5:21.01 X
3006 jdoe   15   0 30928 15m 27m S  0.3  3.0   0:22.40 kdeinit
3008 jdoe   16   0 5624 892 4468 S  0.3  0.2   0:06.59 autorun
3034 jdoe   15   0 26764 12m 24m S  0.3  2.5   0:12.68 kscd
3810 jdoe   16   0 2892 916 1620 R  0.3  0.2   0:00.06 top
```

▶ You can change the sorting order by typing
M: Memory usage, P: %CPU, T: Time.

▶ You can kill a task by typing **k** and the process id.

70

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Killing processes (1)

▶ kill <pids>

Sends an abort signal to the given processes. Lets processes save data and exit by themselves. Should be used first.

Example:

```
kill 3039 3134 3190 3416
```

▶ kill -9 <pids>

Sends an immediate termination signal. The system itself terminates the processes. Useful when a process is really stuck (doesn't answer to kill -1).

▶ kill -9 -1

Kills all the processes of the current user. -1: means all processes.

71

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Killing processes (2)

▶ killall [-<signal>] <command>

Kills all the jobs running <command>. Example:
killall bash

▶ xkill

Lets you kill a graphical application by clicking on it!
Very quick! Convenient when you don't know the application command name.

72

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Recovering from stuck graphics

- ▶ If your graphical session is stuck and you can no longer type in your terminals, don't reboot!
- ▶ It is very likely that your system is still fine. Try to access a text console by pressing the [Ctrl][Alt][F1] keys (or [F2],[F3] for more text consoles)
- ▶ In the text console, you can try to kill the guilty application.
- ▶ Once this is done, you can go back to the graphic session by pressing [Ctrl][Alt][F5] or [Ctrl][Alt][F7] (depending on your distribution)
- ▶ If you can't identify the stuck program, you can also kill all your processes: `kill -9 -1`
You are then brought back to the login screen.

73



Sequential commands

- ▶ Can type the next command in your terminal even when the current one is not over.
 - ▶ Can separate commands with the ; symbol:
`echo "I love thee"; sleep 10; echo " not"`
 - ▶ Conditionals: use || (or) or && (and):
`more God || echo "Sorry, God doesn't exist"`
Runs `echo` only if the first command fails
- `ls ~sd6 && cat ~sd6/* > ~sydney/recipes.txt`
Only cats the directory contents if the `ls` command succeeds (means read access).

74



Quoting (1)

Double (") quotes can be used to prevent the shell from interpreting spaces as argument separators, as well as to prevent file name pattern expansion.

```
> echo "Hello World"
Hello World

> echo "You are logged as $USER"
You are logged as bgates

> echo *.log
find_prince_charming.log cosmetic_buys.log

> echo "*.log"
*.log
```

75



Quoting (2)

Single quotes bring a similar functionality, but what is between quotes is never substituted

```
> echo 'You are logged as $USER'
You are logged as $USER
```

Back quotes (`) can be used to call a command within another

```
> cd /lib/modules/`uname -r`; pwd
/lib/modules/2.6.9-1.6_FC2
```

Back quotes can be used within double quotes

```
> echo "You are using Linux `uname -r`"
You are using Linux 2.6.9-1.6_FC2
```

76



Measuring elapsed time

```
▶ time find_expensive_housing --near
<...command output...>
real 0m2.304s (actual elapsed time)
user 0m0.449s (CPU time running program code)
sys 0m0.106s (CPU time running system calls)
```

`real = user + sys + waiting`
`waiting = I/O waiting time + idle time (running other tasks)`

77



Environment variables

- ▶ Shells let the user define *variables*. They can be reused in shell commands. Convention: lower case names
- ▶ You can also define *environment variables*: variables that are also visible within scripts or executables called from the shell. Convention: upper case names.
- ▶ `env`
Lists all defined environment variables and their value.

78



Shell variables examples

Shell variables (bash)

```
▶ projdir=/home/marshall/coolstuff
ls -la $projdir; cd $projdir
```

Environment variables (bash)

```
▶ cd $HOME
▶ export DEBUG=1
./find_extraterrestrial_life
(displays debug information if DEBUG is set)
```

79



Main standard environment variables

Used by lots of applications!

```
▶ LD_LIBRARY_PATH
Shared library search path
▶ DISPLAY
Screen id to display X
(graphical) applications on.
▶ EDITOR
Default editor (vi, emacs...)
▶ HOME
Current user home directory
▶ HOSTNAME
Name of the local machine
```

```
▶ MANPATH
Manual page search path
▶ PATH
Command search path
▶ PRINTER
Default printer name
▶ SHELL
Current shell name
▶ TERM
Current terminal type
▶ USER
Current user name
```

80



PATH environment variables

▶ PATH

Specifies the shell search order for commands

```
/home/acox/bin:/usr/local/bin:/usr/kerberos/bin:/usr/bin:/bin:/usr/
X11R6/bin:/bin:/usr/bin
```

▶ LD_LIBRARY_PATH

Specifies the shared library (binary code libraries shared by applications, like the C library) search order for ld

```
/usr/local/lib:/usr/lib:/lib:/usr/X11R6/lib
```

▶ MANPATH

Specifies the search order for manual pages

```
/usr/local/man:/usr/share/man
```

81



PATH usage warning

It is strongly recommended not to have the "." directory in your PATH environment variable, in particular not at the beginning:

▶ A cracker could place a malicious `ls` file in your directories. It would get executed when you run `ls` in this directory and could do naughty things to your data.

▶ If you have an executable file called `test` in a directory, this will override the default test program and some scripts will stop working properly.

▶ Each time you `cd` to a new directory, the shell will waste time updating its list of available commands.

Call your local commands as follows: `./test`

82



Alias

Shells let you define command *aliases*: shortcuts for commands you use very frequently.

Examples

```
▶ alias ls='ls -la'
Useful to always run commands with default arguments.
▶ alias rm='rm -i'
Useful to make rm always ask for confirmation.
▶ alias frd='find_rambaldi_device --asap --risky'
Useful to replace very long and frequent commands.
▶ alias cia='. /home/sydney/env/cia.sh'
Useful to set an environment in a quick way
(. is a shell command to execute the content of a shell script).
```

83



The which command

Before you run a command, *which* tells you where it is found

```
▶ bash> which ls
alias ls='ls --color=tty'
/bin/ls
▶ tcsh> which ls
ls: aliased to ls --color=tty
▶ bash> which alias
/usr/bin/which: no alias in
(/usr/local/bin:/usr/bin:/bin:/usr/X11R6/bin)
▶ tcsh> which alias
alias: shell built-in command.
```

84



~/.bashrc file

▶ ~/.bashrc

Shell script read each time a `bash` shell is started

- ▶ You can use this file to define
- ▶ Your default environment variables (`PATH`, `EDITOR`...).
- ▶ Your aliases.
- ▶ Your prompt (see the `bash` manual for details).
- ▶ A greeting message.

85

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Command editing

- ▶ You can use the left and right arrow keys to move the cursor in the current command.
- ▶ You can use `[Ctrl][a]` to go to the beginning of the line, and `[Ctrl][e]` to go to the end.
- ▶ You can use the up and down arrows to select earlier commands.
- ▶ You can use `[Ctrl][r]` to search inside the history of previous commands.

86

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Command history (1)

▶ history

Displays the latest commands that you ran and their number.
You can copy and paste command strings.

- ▶ You can recall the latest command:

!!

- ▶ You can recall a command by its number

!1003

- ▶ You can recall the latest command matching a starting string:

!cat

87

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Command history (2)

- ▶ You can make substitutions on the latest command:

`^more^less`

- ▶ You can run another command with the same arguments:

`more !*`

88

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The Unix and GNU / Linux command line

Miscellaneous Text editors



Text editors

Graphical text editors
Fine for most needs

- ▶ `nedit`
- ▶ `Emacs`, `Xemacs`
- ▶ `Kate`, `Gedit`

Text-only text editors
Often needed for sysadmins and great for power users

- ▶ `vi`, `vim`
- ▶ `nano`

89

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

90

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The nedit text editor

<http://www.nedit.org/>

- Best text editor for non vi or emacs experts
- Feature highlights:
 - Very easy text selection and moving
 - Syntax highlighting for most languages and formats. Can be tailored for your own log files, to highlight particular errors and warnings.
 - Easy to customize through menus
 - Not installed by default by all distributions

91

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



nedit screenshot

92

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Emacs / Xemacs

- Emacs and Xemacs are pretty similar (up to your preference)
- Extremely powerful text editor features
- Great for power users
- Less ergonomic than nedit
- Non standard shortcuts
- Much more than a text editor (games, e-mail, shell, browser).
- Some power commands have to be learnt.

93

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Kate and gedit

- Kate is a powerful text editor dedicated to programming activities, for KDE

<http://kate.kde.org>

- Gedit is a text editor for the Gnome environment

<http://projects.gnome.org/gedit/>

94

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



vi

Text-mode text editor available in all Unix systems. Created before computers with mice appeared.

- Difficult to learn for beginners used to graphical text editors.
- Very productive for power users.
- Often can't be replaced to edit files in system administration or in Embedded Systems, when you just have a text console.

95

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



vim - vi improved

- vi implementation now found in most GNU / Linux host systems
- Implements lots of features available in modern editors: syntax highlighting, command history, help, unlimited undo and much much more.
- Cool feature example: can directly open compressed text files.
- Comes with a GTK graphical interface (gvim)
- Unfortunately, not free software (because of a small restriction in freedom to make changes)

96

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



vi basic commands



Though vi is extremely powerful, its main 30 commands are easy to learn and are sufficient for 99% of everyone's needs!

You can also take the quick tutorial by running `vimtutor`.

Get our vi memento sheet if you didn't get it with this course:

<http://free-electrons.com/docs/command-line>

97

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



GNU nano

<http://www.nano-editor.org/>

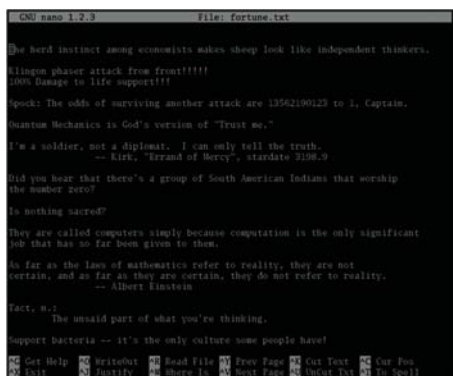
- ▶ Another small text-only, mouse free text editor.
- ▶ An enhanced Pico clone (non free editor in Pine)
- ▶ Friendly and easier to learn for beginners thanks to on screen command summaries.
- ▶ Available in binary packages for several platforms.
- ▶ An alternative to vi in embedded systems. However, not available as a busybox built-in.

98

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



GNU nano screenshot



99

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The Unix and GNU / Linux command line

Miscellaneous
Compressing and archiving

100

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Measuring disk usage

Caution: different from file size!

▶ `du -h <file>` (disk usage)

-h: returns size on disk of the given file, in human readable format: K (kilobytes), M (megabytes) or G (gigabytes). Without -h, du returns the raw number of disk blocks used by the file (hard to read).

Note that the -h option only exists in GNU du.

▶ `du -sh <dir>`

-s: returns the sum of disk usage of all the files in the given directory.

101

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Measuring disk space

▶ `df -h <dir>`

Returns disk usage and free space for the filesystem containing the given directory. Similarly, the -h option only exists in GNU df.

▶ Example:

```
> df -h
Filesystem      Size  Used Avail Use% Mounted on
/dev/hda5       9.2G  7.1G  1.8G  81% /
```

▶ `df -h`

Returns disk space information for all filesystems available in the system. When errors happen, useful to look for full filesystems.

102

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Compressing and decompressing

Very useful for shrinking huge files and saving space

▶ `g[un]zip <file>`

GNU zip compression utility. Creates `.gz` files.
Ordinary performance (similar to Zip).

▶ `b[un]zip2 <file>`

More recent and effective compression utility.
Creates `.bz2` files. Usually 20-25% better than `gzip`.

▶ `[un]lzma <file>`

Much better compression ratio than `bzip2` (up to 10 to 20%).
Compatible command line options.

103

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Archiving (1)

Useful to backup or release a set of files within 1 file

▶ `tar`: originally "tape archive"

▶ Creating an archive:

`tar cvf <archive> <files or directories>`

`c`: create

`v`: verbose. Useful to follow archiving progress.

`f`: file. Archive created in file (tape used otherwise).

▶ Example:

`tar cvf /backup/home.tar /home`

`bzip2 /backup/home.tar`

104

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Archiving (2)

▶ Viewing the contents of an archive or integrity check:

`tar tvf <archive>`

`t`: test

▶ Extracting all the files from an archive:

`tar xvf <archive>`

▶ Extracting just a few files from an archive:

`tar xvf <archive> <files or directories>`

Files or directories are given with paths relative to the archive root directory.

105

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Extra options in GNU tar

`tar = gtar` = GNU `tar` on GNU / Linux

Can compress and uncompress archives on the fly.

Useful to avoid creating huge intermediate files

Much simpler to do than with `tar` and `bzip2`!

▶ `j` option: [un]compresses on the fly with `bzip2`

▶ `z` option: [un]compresses on the fly with `gzip`

▶ `--lzma` option: [un]compresses on the fly with `lzma`

▶ Examples (which one will you remember?)



`gtar jcvf bills_bugs.tar.bz2 bills_bugs`

`tar cvf - bills_bugs | bzip2 > bills_bugs.tar.bz2`

106

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Checking file integrity

Very low cost solution to check file integrity

▶ `md5sum FC3-i386-disk*.iso > MD5SUM`

Computes a MD5 (Message Digest Algorithm 5) 128 bit checksum of the given files. Usually redirected to a file.

▶ Example output:

```
db8c7254beeb4f6b891d1ed3f689b412 FC3-i386-disc1.iso
2c11674cf429fe570445afd9d5ff564e FC3-i386-disc2.iso
f88f6ab5947ca41f3cf31db04487279b FC3-i386-disc3.iso
6331c00aa3e8c088cc365eeb7ef230ea FC3-i386-disc4.iso
```

▶ `md5sum -c MD5SUM`

Checks the integrity of the files in `MD5SUM` by comparing their actual MD5 checksum with their original one.

107

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



The Unix and GNU / Linux command line

Miscellaneous Printing

108

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Unix printing

- ▶ Multi-user, multi-job, multi-client, multi-printer
In Unix / Linux, printing commands don't really print. They send jobs to printing queues, possibly on the local machine, on network printing servers or on network printers.
- ▶ Printer independent system:
Print servers only accept jobs in PostScript or text. Printer drivers on the server take care of the conversion to each printers own format.
- ▶ Robust system:
Reboot a system, it will continue to print pending jobs.



109

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Printing commands

- ▶ Useful environment variable: **PRINTER**
Sets the default printer on the system. Example:
`export PRINTER=lp`
- ▶ `lpr [-P<queue>] <files>`
Sends the given files to the specified printing queue
The files must be in text or PostScript format. Otherwise, you only print garbage.
- ▶ `a2ps [-P<queue>] <files>`
“Any to PostScript” converts many formats to PostScript and send the output to the specified queue. Useful features: several pages / sheet, page numbering, info frame...

110

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



Print job control

- ▶ `lpq [-P<queue>]`
Lists all the print jobs in the given or default queue.
- ```
lp is not ready
Rank Owner Job File(s) Total Size
1st asloane 84 nsa_windows_backdoors.ps 60416 bytes
2nd amoore 85 gw_bush_iraq_mistakes.ps 65024000 bytes
```
- ▶ `cancel <job#> [<queue>]`  
Removes the given job number from the default queue.

111

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## Using PostScript and PDF files

- Viewing a PostScript file
- ▶ PostScript viewers exist, but their quality is pretty poor.
- ▶ Better convert to PDF with `ps2pdf`:  
`ps2pdf decss_algorithm.ps`  
`xpdf decss_algorithm.pdf &`
- Printing a PDF file
- ▶ You don't need to open a PDF reader!
- ▶ Better convert to PostScript with `pdf2ps`:  
`pdf2ps rambaldi_artifacts_for_dummies.pdf`  
`lpr rambaldi_artifacts_for_dummies.ps`

112

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## The Unix and GNU / Linux command line

### Miscellaneous Synchronizing files



## Smart directory copy with rsync

- `rsync` (remote sync) has been designed to keep in sync directories on 2 machines with a low bandwidth connection.
- ▶ Only copies files that have changed. Files with the same size are compared by checksums.
- ▶ Only transfers the blocks that differ within a file!
- ▶ Can compress the transferred blocks
- ▶ Preserves symbolic links and file permissions: also very useful for copies on the same machine.
- ▶ Can work through ssh (secure remote shell). Very useful to update the contents of a website, for example.

113

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

114

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## rsync examples (1)

- ▶ `rsync -a /home/arvin/sd6_agents/ /home/sydney/misc/`  
-a: archive mode. Equivalent to `-rlptgoD...` easy way to tell you want recursion and want to preserve almost everything.
  - ▶ `rsync -Pav --delete /home/steve/ideas/ /home/bill/my_ideas/`  
-P: `--partial` (keep partially transferred files) and `--progress` (show progress during transfer)  
`--delete`: delete files in the target which don't exist in the source.
- Caution: directory names should end with `/`. Otherwise, you get a `my_ideas/ideas/` directory at the destination.

115

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

## rsync examples (2)

- ▶ Copying to a remote machine  
`rsync -Pav /home/bill/legal/arguments/ \`  
`bill@www.sco.com:/home/legal/arguments/`  
User `bill` will be prompted for a password.
- ▶ Copying from a remote machine through ssh  
`rsync -Pav -e ssh homer@tank.duff.com:/prod/beer/ \`  
`fridge/homer/beer/`  
User `homer` will be prompted for his ssh key password.

116

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

## The Unix and GNU / Linux command line

### Miscellaneous Comparing files and directories

117

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

## Comparing files and directories

- ▶ `diff file1 file2`  
Reports the differences between 2 files, or nothing if the files are identical.
- ▶ `diff -r dir1/ dir2/`  
Reports all the differences between files with the same name in the 2 directories.
- ▶ These differences can be saved in a file using the redirection, and then later re-applied using the `patch` command.
- ▶ To investigate differences in detail, better use graphical tools!

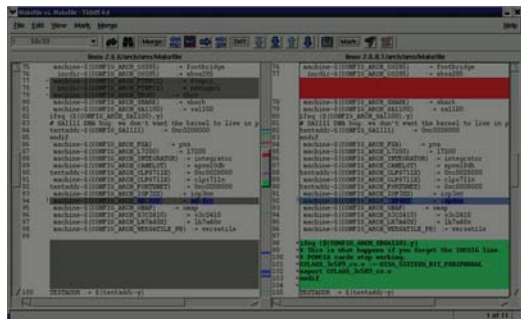
118

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

## tkdiff

<http://tkdiff.sourceforge.net/>

Useful tool to compare files and merge differences

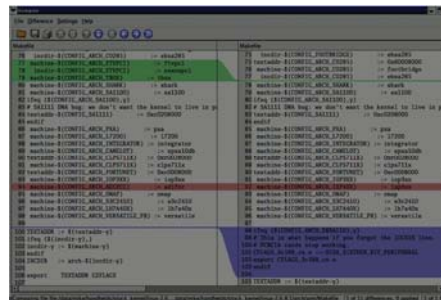


119

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

## kompare

Another nice tool to compare files and merge differences  
Part of the `kdesdk` package (Fedora Core)



120

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>






## gvimdiff

Another nice tool to view differences in files

Available in most distributions with **gvim**  
Apparently not using **diff**.

No issue with files with binary sections!



121

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## The Unix and GNU / Linux command line

Miscellaneous  
Looking for files

122

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## The find command

Better explained by a few examples!

► **find . -name "\*.pdf"**

Lists all the \*.pdf files in the current (.) directory or subdirectories. You need the double quotes to prevent the shell from expanding the \* character.

► **find docs -name "\*.pdf" -exec xpdf {} \;**

Finds all the \*.pdf files in the docs directory and displays one after the other.

► Many more possibilities available! However, the above 2 examples cover most needs.

123

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## The locate command

Much faster regular expression search alternative to **find**

► **locate keys**

Lists all the files on your system with keys in their name.

► **locate "\*.pdf"**

Lists all the \*.pdf files available on the whole machine

► **locate "/home/fridge/\*beer\*"**

Lists all the \*beer\* files in the given directory (absolute path)

► **locate** is much faster because it indexes all files in a dedicated database, which is updated on a regular basis.

► **find** is better to search through recently created files.

124

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## The Unix and GNU / Linux command line

Miscellaneous  
Various commands



## Getting information about users

► **who**

Lists all the users logged on the system.

► **whoami**

Tells what user I am logged as.

► **groups**

Tells which groups I belong to.

► **groups <user>**

Tells which groups <user> belongs to.

► **finger <user>**

Tells more details (real name, etc) about <user>

Disabled in some systems (security reasons).

125

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>

126

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## Changing users

You do not have to log out to log on another user account!

### ► `su hyde`

(Rare) Change to the `hyde` account, but keeping the environment variable settings of the original user.

### ► `su - jekyll`

(More frequent) Log on the `jekyll` account, with exactly the same settings as this new user.

### ► `su -`

When no argument is given, it means the `root` user.

127

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## The wget command

Instead of downloading files from your browser, just copy and paste their URL and download them with `wget`!

### `wget` main features

#### ► http and ftp support

#### ► Can resume interrupted downloads

#### ► Can download entire sites or at least check for bad links

#### ► Very useful in scripts or when no graphics are available (system administration, embedded systems)

#### ► Proxy support (`http_proxy` and `ftp_proxy` env. variables)

128

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## wget examples

### ► `wget -c \ http://microsoft.com/customers/dogs/winxp4dogs.zip`

Continues an interrupted download.

### ► `wget -m http://lwn.net/`

Mirrors a site.

### ► `wget -r -np http://www.xml.com/ldd/chapter/book/`

Recursively downloads an on-line book for off-line access.

`-np`: "no-parent". Only follows links in the current directory.

129

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## Misc commands (1)

### ► `sleep 60`

Waits for 60 seconds

(doesn't consume system resources).

### ► `wc report.txt` (word count)

438 2115 18302 `report.txt`

Counts the number of lines, words and characters in a file or in standard input.

130

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## Misc commands (2)

### ► `bc` ("basic calculator?")

`bc` is a handy but full-featured calculator. Even includes a programming language! Use the `-l` option to have floating point support.

### ► `date`

Returns the current date. Useful in scripts to record when commands started or completed.

131

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## Checksum commands

► A checksum or hash sum is a fixed-size datum computed from an arbitrary block of digital data for the purpose of detecting accidental errors that may have been introduced during its transmissions or storage.

<http://en.wikipedia.org/wiki/Checksum>

### ► The MD5 hash algorithm is implemented in the `md5sum` command

\$ `md5sum patch-2.6.24.7.bz2`

0c1c5d6d8cd82e18d62406d2f34d1d38 patch-2.6.24.7.bz2

### ► The SHA algorithm is implemented in the `shaXsum` (`sha1sum`, `sha256sum`, etc.)

### ► The integrity of several files can be verified against a file listing the checksums using the `-c` option.

132

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## System administration

See our presentation about system administration basics:

- ▶ Network setup
- ▶ Creating and mounting filesystems
- ▶ Accessing administrator (root) privileges
- ▶ Package management

Also available on <http://free-electrons.com/docs/command-line>

133

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## The Unix and GNU / Linux command line

### Application development

134

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## Compiling simple applications

- ▶ The compiler used for all Linux systems is GCC  
<http://gcc.gnu.org>
- ▶ To compile a single-file application, developed in C :  
`gcc -o test test.c`
- ▶ Will generate a test binary, from the test.c source file
- ▶ For C++ :  
`g++ -o test test.cc`
- ▶ The -Wall option enables more warnings
- ▶ To compile sources files to object files and link the application :  
`gcc -c test1.c`  
`gcc -c test2.c`  
`gcc -o test test1.o test2.o`
- ▶ gcc automatically calls the linker ld

135

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## Using libraries (1)

- ▶ On any Linux system, a C library is available and offers a large set of APIs for application development.
- ▶ See <http://www.gnu.org/software/libc/manual/>
- ▶ Outside of the C library, thousands of other libraries are available for graphic programming, multimedia, networking, scientific computations, and more.
- ▶ Most libraries are already available as packages in your distribution, in general in two packages
- ▶ **libfoo** is the package containing the library itself. This package is required to **execute** already compiled applications, but not sufficient to build new applications
- ▶ **libfoo-dev** is the package containing the headers and other configurations files and tools needed to **build** new applications relying on libfoo.

136

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## Using libraries (2)

- ▶ In your source code, include the proper header files of the library
- ▶ Usually `#include <foo.h>` or `#include <foo/foo.h>`
- ▶ These headers are present in `/usr/include/`
- ▶ Refer to the documentation of the library for details, available in `/usr/share/doc/<package>/`, on the Web... or in the header files !
- ▶ To compile your application with the library, the easiest solution is to use `pkg-config`, which is supported by most libraries today :  
`gcc -o test test.c $(pkg-config --cflags --libs)`
- ▶ By default, the application are dynamically linked with the libraries
- ▶ The libraries must be present in `/lib/` for the application to work
- ▶ Use the `ldd` command to see which libraries are needed by an application

137

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## Make and Makefiles

- ▶ The compilation process can be automated using the make tool.
- ▶ make reads a file called Makefile from the current directory, and executes the rules described in this file
- ▶ Every rule has a target name, a colon, and a list of dependencies, and the list of commands to generate the target from the dependencies  
target: dep1 dep2  
    command1  
    command2  
    command3
- ▶ When simply running make, the default target that is generated is "all". A target is only re-generated if dependencies have changed.
- ▶ See <http://www.gnu.org/software/make/manual/>

138

Free Electrons. Kernel, drivers and embedded Linux development, consulting, training and support. <http://free-electrons.com>



## Simple Makefile example



Variables can be defined and later expanded with \$(VARNAME)

The default target "all" simply depends on the "test" target.

The "test" target depends on t1.o and t2.o. Once these files are generated, the gcc command is executed.  
\$@ is the target name  
\$^ is the name of all dependencies.  
The .o files are generated using implicit dependencies, known by make.

```

CFLAGS = -Wall

all: test

test: t1.o t2.o
 gcc -o $@ $^ $(CFLAGS)

clean:
 $(RM) -f test

install:
 $(CP) test /usr/bin

```

These targets are executed by running "make clean" and "make install"

139

140



## Build systems

- ▶ Makefiles are nice, but they don't easily allow easy adaptation to the different build environment and different build options
- ▶ More elaborated build systems have been developed
- ▶ Autotools (automake, autoconf), based on Makefiles and shell scripts. Even though they are old and a little bit difficult to understand, they are the most popular build system for free software packages.
- ▶ CMake, a newer, cleaner build system
- ▶ Sconcs and Waf, other build systems based on Python
- ▶ The typical steps to compile a autotools based package are  
./configure  
make  
sudo make install

141



## Debugging

- ▶ The official debugger that comes with the GNU development tools is gdb.
- ▶ See <http://sourceware.org/gdb/download/onlinedocs/gdb.html>
- ▶ An application must be compiled with the -g option to be properly debugged. This option adds debugging information to the application binary  
gcc -o test test.c -g
- ▶ The application can then be run inside the gdb debugger :  
gdb test
- ▶ Or the debugger can be attached to the application while it is running :  
gdb test -p PID
- ▶ Where PID is the process ID of the running application

142



## Using gdb

- ▶ **gdb** is a text-based debugger, with a command-line interface like a shell, providing dedicated commands. Some of the important commands are :
- ▶ **break (b)** to set a breakpoint in the code. Can be used with a function name or a location in the source code, or an absolute memory address.
- ▶ **print (p)** to print the value of a variable. Used with a variable name, even if it's a complex one (which involves dereferencing structures, for example)
- ▶ **c** to continue the execution until the next breakpoint.
- ▶ **next (n)** to execute only the next line of code (step over any function call) and **step (s)** to execute only the next line of code (step into any function call)
- ▶ **backtrace (bt)** to display the function call stack

143



## gdb sample session

```

thomas@surf:/tmp$ gcc -o test test.c -g
thomas@surf:/tmp$ gdb test
GNU gdb 6.8-debian
[...]
(gdb) break foo
Breakpoint 1 at 0x80483c7: file test.c, line 5.
(gdb) run
Starting program: /tmp/test2

Breakpoint 1, foo (a=2, b=3) at test.c:5
5 return a + b;
(gdb) p a
$1 = 2
(gdb) p b
$2 = 3
(gdb) c
Continuing.
foo=5

Program exited normally.

```

144



## ddd, the graphical gdb

- ▶ There are several graphical front-ends to gdb. A popular one is **ddd**.
- ▶ Allows to navigate in the source code while debugging
- ▶ Allows to set break points, inspect and change variable value directly by looking at the source code
- ▶ <http://www.gnu.org/software/ddd/>

