

07-3 Using the DSP on the BeagleBoard via c6run

Mark A. Yoder

So how do I use the DSP?

- Skip the DSP altogether
 - Use Arm Neon floating-point
- Bare Metal Approach
 - Get out the hardware manual and learn where the shared memory is...
- Something in between – Use a framework
 - c6run
- Find prepackaged software that uses the DSP
 - Gstreamer (<http://www.gstreamer.net/>)
 - open source multimedia framework

c6run Overview


- Adapted from:
http://processors.wiki.ti.com/index.php/Introduction_to_C6Run
- Open Source Project, hosted on <http://gforge.ti.com>
- Intends to provide an abstracted mechanism for getting code running on the DSP
- Project Goals
 - Introduce DSP Development to ARM/Linux developer
 - Simplify simple application/function offloading to the DSP
 - Start getting Linux and open-source community using the DSP

Project Details

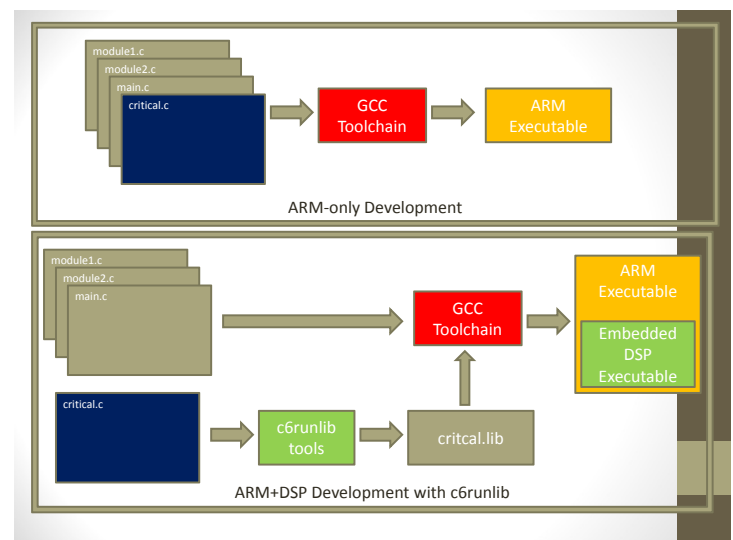
- Currently consists of several components:
 - C6RunApp: Run an entire application on the DSP
 - C6RunLib: Partition an application between the ARM and DSP

C6RunLib

- Goal: automate building an ARM-side library from user's C code of critical functions
- User app can call into library, and calls will be executed on the DSP
- Consists of automating creation of framework, hiding other TI specific bits



5



Example C6RunLib Usage

```
$ c6runlib-cc -c -O2 -o dummy.o dummy.c
```

- Above converts C code containing critical functions to C6000 object file
- Also analyzes global C functions and generates ARM-side remote procedure call stubs


```
$ c6runlib-ar rcs dummy_dsp.lib dummy.o
```

- Add object file to library dummy_dsp.lib
- Underneath, the dummy.o object file is linked to a DSP executable and compiled into the framework
- Framework object file and stubs object file is archived into the lib
- ARM-side stubs resolve symbols for ARM application when built against the library

9/17/2018

C6RunApp

- Entire application retargets to DSP, but with C I/O access to ARM/Linux



8

Example C6RunApp Usage

```
$ c6runapp-cc -o hello_world hello_world.c
```


- Compiles hello_world.c to C6000 object file, which is then linked into a DSP executable
- Executable is compiled into the ARM side framework, which is used to build an ARM-side executable called hello_world

- Notes on the c6runapp-cc cross-compiler tool:
 - Front end script wraps the TI C6000 Codegen tools (specifically cl6x)
 - Supports many GCC options and translates them to appropriate cl6x options
 - Many GCC optimization/warning options are silently ignored
 - Unknown options passed directly to cl6x command-line

9/17/2018

11

Complex FFT Performance



- ▶ FFT runs ~10x faster on DSP than on ARM
- ▶ Small FFT size, overhead dominates, running on DSP does not provide advantage
- ▶ Larger FFT size, overhead absorbed, running on DSP provides advantage

SoC: ARM9 + Floating-Point DSP
 CPU Frequency: 300MHz
 Code & Data Location: External
 DDR2 Memory
 Instruction and Data Cache:
 Enabled
 Single-precision floating-point data
 buffers.

c6run Hands on

- Once compiled, binaries are copied to the Beagle
- Try:

```
beagle$ cd ~/exercises/audioThru/lab06d_audio_c6run
beagle$ git clone git://github.com/MarkAYoder/c6run_build.git
beagle$ cd c6run_build
$ ./unloadmodules.sh
$ ./loadmodules.sh
$ source environment.sh
$ cd examples/c6runapp/hello_world/
$ ./hello_world_arm
Hello SPED world, from C6RunApp!
$ ./hello_world_dsp
Hello SPED world, from C6RunApp!
```

c6run Hands on

- Explore the other examples

```
./c6runapp:
cio_example  emqbit  hello_world
./c6runapp/cio_example:
cio_example.c  cio_example_arm  cio_example_dsp
./c6runapp/emqbit:
bench_arm  bench_dsp  cfft.c  cfft.h  cfft_arm  cfft_dsp
./c6runapp/hello_world:
hello_world.c  hello_world_arm  hello_world_dsp
./c6runlib:
emqbit
./c6runlib/emqbit:
bench_arm  bench_dsp  cfft.c  cfft_arm  cfft_dsp
bench_arm.lib  bench_dsp.lib  cfft.h  cfft_arm.lib  cfft_dsp.lib
```

Details — `cd c6run_target/examples/c6runlib/emqbit`

| # ./cfft_arm | # ./cfft_dsp |
|-------------------------------|-------------------------------|
| N=16,nTimes=100: 0.000458 s | N=16,nTimes=100: 0.015533 s |
| N=32,nTimes=100: 0.00119 s | N=32,nTimes=100: 0.022827 s |
| N=64,nTimes=100: 0.004486 s | N=64,nTimes=100: 0.036712 s |
| N=128,nTimes=100: 0.00708 s | N=128,nTimes=100: 0.07547 s |
| N=256,nTimes=100: 0.018005 s | N=256,nTimes=100: 0.154419 s |
| N=512,nTimes=100: 0.034149 s | N=512,nTimes=100: 0.331909 s |
| N=1024,nTimes=100: 0.076507 s | N=1024,nTimes=100: 0.731232 s |
| N=2048,nTimes=100: 0.1651 s | N=2048,nTimes=100: 1.59704 s |
| N=4096,nTimes=100: 0.364411 s | N=4096,nTimes=100: 3.48624 s |
| N=8192,nTimes=100: 0.79776 s | N=8192,nTimes=100: 7.62561 s |
| N=16384,nTimes=100: 1.72101 s | N=16384,nTimes=100: 16.6948 s |

Details – `main_cfft.c`

```
int main () {
    t1 = get_timestamp();
    secs = (t1 - t0) / 1000000.0L;

    for (N = (1 << MINPOW2), n = 0; N < (1
    << MAXPOW2); N = N << 1, n++)
    {
        complex *in = new_complex_vector(N);
        complex *out = new_complex_vector(N);
        fft_init (N);
        // Copy input data and do one FFT
        memcpy (out, in, (N) * _);
        fft_exec (N, out);
        nTimes = ITERATIONS;
        t0 = get_timestamp();
        for (i = 0; i < nTimes; i++) {
            memcpy (out, in, (N) * _);
            fft_exec (N, out);
        }

        free (in);
        free (out);
        fft_end ();

        fprintf (stderr,
        "N=%d,nTimes=%d: %g s\n", N,
        nTimes, secs);
    }

    return 0;
}
```

`c6run_build/examples/c6runlib/emqbit`

Details – `main_cfft.c`

```
static complex *new_complex_vector(int size) {
    int i;
    complex *new;
    new = (complex *) malloc(sizeof(complex) * size);
    for(i = 0; i < size; ++i)
    {
        new[i].r = (float)rand()/(float)RAND_MAX - 0.5;
        new[i].i = (float)rand()/(float)RAND_MAX - 0.5;
    }
    return new;
}
```

Details – `main_cfft.c`

```
int main () {
    t1 = get_timestamp();
    secs = (t1 - t0) / 1000000.0L;

    for (N = (1 << MINPOW2), n = 0; N < (1
    << MAXPOW2); N = N << 1, n++)
    {
        complex *in = new_complex_vector(N);
        complex *out = new_complex_vector(N);
        fft_init (N);
        // Copy input data and do one FFT
        memcpy (out, in, (N) * _);
        fft_exec (N, out);
        nTimes = ITERATIONS;
        t0 = get_timestamp();
        for (i = 0; i < nTimes; i++) {
            memcpy (out, in, (N) * _);
            fft_exec (N, out);
        }

        free (in);
        free (out);
        fft_end ();

        fprintf (stderr,
        "N=%d,nTimes=%d: %g s\n", N,
        nTimes, secs);
    }

    return 0;
}
```

Where's the DSP?

- **malloc** is overridden to use **cmem**
- **Makefile** decides what is run on the DSP and ARM

Details – Makefile – ARM

```
ARM_TOOLCHAIN_PREFIX ?= arm-none-linux-gnueabi-
ifdef ARM_TOOLCHAIN_PATH
    ARM_CC := $(ARM_TOOLCHAIN_PATH)/bin/$(ARM_TOOLCHAIN_PREFIX)gcc
    ARM_AR := $(ARM_TOOLCHAIN_PATH)/bin/$(ARM_TOOLCHAIN_PREFIX)ar
else
    ARM_CC := $(ARM_TOOLCHAIN_PREFIX)gcc
    ARM_AR := $(ARM_CROSS_COMPILE)ar
endif
# Get any compiler flags from the environment
ARM_CFLAGS = $(CFLAGS)
ARM_CFLAGS += -std=gnu99 \
    -Wdeclaration-after-statement -Wall -Wno-trigraphs \
    -fno-strict-aliasing -fno-common -fno-omit-frame-pointer \
    -c -O3
ARM_LDFLAGS = $(LDFLAGS)
ARM_LDFLAGS += -lm -lpthread
ARM_ARFLAGS = rcs
```

Details – Makefile – 'C6x

```
C6RUN_TOOLCHAIN_PREFIX=c6runlib-
ifdef C6RUN_TOOLCHAIN_PATH
    C6RUN_CC :=
        $(C6RUN_TOOLCHAIN_PATH)/bin/$(C6RUN_TOOLCHAIN_PREFIX)cc
    C6RUN_AR :=
        $(C6RUN_TOOLCHAIN_PATH)/bin/$(C6RUN_TOOLCHAIN_PREFIX)ar
else
    C6RUN_CC := $(C6RUN_TOOLCHAIN_PREFIX)cc
    C6RUN_AR := $(C6RUN_TOOLCHAIN_PREFIX)ar
endif

C6RUN_CFLAGS = -c -O3
C6RUN_ARFLAGS = rcs --C6Run:replace_malloc
```

Sharing Memory

- The ARM and the DSP share memory
 - Pass pointers to the buffers
- No need to copy from DSP to ARM
 - Very little overhead
- ARM uses a memory management unit (MMU)
 - maps virtual addr to physical addr
- DSP doesn't have an MMU!

Sharing Memory

- ARM pointers (*outputBuffer*, *inputBuffer*) are virtual addresses
- DSP pointers (*outputBuffer*, *inputBuffer*) are physical addresses
- C6Run automatically provides the code needed to map from the virtual address space to the physical

Sharing Memory

- Bigger problem
 - *outputBuffer* and *inputBuffer* are allocated at run time using the standard C routine **malloc**
 - malloc allocates contiguous memory of the desired size
- Contiguous in the *virtual* space
 - probably not contiguous in the *physical* space
- Causes problems for the DSP

Sharing Memory

- Solution:
 - loader flag **--C6Run:replace_malloc**
 - Tells the loader to replace malloc with **cmem**
- **cmem** is an API and library for managing one or more blocks of physically contiguous memory
- Provides address translation services
 - e.g. virtual to physical translation

Sharing Memory

- If you are uncomfortable with replacing all `mallocs` with `cmem`
 - Remove `--C6Run:replace_malloc`
 - Call `C6RUN_MEM_malloc(N * sizeof(short))` when you have memory to share with the DSP

cmem

- Where does `cmem` allocate the memory?
 - A section of the Beagle's RAM that Linux doesn't control

```
$ cat /proc/cmdline
console=tty0 console=ttyS2,115200n8
consoleblank=0 mpu=auto buddy=none
camera=lbc3m1 vram=24M omapfb.mode=dvi:hd720
mem=99M@0x80000000 mem=384M@0x88000000
omapfb.vram=0:12M,1:8M,2:4M
omapdss.def_disp=dvi root=/dev/mmcblk0p2 rw
rootfstype=ext3 rootwait
```

cmem

- 99M bytes start at 0x8000 0000 and 384M start at 0x8800 0000
- Where does the first block of addresses end?

```
$ bc
obase=16
99*1024*1024      6300000
ibase=16
6300000+80000000  86300000
```

- Block ends at 0x8630 0000

cmem

- How does `cmem` know that?
- Look in `~/c6run_build/loadmodules.sh`

```
DSP_REGION_START_ADDR="0x86300000"
DSP_REGION_END_ADDR="0x88000000"
CMEMK_OPTS="phys_start=$DSP_REGION_START_ADDR phys_end=$DSP_REGION_END_ADDR
allowOverlap=1"
modprobe cmemk ${CMEMK_OPTS}
```

cmem

- How does C6Run know what memory is used?
 - Look in `~/c6run_build/environment.sh`
 - At the bottom you'll see:

```
PLATFORM_CFLAGS='
-DDSP_REGION_BASE_ADDR=0x86300000
-DDSP_REGION_CMEM_SIZE=0x01000000
-DDSP_REGION_CODE_SIZE=0x00D00000
-DLPM_REQUIRED
-DDSP_HAS_MMU'
```

Looking inside c6run

- Edit Makefile
 - add `--C6Run:debug` to the FLAGS:

```
C6RUN_CFLAGS = -c -O3 -D_DEBUG --C6Run:debug
C6RUN_ARFLAGS = rcs --C6Run:replace_malloc --C6Run:debug
```

- Then recompile everything:

```
$ make dsp_clean
$ make dsp_exec
$ ls -sh
```

- Many interesting files appear

ARM vs. DSP

- How do I specify what runs on the ARM and what runs on the DSP?
- **Makefile**
 - `main_cfft.c` is run on the ARM
 - `cfft.c` is run on the DSP

Details – make dsp_exec

```
EXEC_SRCS := main_cfft.c
EXEC_DSP_OBJS := $(EXEC_SRCS:%.c=dsp/%.o)

dsp_exec: dsp/.created dsp_lib $(EXEC_DSP_OBJS)
    $(ARM_CC) $(ARM_LDFLAGS) $(CINCLUDES) -o cfft_dsp \
        dsp/main_cfft.o cfft_dsp.lib
dsp_lib: dsp_lib/.created $(LIB_DSP_OBJS)
    $(C6RUN_AR) $(C6RUN_ARFLAGS) cfft_dsp.lib dsp_lib/cfft.o
dsp/%.o : %.c
    $(ARM_CC) $(ARM_CFLAGS) $(CINCLUDES) -o $$ $<
dsp_lib/%.o : %.c
    $(C6RUN_CC) $(C6RUN_CFLAGS) $(CINCLUDES) -o $$ $<
dsp/.created:
    @mkdir -p dsp
    @touch dsp/.created
dsp_lib/.created:
    @mkdir -p dsp_lib
    @touch dsp_lib/.created
```

Details – make gpp_exec

```
EXEC_SRCS := main_cfft.c
EXEC_DSP_OBJS := $(EXEC_SRCS:%.c=gpp/%.o)

gpp_exec: gpp/.created gpp_lib $(EXEC_ARM_OBJS)
    $(ARM_CC) $(ARM_LDFLAGS) $(CINCLUDES) -o cfft_arm \
        gpp/main_cfft.o cfft_arm.lib
gpp_lib: gpp_lib/.created $(LIB_ARM_OBJS)
    $(ARM_AR) $(ARM_ARFLAGS) cfft_arm.lib gpp_lib/cfft.o
gpp/%.o : %.c
    $(ARM_CC) $(ARM_CFLAGS) $(CINCLUDES) -o $$ $<
gpp_lib/%.o : %.c
    $(ARM_CC) $(ARM_CFLAGS) $(CINCLUDES) -o $$ $<
gpp/.created:
    @mkdir -p gpp
    @touch gpp/.created
gpp_lib/.created:
    @mkdir -p gpp_lib
    @touch gpp_lib/.created
```

Conclusions

- Using the DSP is easy as
 - Moving the functions to run on the DSP into their own files
 - Editing the Makefile to show what runs where